



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Computer Vision

Chapter 4: Feature detection and Image
matching

Plan

- Edges
 - Detection
 - linking
- Feature extraction
 - Global features
 - Local features
 - Matching
- Applications

Local features vs Global features

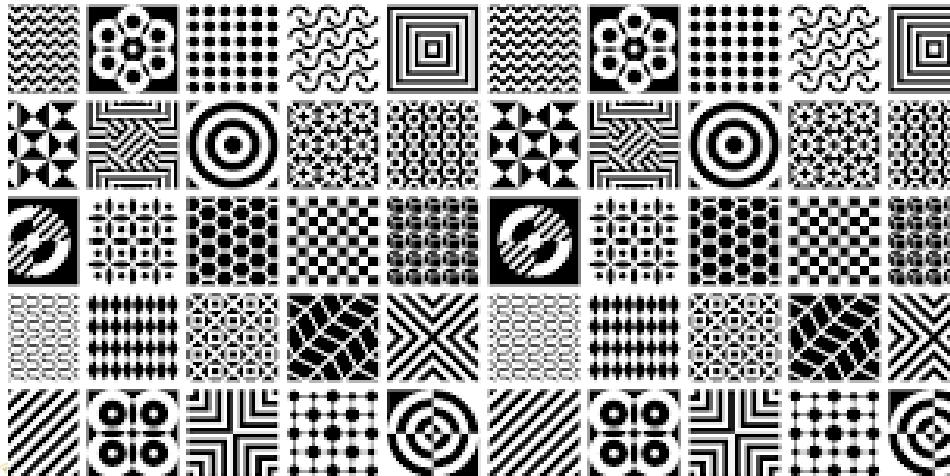
- Two types of features are extracted from the image:
 - local and global features (descriptors)
- **Global features**
 - Describe the **image as a whole** to generalize the entire object
 - Include contour representations, shape descriptors, and texture features
 - Examples: Invariant Moments (Hu, Zernike), Histogram Oriented Gradients (HOG), PHOG, and Co-HOG,...
- **Local feature:**
 - the local features **describe the image patches** (key points in the image) of an object
 - represents the texture/color in an image patch
 - Examples: SIFT, SURF, LBP, BRISK, MSER and FREAK, ...

Feature extraction

- **Global features**
 - Color / Shape / Texture
- Local features

Global features?

How to distinguish these objects?



Types of features

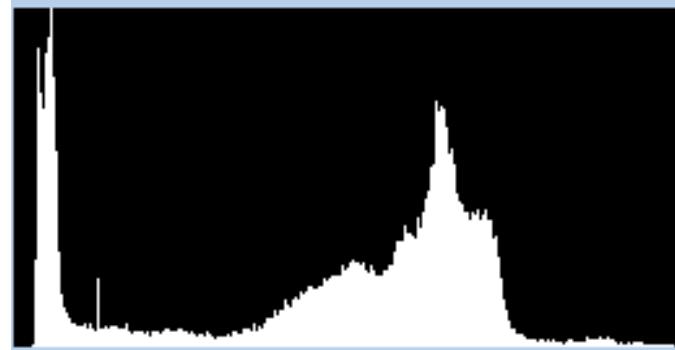
- Contour representation, Shape features
- Color descriptors
- Texture features

Color features

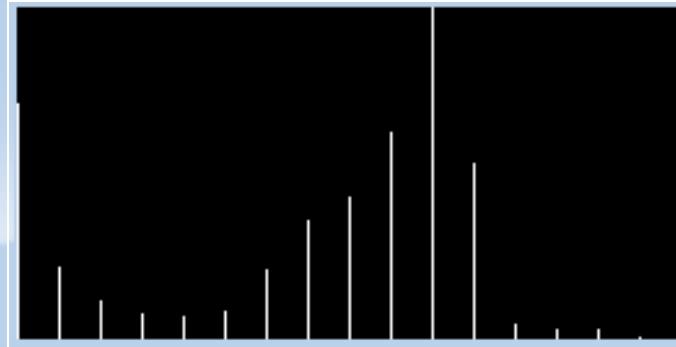
- Histogram



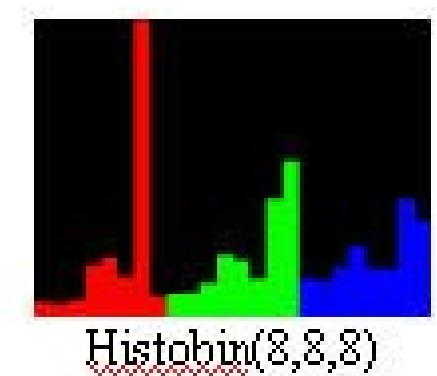
256 bins intensity histogram



16 bins intensity histogram



Image



Distance / Similarity

- L1 ou L2 (euclidian) distances are often used

$$d_{L1}(H, G) = \sum_{i=1}^N |h_i - g_i|$$

- Histogram intersection

$$\cap(H, G) = \frac{\sum_i \min(h_i, g_i)}{\sum_i g_i}$$

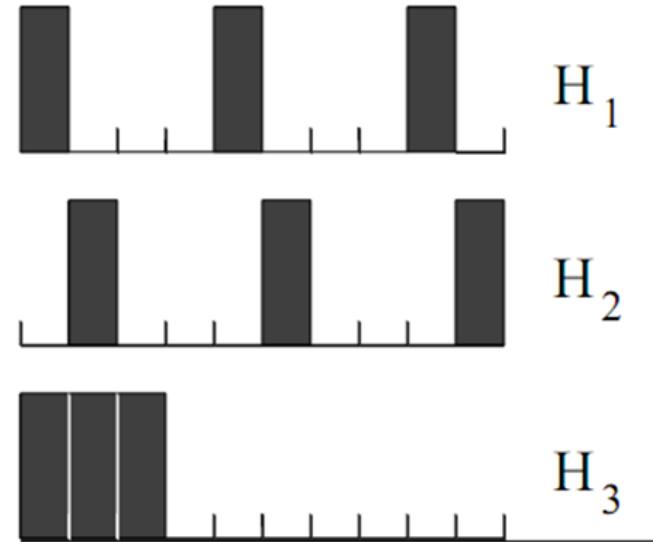
Advantages of histogram

- Invariant to basic geometric transformations:
 - Rotation
 - Translation
 - Zoom (Scaling)

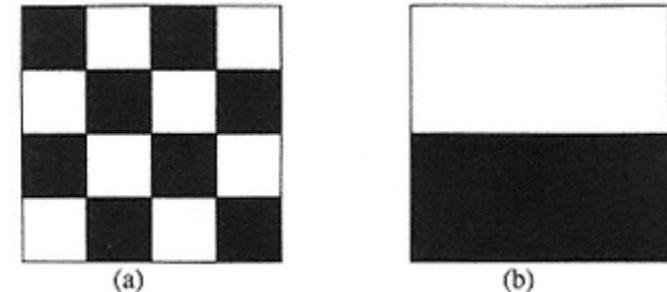


Some inconveniences

- The similarity between colors in adjacent colors (bin) is not taken into account

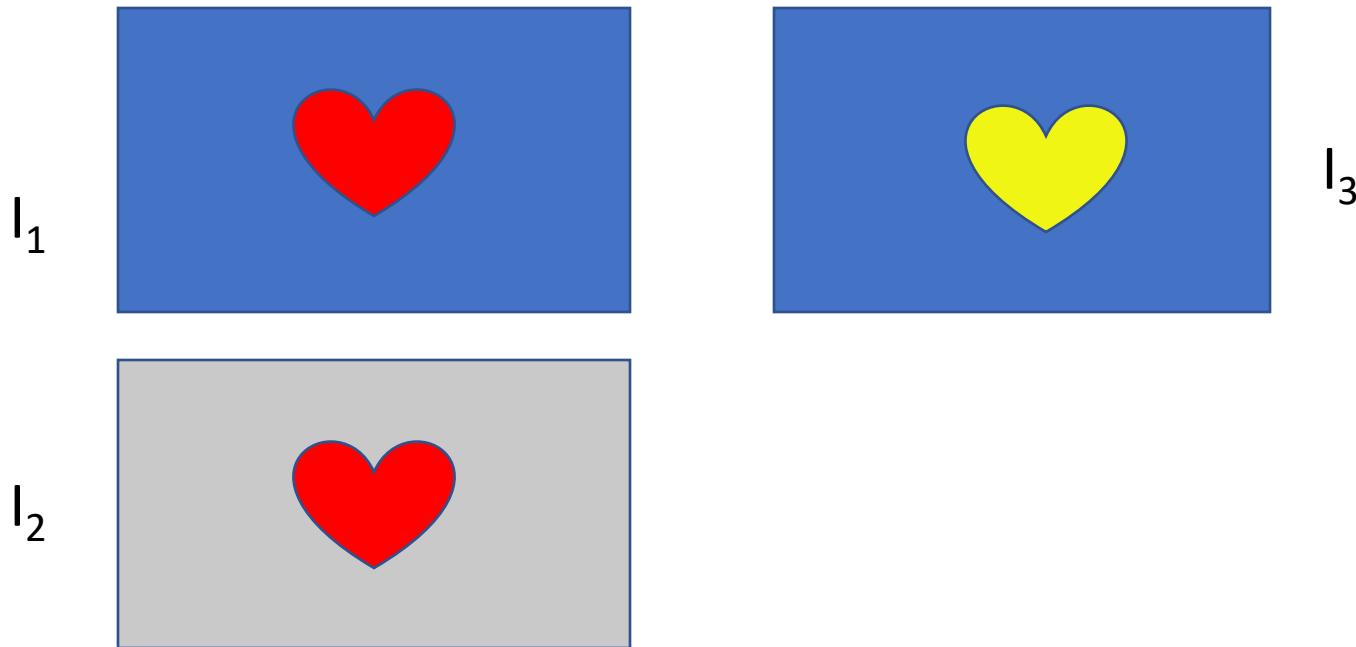


- The spatial distribution of pixel values is not considered: 2 different images may have the same histogram



Some inconveniences

- Background effect: $d(I_1, I_2) \neq d(I_1, I_3)$

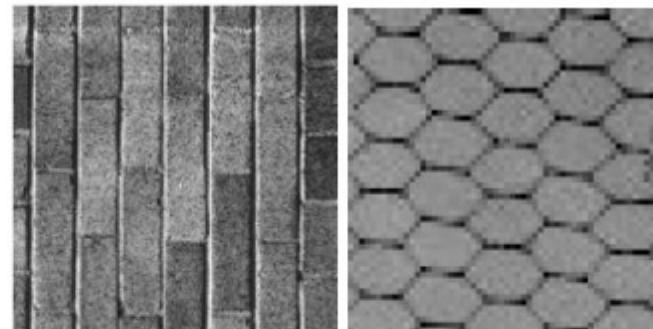


- Color representation dependency (color space), device dependency, ...

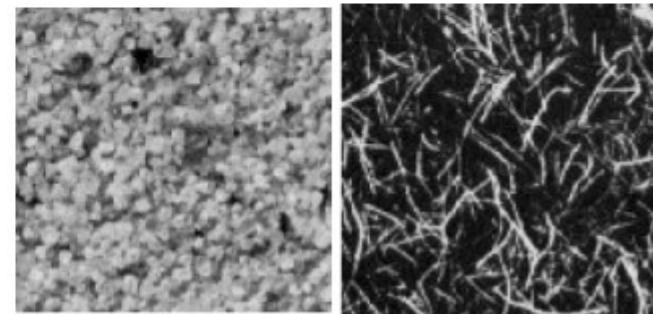
Texture features

- A **texture** can be defined as
 - a region with variation of intensity
 - as a spatial organization of pixels

A texture can be
periodic (a pattern that
repeats itself) ...



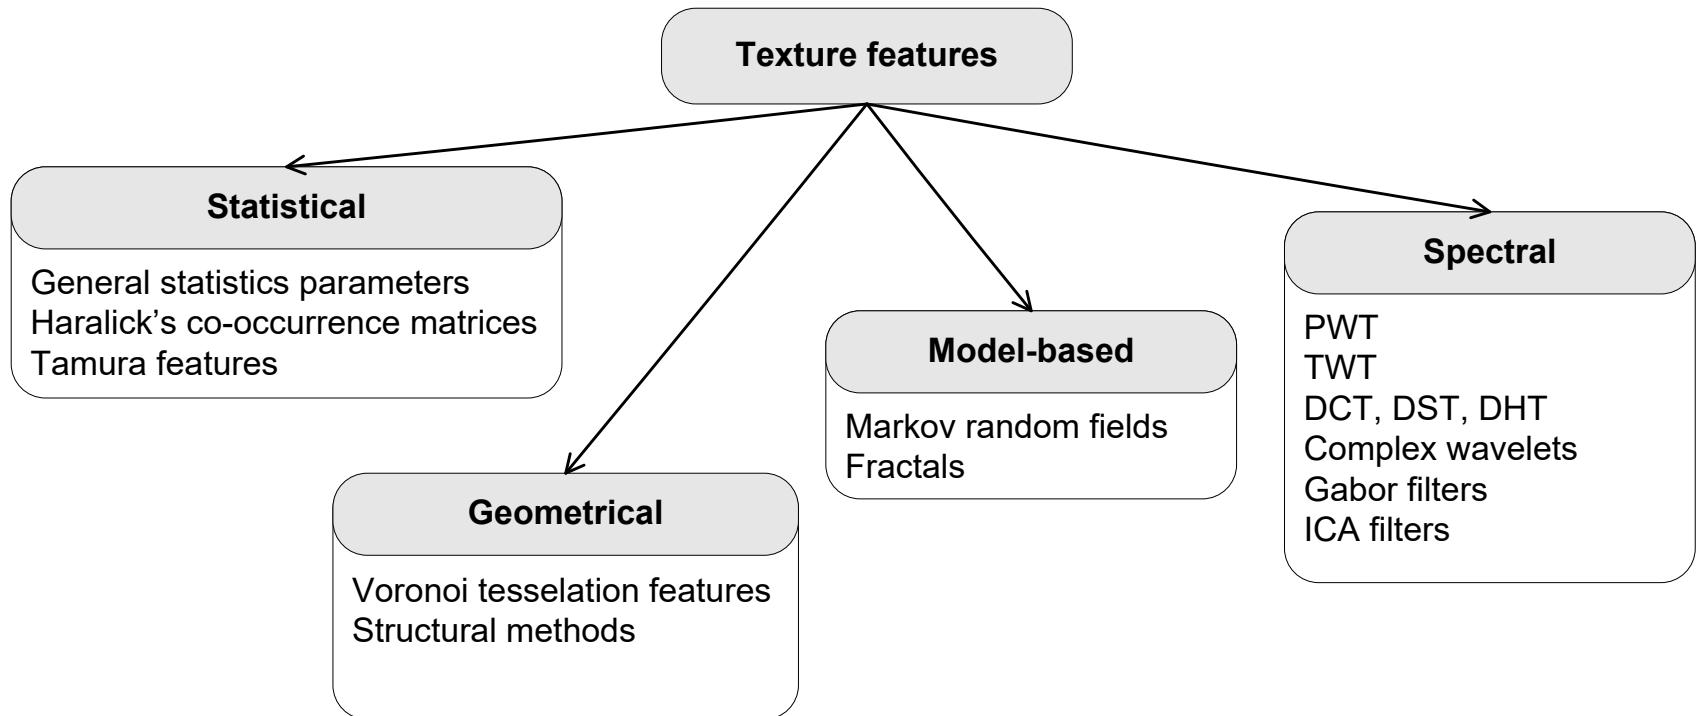
...or **non-periodic** (no
pattern, more
disorganized)



Texture features

- There are several methods for analyzing textures:
 - First order statistics
 - Statistics on histogram
 - Co-occurrence matrices
 - Searching patterns
 - Frequential analysis
 - Gabor filter
 - ...
- The most difficult is to find a good representation (good parameters) for each texture

Texture features



First order statistics

- Histogram-based: mean, variance, skewness, kurtosis, energy, entropy, ...

$$\mu = \frac{1}{n} \sum_{i=0}^{H_g} h(i)$$

$$E = \sum_{i=0}^{H_g} (P(i))^2$$

$$\sigma^2 = \frac{1}{n-1} \sum_{i=0}^{H_g} (h(i) - \mu)^2$$

$$H = - \sum_{i=0}^{H_g} P(i) \log(P(i))$$

$$s = \frac{1}{n\sigma^3} \sum_{i=0}^{H_g} (h(i) - \mu)^3$$

$$P(i) = \frac{h(i)}{n}$$

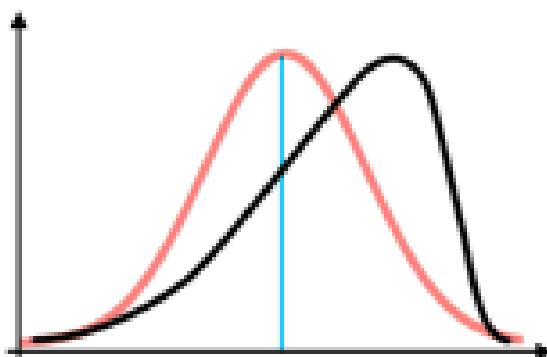
$$k = \frac{1}{n\sigma^4} \sum_{i=0}^{H_g} (h(i) - \mu)^4$$

h(i) : số điểm ảnh ở mức xám i

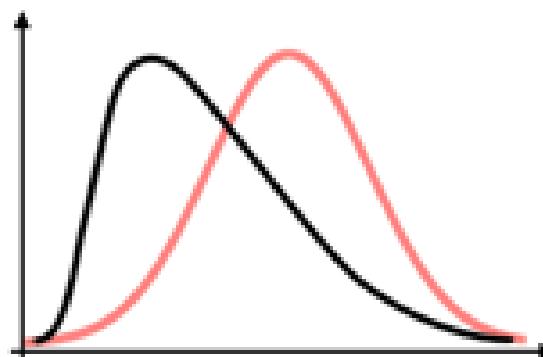
First order statistics

- Skewness vs. kurtosis (red: normal distribution with mean = 5, variance = 4)

© www.scratchapixel.com



Negative Skew
(large tail to the left)



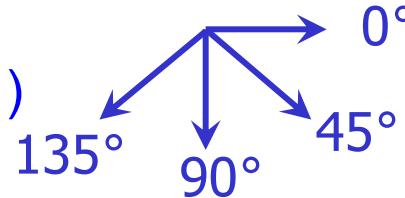
Positive Skew
(large tail to the right)



positive and negative kurtosis

GLCM (Grey Level Co-occurrence Matrices)

- The idea here is to identify **gray level that repeat themselves** given a **distance** and a **direction**
 - *Co-occurrence matrices (Haralick)*
- **Matrix of size $Ng \times Ng$**
 - Ng is the number of gray level in the image (256x256)
 - We often reduce that number to 8x8, 16x16 or 32x32
- Many matrices, one for each distance and direction
 - **Distance** : 1, 2, 3 (,4, ...)
 - **Direction** : 0° , 45° , 90° , 135° (, ...)
- Processing time can be very long



GLCM

$$CM_{d,\beta}(c_i, c_j) = \frac{card(\{p_1, p_2 | I(p_1) = c_i, I(p_2) = c_j, N_{d,\beta}(p_1, p_2) = true\})}{card(\{p1, p2 | N_{d,\beta}(p_1, p_2) = true\})}$$

$N_{d,\beta}(p_1, p_2) = true$ p_2 is a neighbor of p_1 at a distance d direction β

GLCM

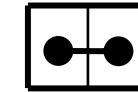
- Example on how to compute these matrices:

1	4	4	3
4	2	3	2
1	2	1	4
1	2	2	3

Image

	1	2	3	4
1	?	?	?	?
2	?	?	?	?
3	?	?	?	?
4	?	?	?	?

*Matrix for distance=1
and direction=0°*



We loop over the image and for each pair of pixels following the given distance and orientation, we increment the co-occurrence matrix

GLCM

- Example on how to compute these matrices:

1	4	4	3
4	2	3	2
1	2	1	4
1	2	2	3

Image

	1	2	3	4
1	0	0	0	1
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

*Matrix for distance=1 and
direction=0°*

Pair of neighbor pixels (1,4)

GLCM

- Example on how to compute these matrices:

1	4	4	3
4	2	3	2
1	2	1	4
1	2	2	3

Image

	1	2	3	4
1	0	0	0	1
2	0	0	0	0
3	0	0	0	0
4	0	0	0	1

Matrix for distance=1 and direction=0°

Pair of neighbor pixels (4,4)

GLCM

- Example on how to compute these matrices:

1	4	4	3
4	2	3	2
1	2	1	4
1	2	2	3

Image

	1	2	3	4
1	0	0	0	1
2	0	0	1	0
3	0	1	0	0
4	0	1	1	1

Matrix for distance=1 and direction=0°

...and so on (after 2 lines)

GLCM

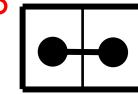
- Example on how to compute these matrices (final):

1	4	4	3
4	2	3	2
1	2	1	4
1	2	2	3

Image

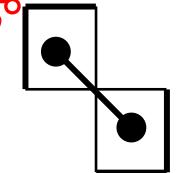
	1	2	3	4
1	0	2	0	2
2	1	1	2	0
3	0	1	0	0
4	0	1	1	1

*Matrix for distance=1
and direction=0°*



	1	2	3	4
1	0	2	1	0
2	1	1	0	0
3	0	0	0	1
4	0	2	1	0

*Matrix for distance=1
and direction=45°*



...and so on for each matrix (several matrices at the end)

GLCM

- Most important/popular parameters computed from GLCM:

$$Energy = \sum_i \sum_j CM_d^2(i, j) \quad \text{minimal when all elements are equal}$$

$$entropy = -\sum_i \sum_j CM_d(i, j) \log(CM_d(i, j)) \quad \begin{array}{l} \text{a measure of chaos,} \\ \text{maximal when all elements are equal} \end{array}$$

$$contrast = \sum_i \sum_j (i - j)^2 CM_d(i, j) \quad \begin{array}{l} \text{small values when big elements} \\ \text{are near the main diagonal} \end{array}$$

$$idm = \sum_i \sum_j \frac{1}{1 + (i - j)^2} CM_d(i, j) \quad \begin{array}{l} idm \text{ (inverse differential moment) has small} \\ \text{values when big elements are far from the} \\ \text{main diagonal} \end{array}$$

GLCM

- Haralick features:

- For each GLCM, we can compute up to **14 (13)** **parameters** characterizing the texture, of which the most important : mean, variance, energy, inertia, entropy, inverse differential moment
- Ref: <http://haralick.org/journals/TexturalFeatures.pdf>

Invariances

- Rotation?
 - Average on all directions
- Scaling?
 - Multi-resolutions



Texture features comparision

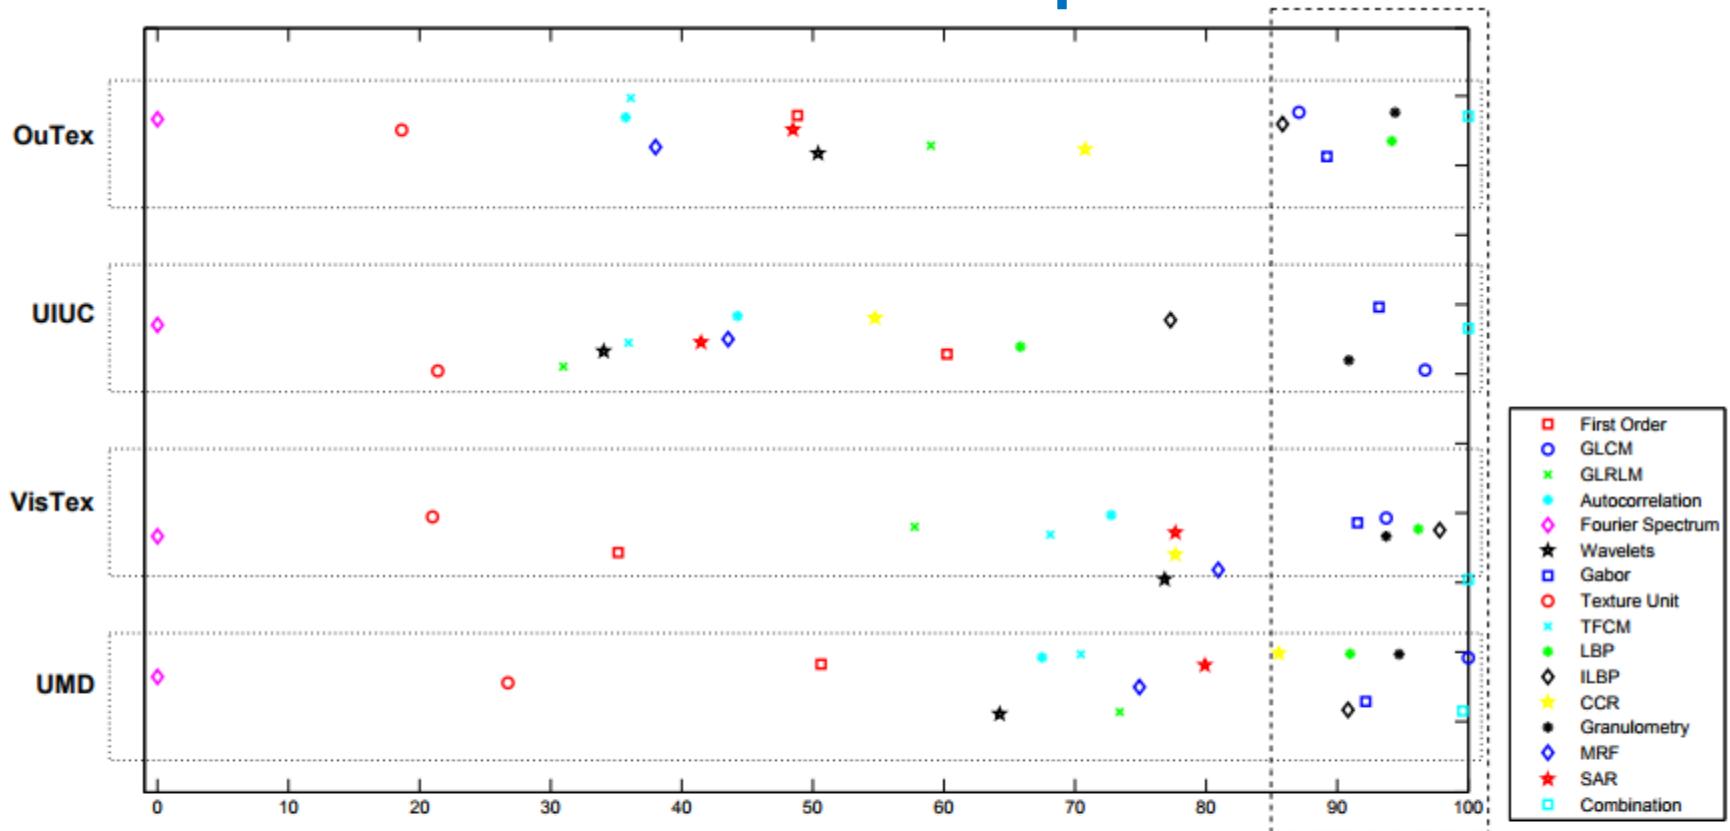
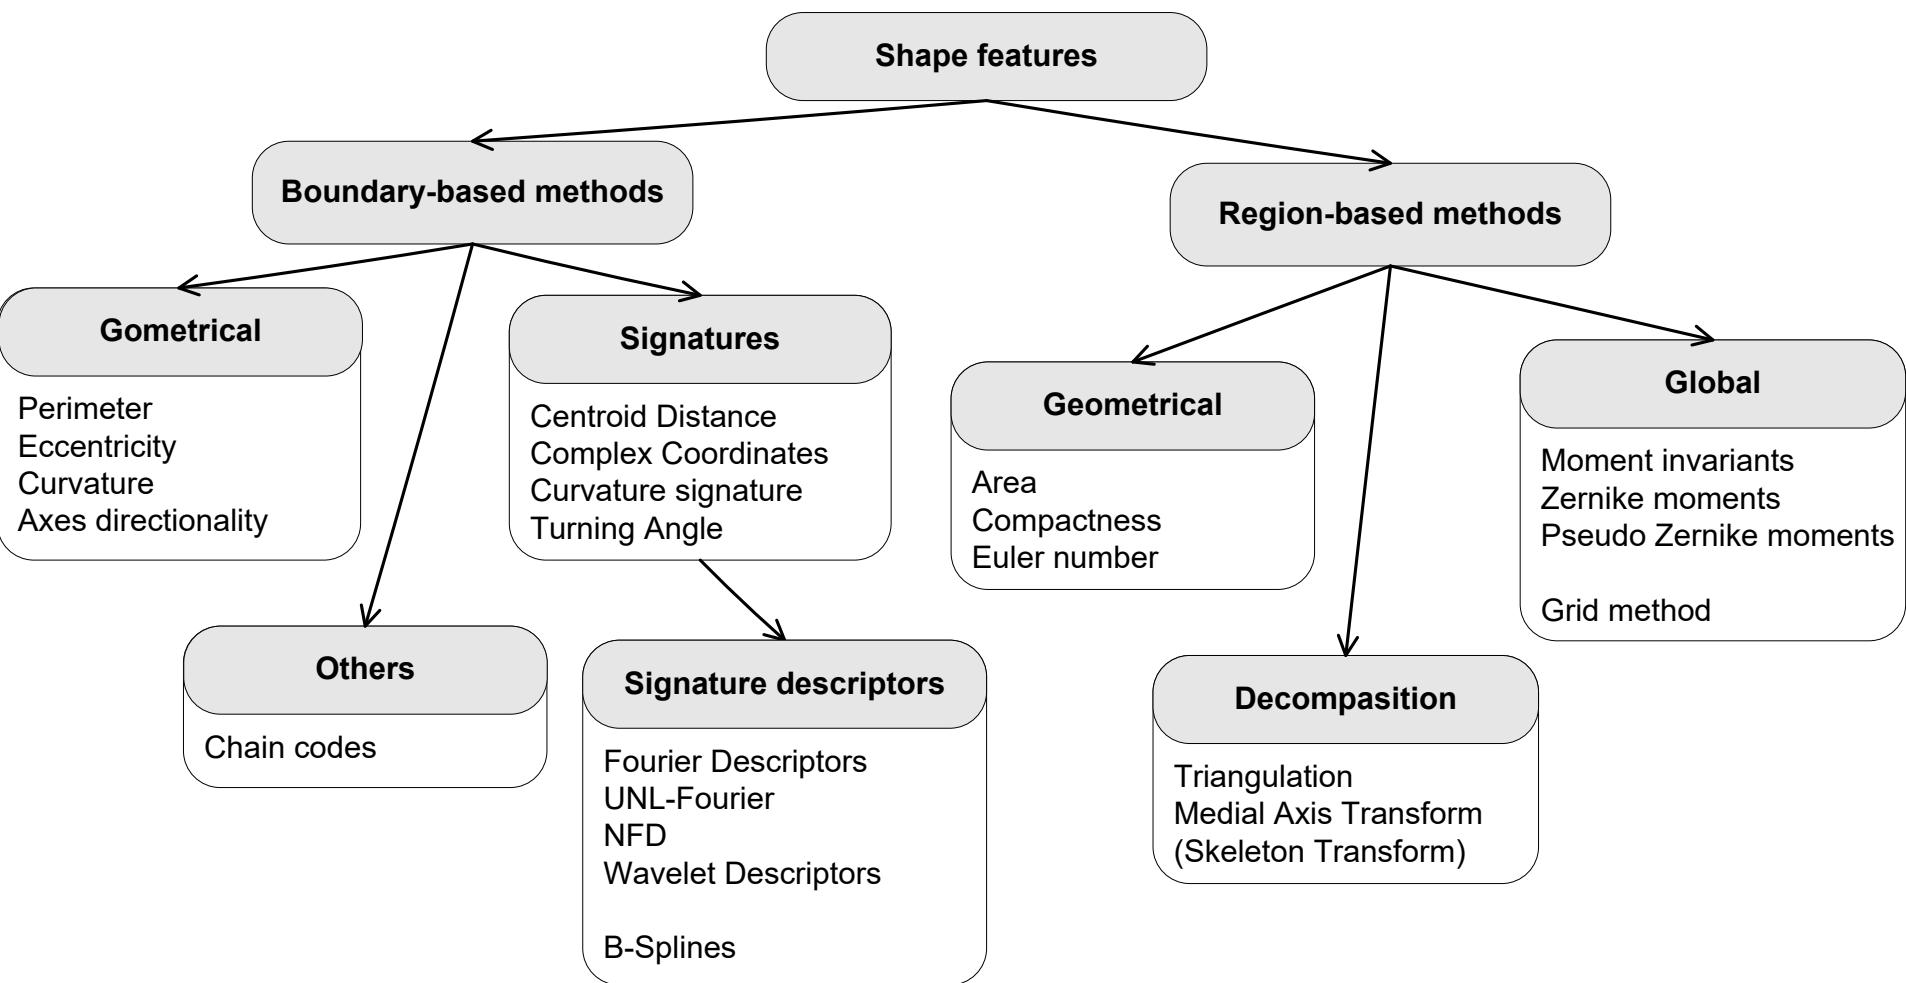


Figure 10. Comparison of results for all considered features on each data set. Each horizontal dotted box displays results for one data set. The x -axis shows normalized results (between 0 and 100) according to the minimum and maximum classification rates achieved for each data set. The vertical slashed box on the right-hand side selects the feature extraction methods that achieved the best results throughout all data sets.

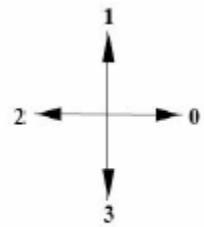
Shape features

- Contour-based features
 - Chain coding, polygon approximation, geometric parameters, angular profile, surface, perimeter, ...
- Region based:
 - Invariant moments, ...

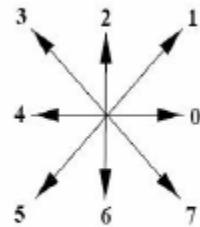
Shape features



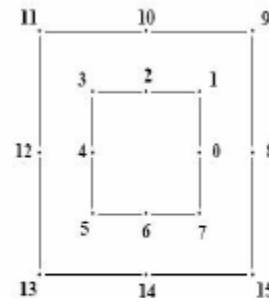
Examples: Freeman chain coding



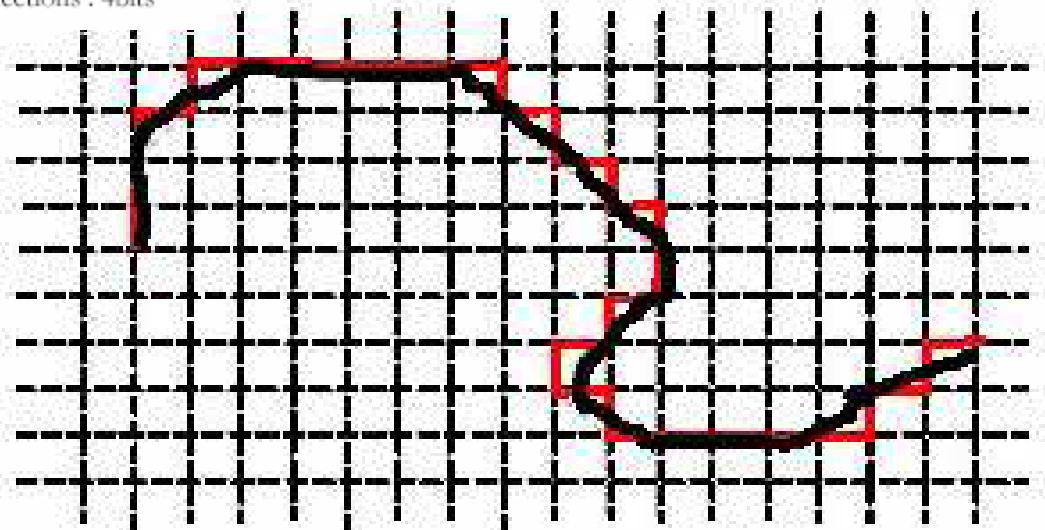
4 directions : 2bits



8 directions : 3bits



16 directions : 4bits



111010000003030303232303000001010

$34 \times 2 = 68\text{bits}$

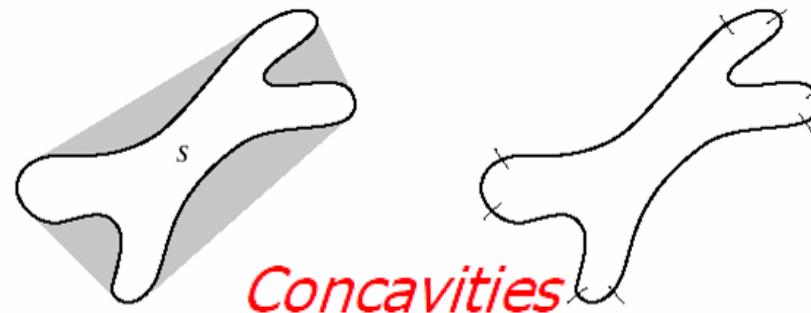
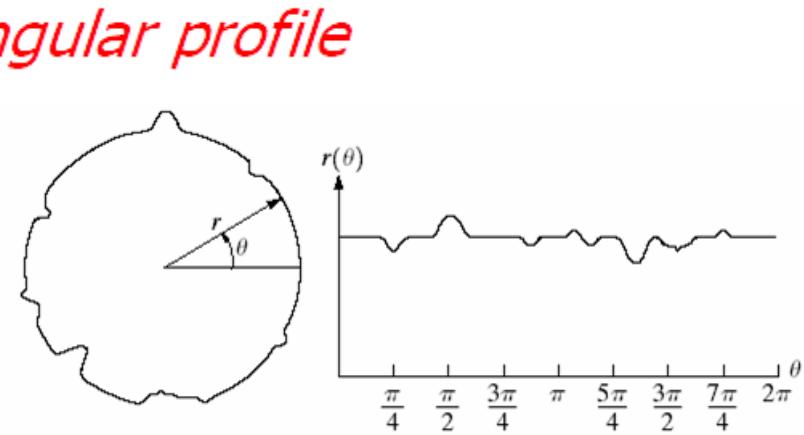
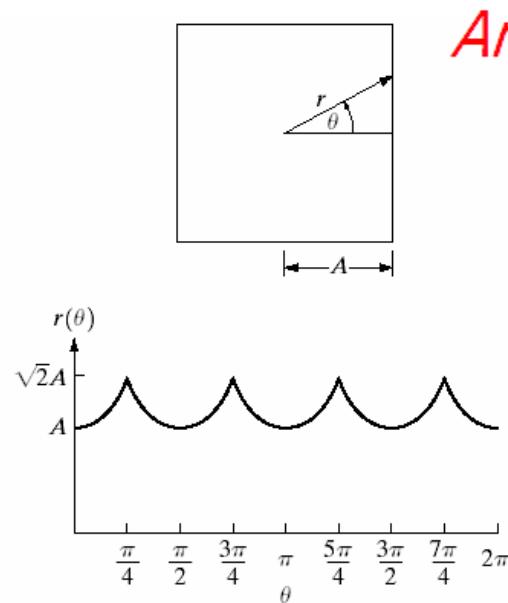
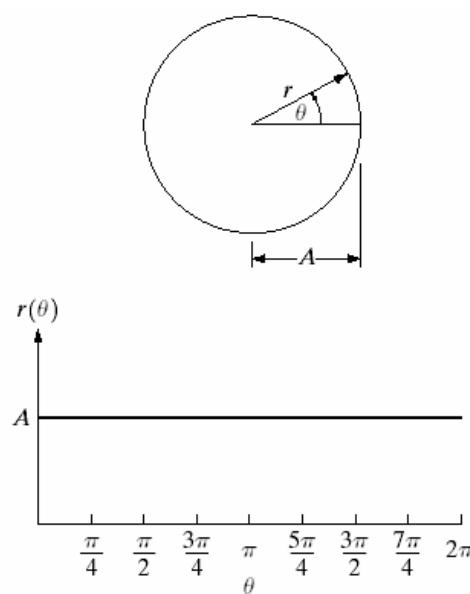
221100007777654670000101

$24 \times 3 = 72\text{bits}$

1098815156546788101

$16 \times 4 = 64\text{bits}$

Examples: angular profile, ...



Concavities

Examples : Image moments

- Moment

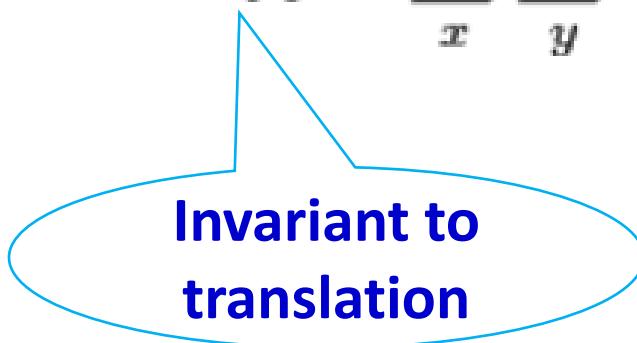
$$M_{ij} = \sum_x \sum_y x^i y^j I(x, y)$$

$M_{0,0}$ = area of the region D

$(M_{0,1}, M_{1,0})$ = centroid of D

- Central moments:

$$\mu_{pq} = \sum_x \sum_y (x - \bar{x})^p (y - \bar{y})^q I(x, y)$$



$$\bar{x} = \frac{M_{10}}{M_{00}} \quad \bar{y} = \frac{M_{01}}{M_{00}}$$

Invariant moments (Hu's moments)

invariant to
translation,
scale, and
rotation, and
reflection

$$\eta_{ij} = \frac{\mu_{ij}}{\mu_{00}^{\left(1+\frac{i+j}{2}\right)}}$$
$$\begin{aligned} I_1 &= \eta_{20} + \eta_{02} \\ I_2 &= (\eta_{20} - \eta_{02})^2 + 4\eta_{11}^2 \\ I_3 &= (\eta_{30} - 3\eta_{12})^2 + (3\eta_{21} - \eta_{03})^2 \\ I_4 &= (\eta_{30} + \eta_{12})^2 + (\eta_{21} + \eta_{03})^2 \\ I_5 &= (\eta_{30} - 3\eta_{12})(\eta_{30} + \eta_{12})[(\eta_{30} + \eta_{12})^2 - 3(\eta_{21} + \eta_{03})^2] + (3\eta_{21} \\ &\quad - \eta_{03})(\eta_{21} + \eta_{03})[3(\eta_{30} + \eta_{12})^2 - (\eta_{21} + \eta_{03})^2] \\ I_6 &= (\eta_{20} - \eta_{02})[(\eta_{30} + \eta_{12})^2 - (\eta_{21} + \eta_{03})^2] \\ I_7 &= (3\eta_{21} - \eta_{03})(\eta_{30} + \eta_{12})[(\eta_{30} + \eta_{12})^2 - 3(\eta_{21} \\ &\quad + \eta_{03})^2] + (\eta_{30} - 3\eta_{12})(\eta_{21} + \eta_{03})[3(\eta_{30} \\ &\quad + \eta_{12})^2 - (\eta_{21} + \eta_{03})^2] \end{aligned}$$

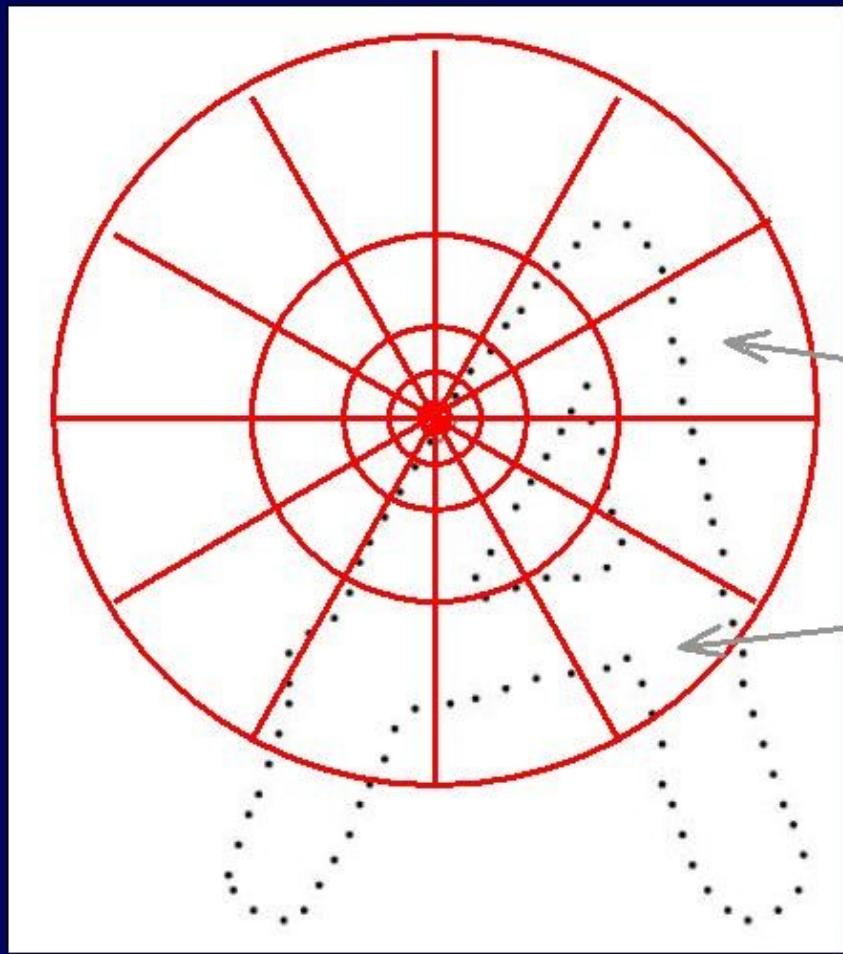
Examples : Hu's moments

6 images and their Hu Moments

id	Image	H[0]	H[1]	H[2]	H[3]	H[4]	H[5]	H[6]
K0		2.78871	6.50638	9.44249	9.84018	-19.593	-13.1205	19.6797
S0		2.67431	5.77446	9.90311	11.0016	-21.4722	-14.1102	22.0012
S1		2.67431	5.77446	9.90311	11.0016	-21.4722	-14.1102	22.0012
S2		2.65884	5.7358	9.66822	10.7427	-20.9914	-13.8694	21.3202
S3		2.66083	5.745	9.80616	10.8859	-21.2468	-13.9653	21.8214
S4		2.66083	5.745	9.80616	10.8859	-21.2468	-13.9653	-21.8214

<https://www.learnopencv.com/wp-content/uploads/2018/12/HuMoments-Shape-Matching.png>

Shape Context



Count the number of points inside each bin, e.g.:

Count = 4

:

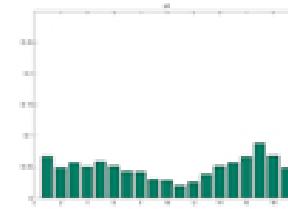
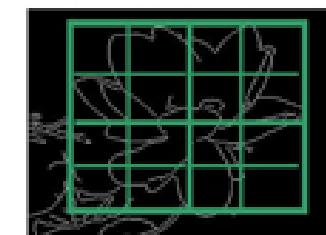
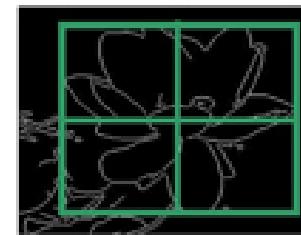
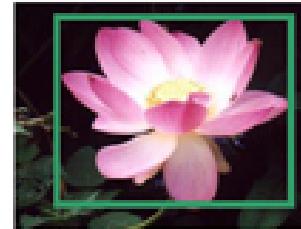
Count = 10

☞ Compact representation of distribution of points relative to each point

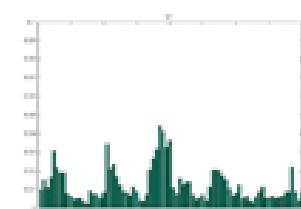
Examples: PHOG

PHOG:
Pyramid Histogram of Oriented Gradients

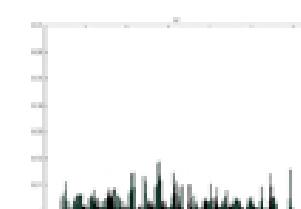
Input Image (image.jpg)



+



+



Output PHOG descriptor (image.jpg.txt)

Feature extraction

- Global features
- **Local features**
 - Interest point detector
 - Local descriptor

Why local features?

- Image matching: a challenging problem



Source: CS131 - Juan Carlos Niebles and Ranjay Krishna

Image matching



by [Diva Sian](#)



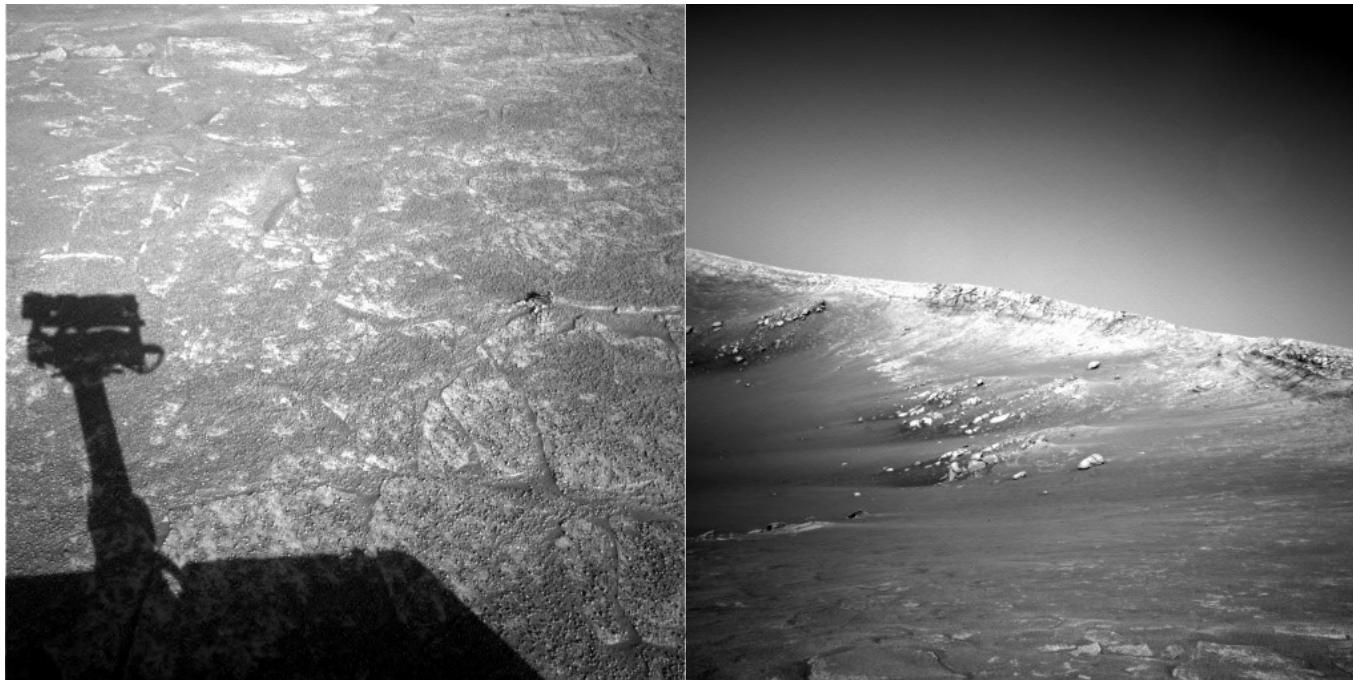
by [swashford](#)



by [scgbt](#)

Slide credit: Steve Seitz

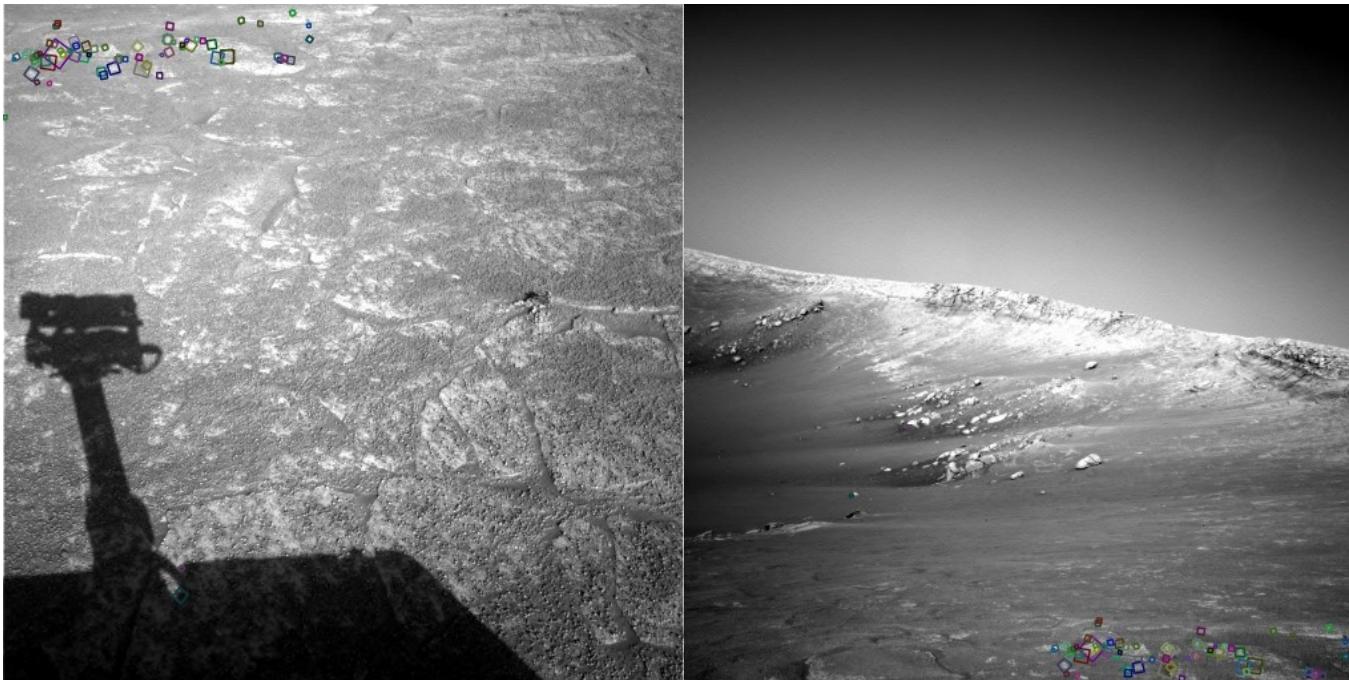
Harder Still?



NASA Mars Rover images

Slide credit: Steve Seitz

Answer Below (Look for tiny colored squares)



NASA Mars Rover images with SIFT feature matches
(Figure by Noah Snavely)

Slide credit: Steve Seitz

- Recognition of specific objects/scenes



Sivic and Zisserman, 2003



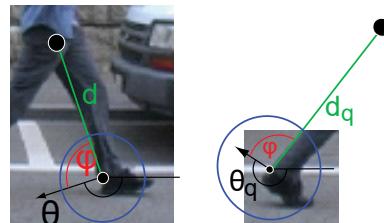
D. Lowe 2002

Motivation for using local features

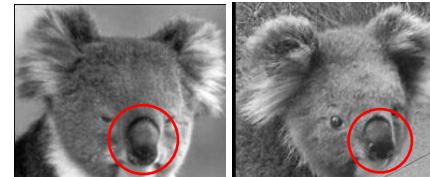
- Global representations have major limitations
- Instead, describe and match only local regions
- Increased robustness to
 - Occlusions



- Articulation



- Intra-category variations



Source: CS131 - Juan Carlos Niebles and Ranjay Krishna

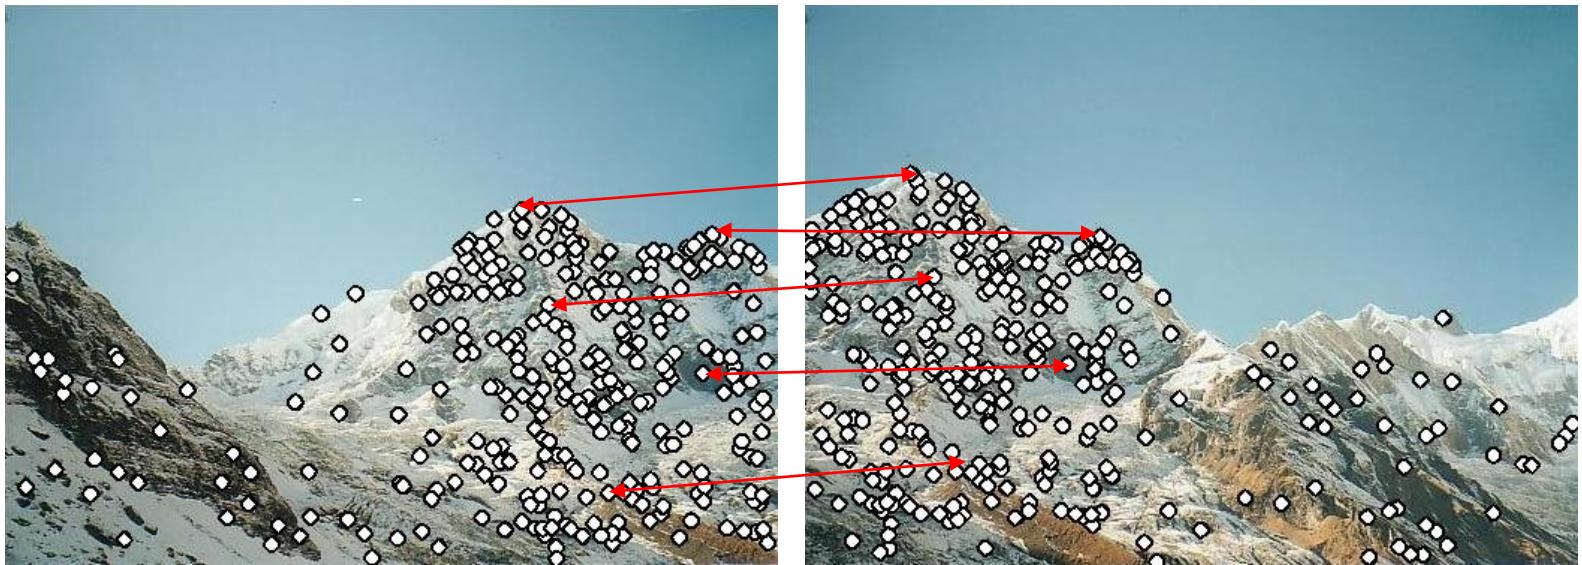
Local features and alignment

- We need to match (align) images
- Global methods sensitive to occlusion, lighting, parallax effects. So look for local features that match well.
- How would you do it by eye?



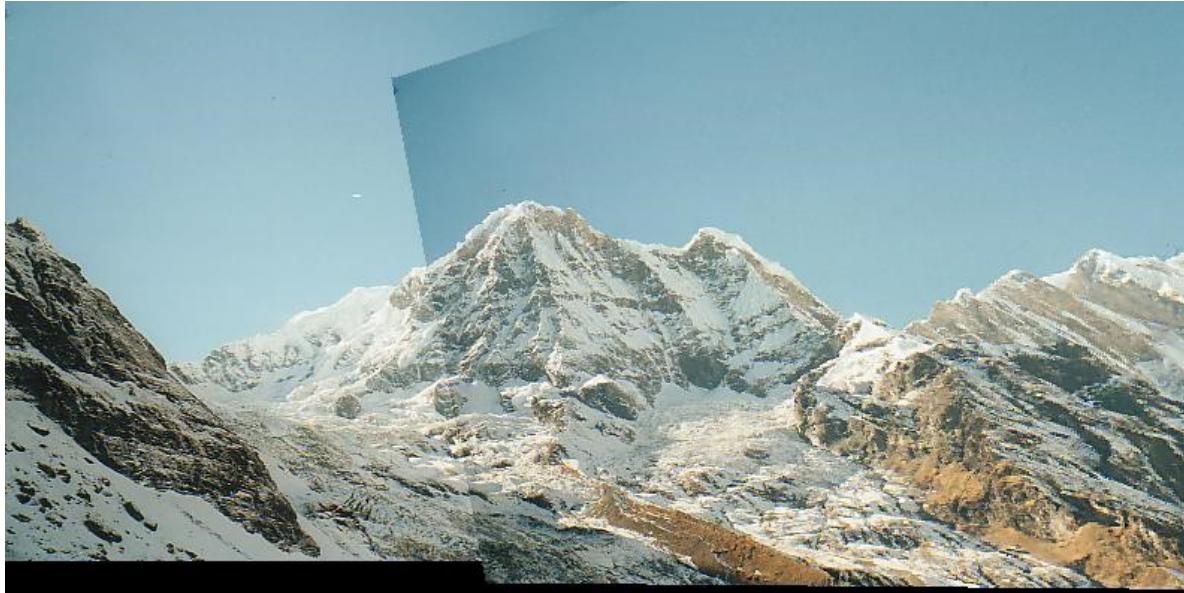
Local features and alignment

- Detect feature points in both images
- Find corresponding pairs

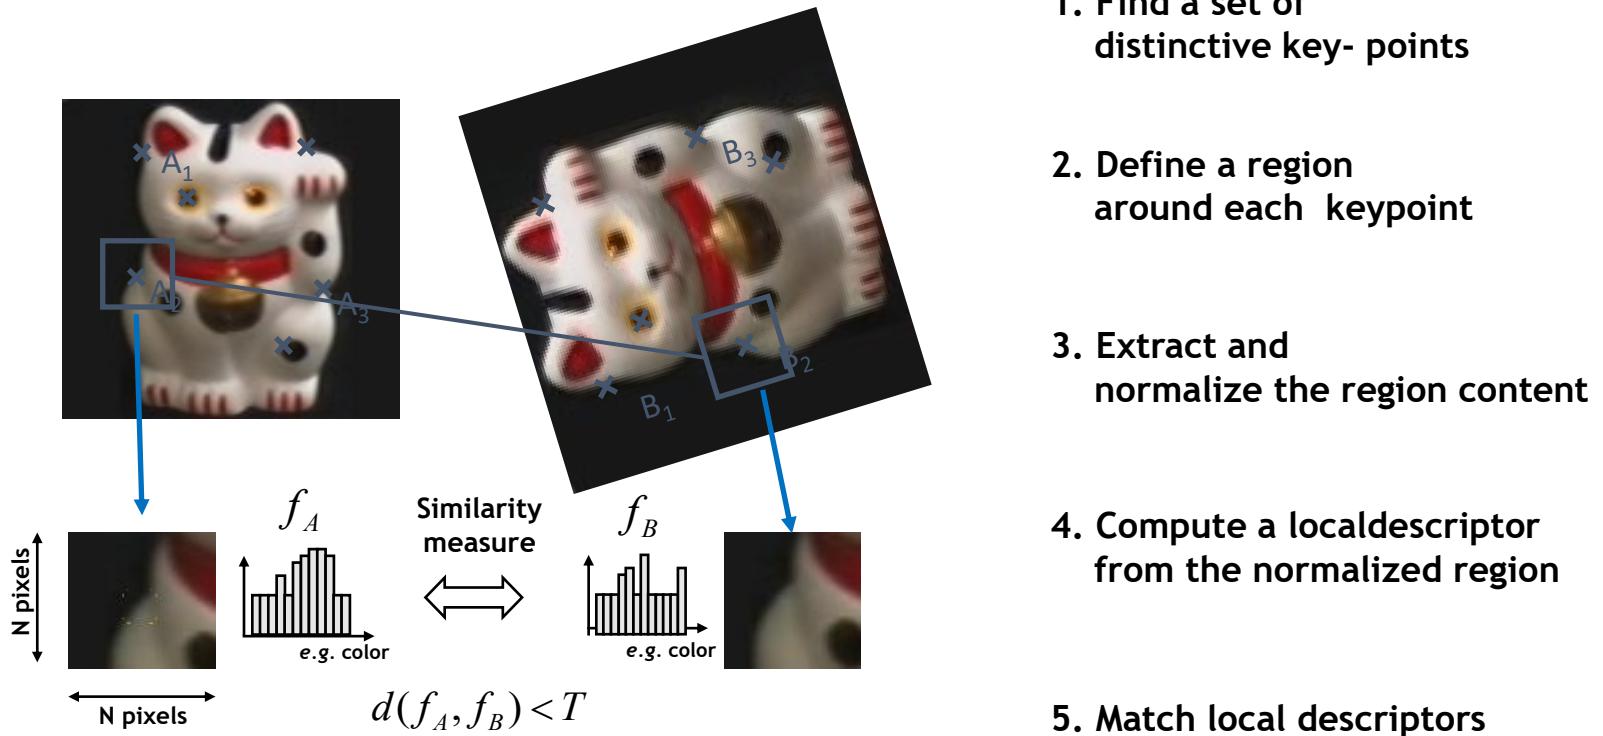


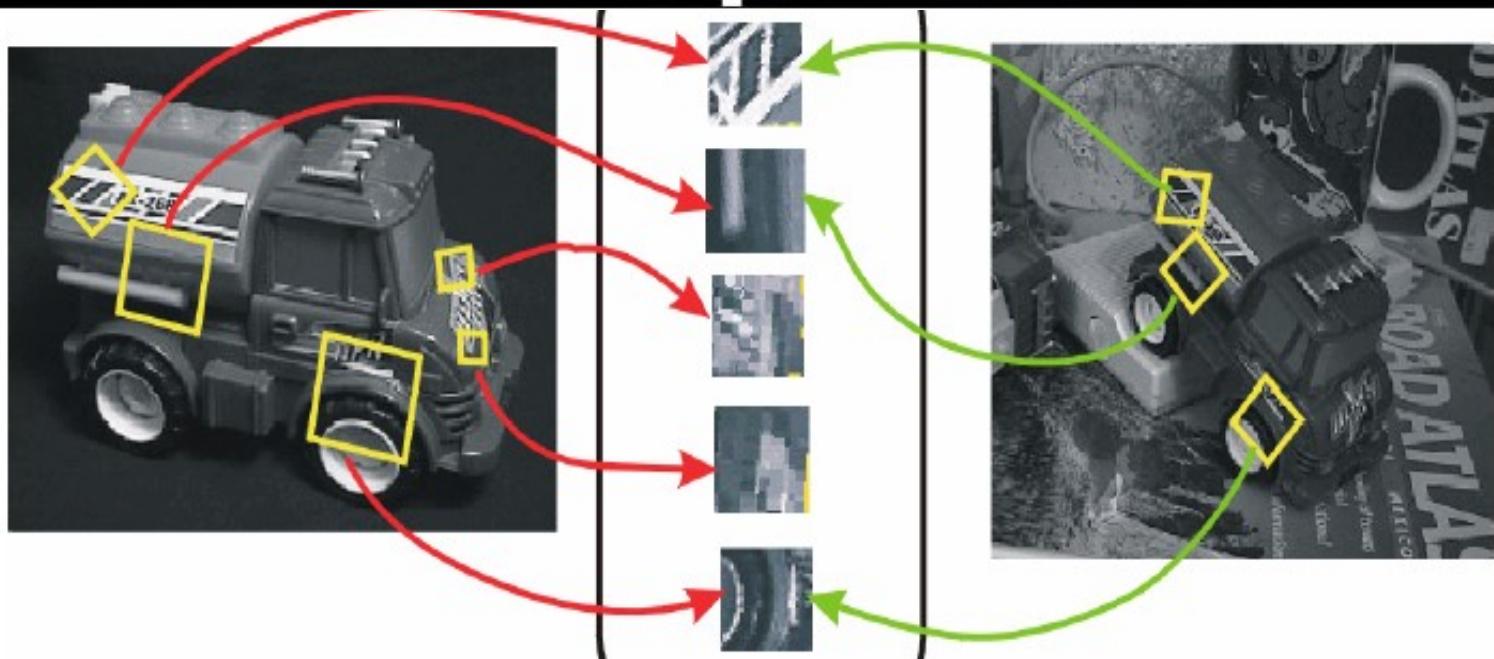
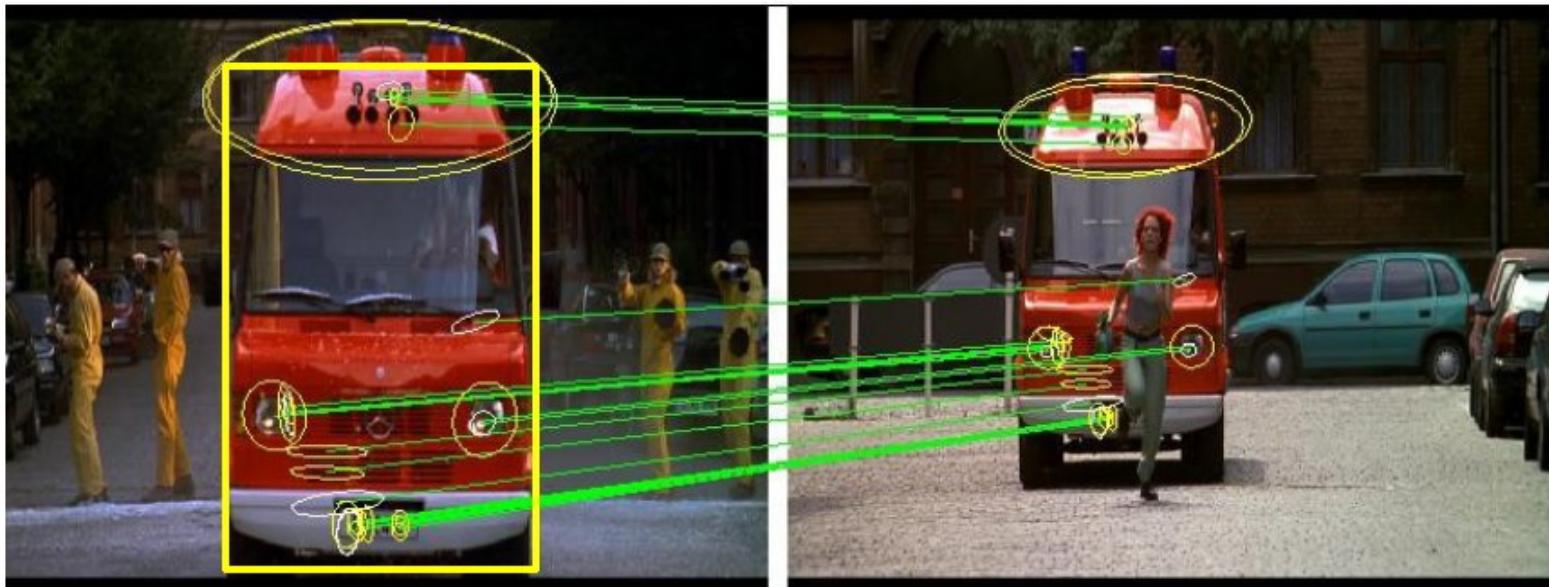
Local features and alignment

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images



Local features for image matching



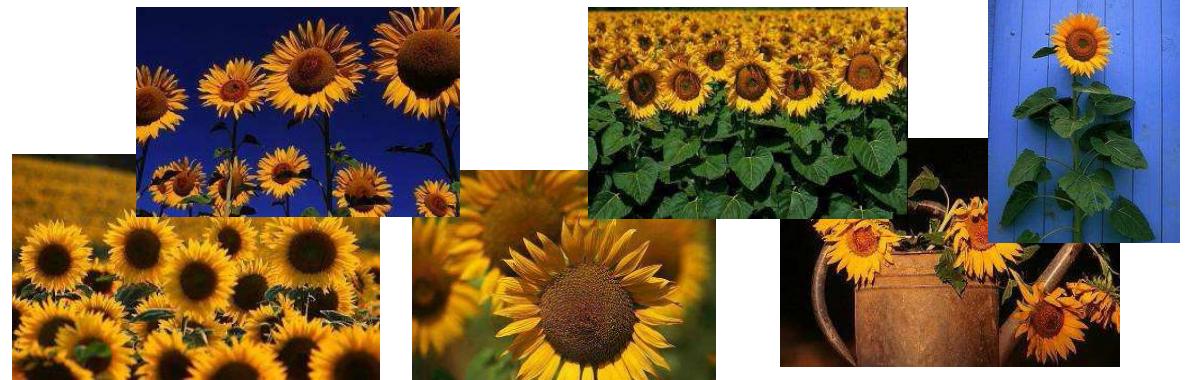


Source : Jim Little, Lowe: features, UBC.

Local features

- Objectifs:
 - Look for similar objects/regions
 - Partial query

Look for pictures that contain sunflowers

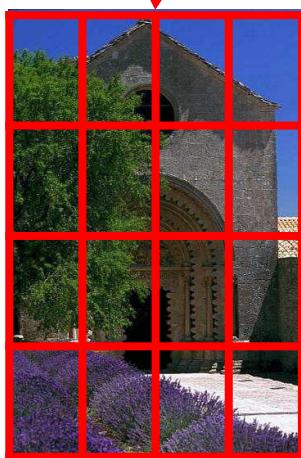


- Solution:
 - Describing **local regions**
 - Adding **spatial constraints** if need

Local feature extraction

- Local features: how to determine image patches / local regions

Dividing into
patches with
regular grid



Without knowledge about
image content

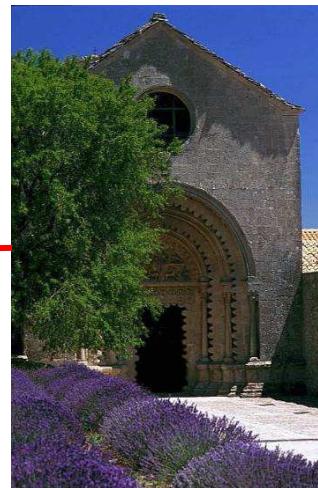


Image segmentation



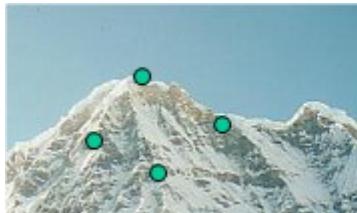
Keypoint detection



Based on the content of image

Common Requirements

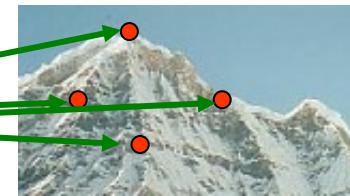
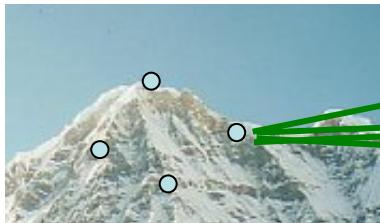
- Problem 1:
 - Detect the same point *independently* in both images



No chance to match!

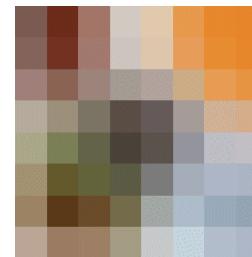
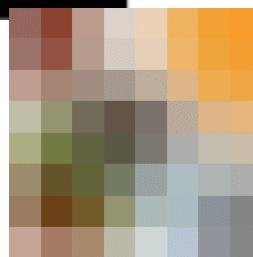
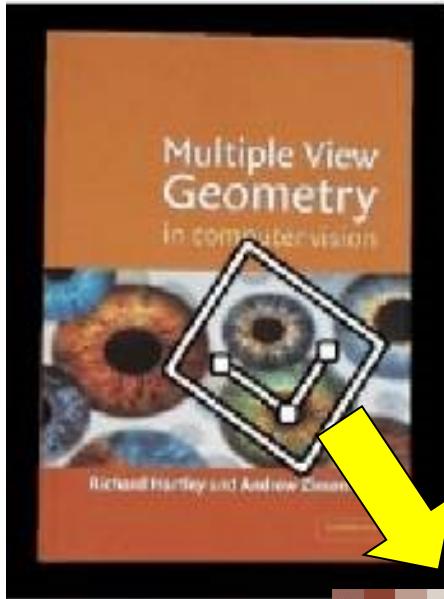
We need a repeatable **detector**!

- Problem 2:
 - For each point correctly recognize the corresponding one



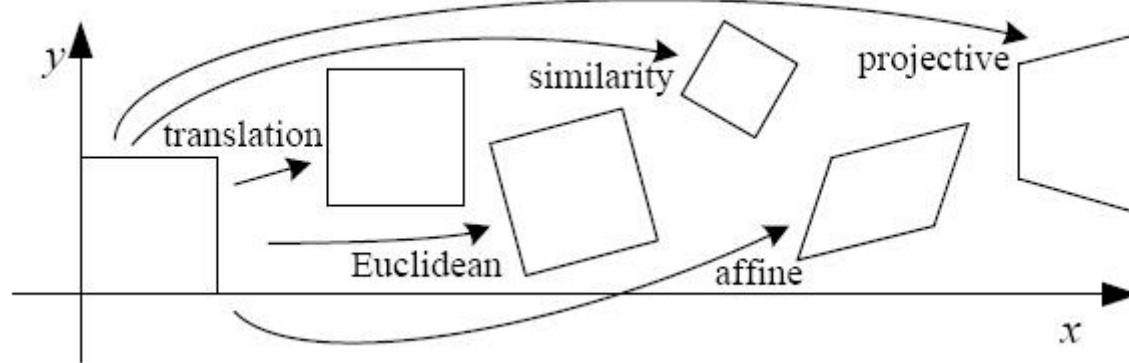
We need a reliable and distinctive **descriptor**!

Invariance: Geometric Transformations



Slide credit: Steve Seitz

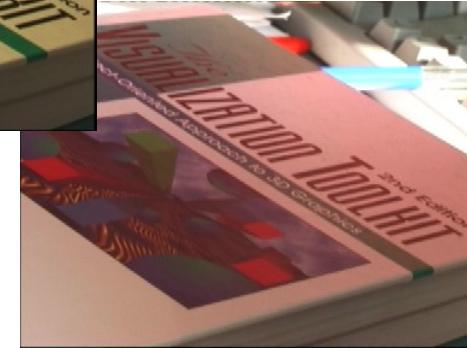
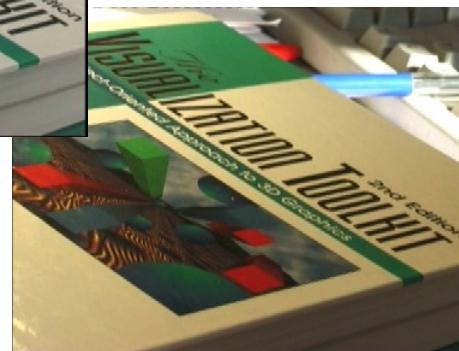
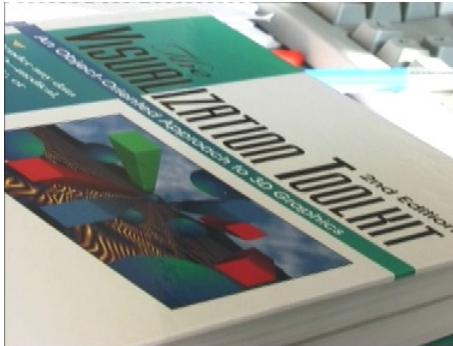
Invariance: Geometric Transformations



Levels of Geometric Invariance

Slide credit: Steve Seitz

Invariance: Photometric Transformations



- Often modeled as a linear transformation:
 - Scaling + Offset

Slide credit: Tinne Tuytelaars

Requirements

- Region extraction needs to be **repeatable** and **accurate**
 - **Invariant** to translation, rotation, scale changes
 - **Robust** or **covariant** to out-of-plane (\approx affine) transformations
 - **Robust** to lighting variations, noise, blur, quantization
- **Locality**: Features are local, therefore robust to occlusion and clutter
- **Quantity**: We need a sufficient number of regions to cover the object
- **Distinctiveness**: The regions should contain “interesting” structure
- **Efficiency**: Close to real-time performance

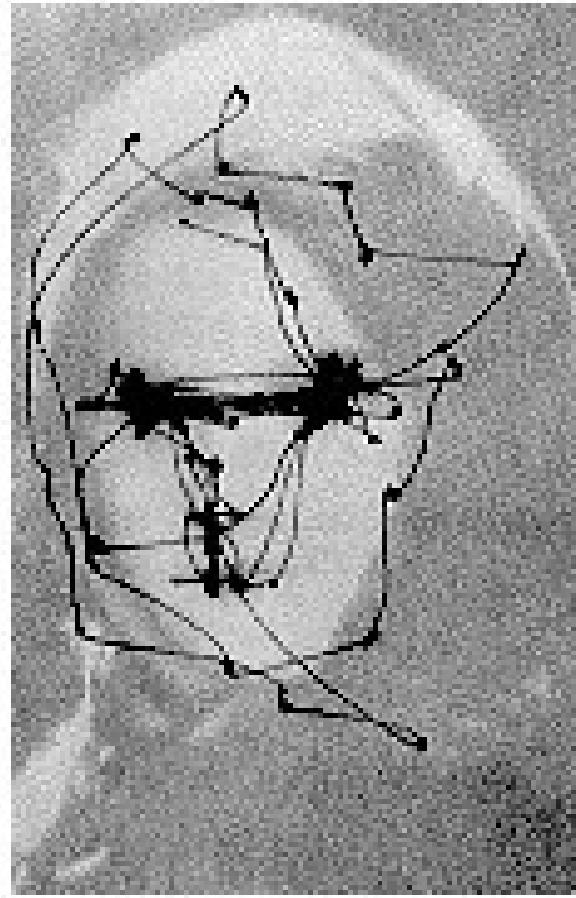
Main questions

- **Where** will the interest points come from?
 - What are salient features that we'll *detect* in multiple views?
- How **to describe** a local region?
- How to **establish correspondences**, i.e., compute matches?

Feature extraction

- Global features
- **Local features**
 - **Interest point detector**
 - Local descriptor
 - Matching

Interest points: why and where?



Yarbus eye tracking

Source : Derek Hoiem, Computer Vision, University of Illinois.

Same image with different questions



Free examination.

1

Estimate material circumstances
of the family



3

Give the ages of the people.



4

Surmise what the family had
been doing before the arrival
of the unexpected visitor.



5

Remember the clothes
worn by the people.



6

Remember positions of people and
objects in the room.



7

Estimate how long the visitor had
been away from the family.

3 min. recordings
of the same
subject

Interest points: why and where?

- Where will the interest points come from?

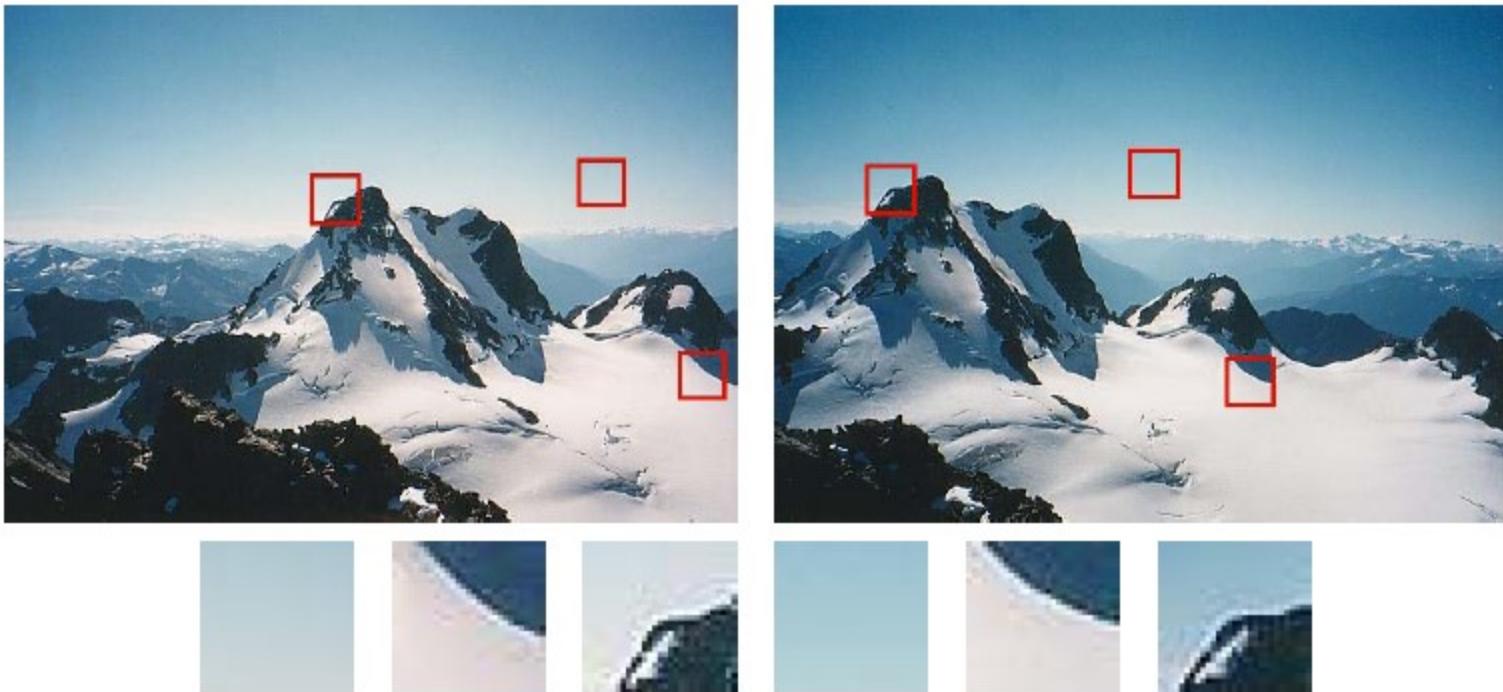


Figure 4.3: *Image pairs with extracted patches below. Notice how some patches can be localized or matched with higher accuracy than others.*

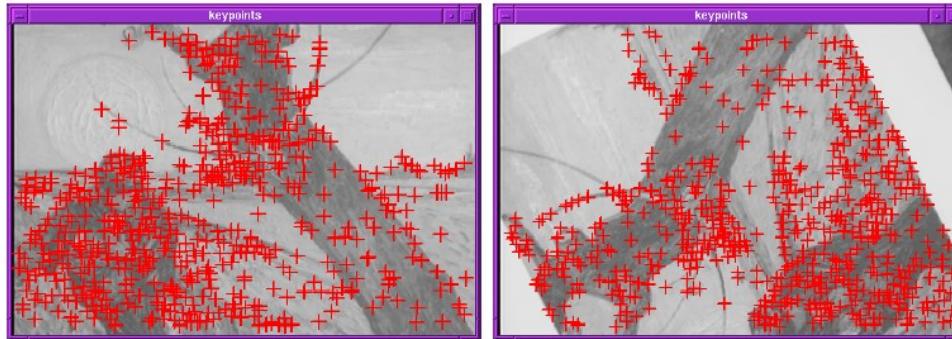
Keypoint Localization

- Goals:
 - Repeatable detection
 - Precise localization
 - Interesting content
- \Rightarrow *Look for two-dimensional signal changes*



Slide credit: Bastian Leibe

Finding Corners

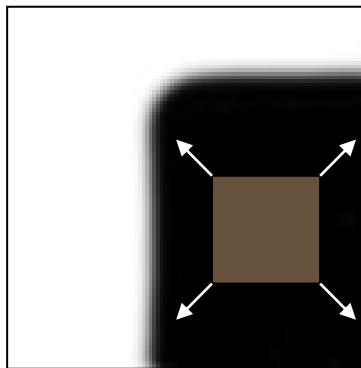


- Key property:
 - In the region around a corner, image gradient has two or more dominant directions
- Corners are *repeatable* and *distinctive*

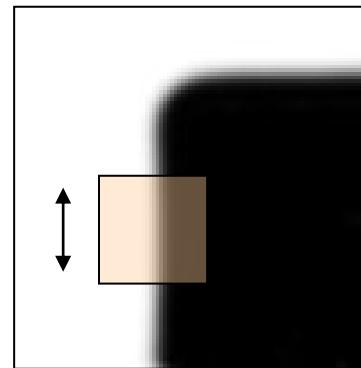
C.Harris and M.Stephens. "[A Combined Corner and Edge Detector](#)." *Proceedings of the 4th Alvey Vision Conference*, 1988.

Corners as distinctive interest points

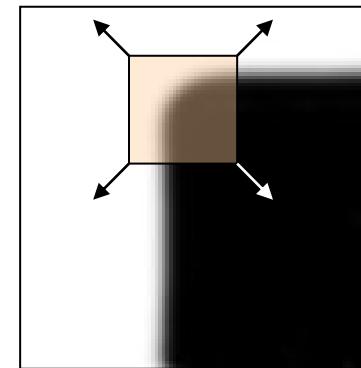
- Design criteria
 - We should easily recognize the point by looking through a small window (*locality*)
 - Shifting the window in *any direction* should give a *large change* in intensity (*good localization*)



“flat” region:
no change in
all directions



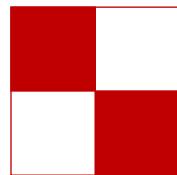
“edge”:
no change along
the edge
direction



“corner”:
significant change
in all directions

Slide credit: Alyosha Efros

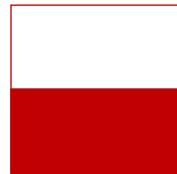
Corners versus edges



$$\sum I_x^2 \longrightarrow \text{Large}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$

Corner



$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$

Edge



$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Small}$$

Nothing

Harris detector formulation

Change of intensity for the shift $[u, v]$:

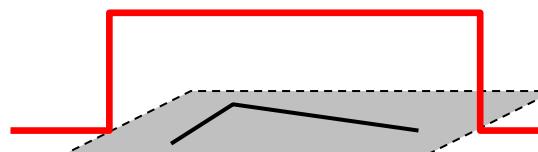
$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function

Shifted
intensity

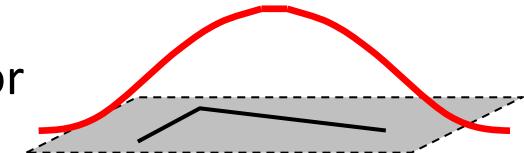
Intensity

Window function $w(x, y) =$



1 in window, 0 outside

or



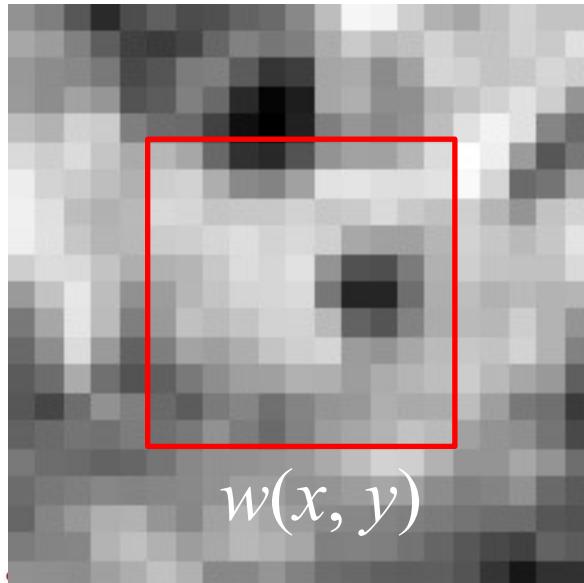
Gaussian

Corner Detection by Auto-correlation

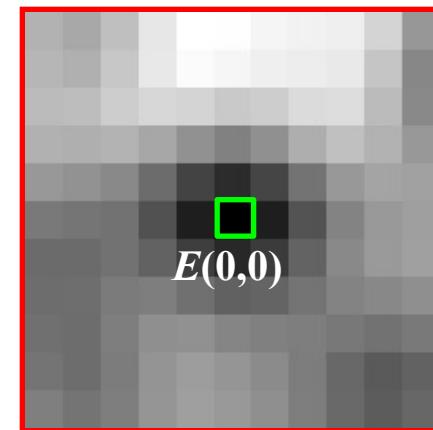
Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



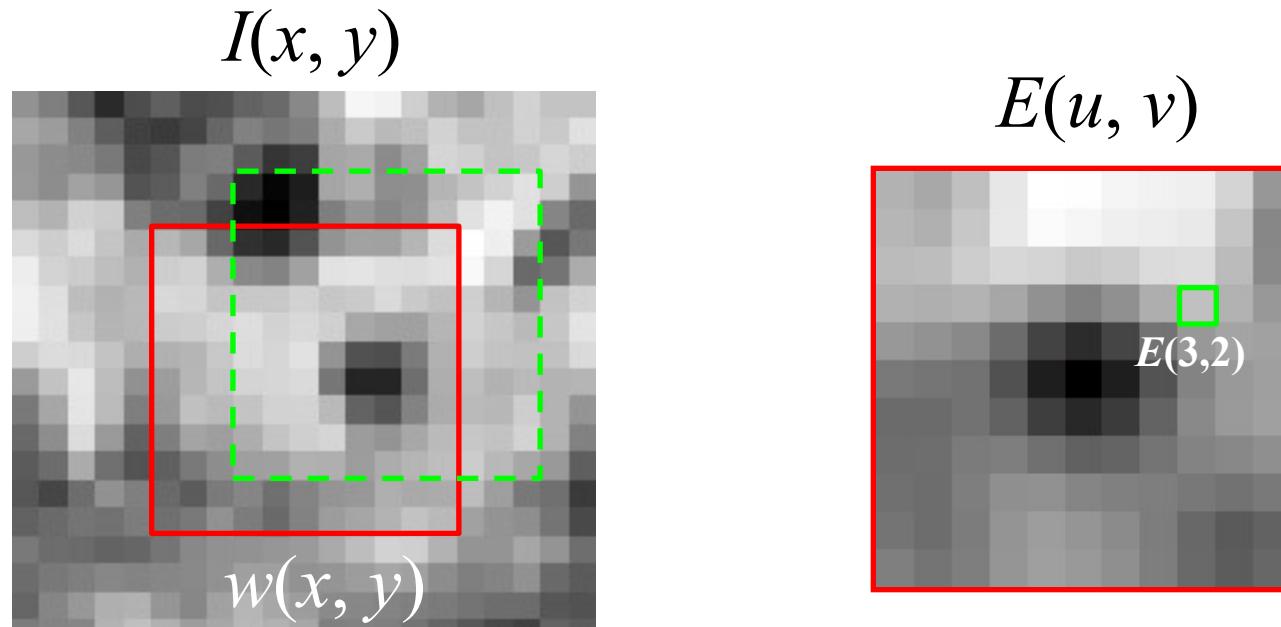
$E(u, v)$



Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to discover how E behaves for small shifts

But this is **very slow** to compute naively.

$O(\text{window_width}^2 * \text{shift_range}^2 * \text{image_width}^2)$

$O(11^2 * 11^2 * 600^2) = 5.2 \text{ billion of these}$

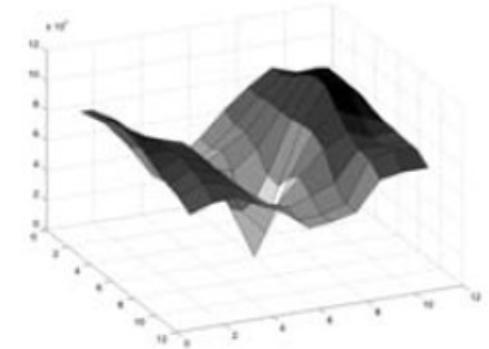
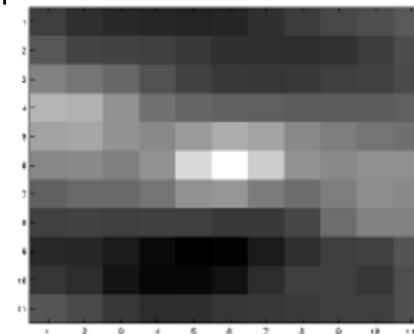
Corner Detection by Auto-correlation

Change in appearance of window $w(x,y)$ for shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to discover how E behaves for small shifts

But we know the response in E that we are looking for – **strong peak**. → Approximation



Local quadratic approximation of $E(u,v)$ in the neighborhood of $(0,0)$ is given by the second-order Taylor expansion:

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

Notation: partial derivative



we are looking for **strong peak**

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



Ignore function
value; set to 0



Ignore first
derivative,
set to 0



Just look at
shape of
second
derivative

Corner Detection: Mathematics

Second-order Taylor expansion of $E(u,v)$ about (0,0):

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_u(u,v) = \sum_{x,y} 2w(x,y) [I(x+u, y+v) - I(x, y)] I_x(x+u, y+v)$$

$$\begin{aligned} E_{uu}(u,v) = & \sum_{x,y} 2w(x,y) I_x(x+u, y+v) I_x(x+u, y+v) \\ & + \sum_{x,y} 2w(x,y) [I(x+u, y+v) - I(x, y)] I_{xx}(x+u, y+v) \end{aligned}$$

$$\begin{aligned} E_{uv}(u,v) = & \sum_{x,y} 2w(x,y) I_y(x+u, y+v) I_x(x+u, y+v) \\ & + \sum_{x,y} 2w(x,y) [I(x+u, y+v) - I(x, y)] I_{xy}(x+u, y+v) \end{aligned}$$

Corner Detection: Mathematics

Second-order Taylor expansion of $E(u,v)$ about (0,0):

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0$$

$$E_u(0,0) = 0$$

$$E_v(0,0) = 0$$

$$E_{uu}(0,0) = \sum_{x,y} 2w(x,y) I_x(x,y) I_x(x,y)$$

$$E_{vv}(0,0) = \sum_{x,y} 2w(x,y) I_y(x,y) I_y(x,y)$$

$$E_{uv}(0,0) = \sum_{x,y} 2w(x,y) I_x(x,y) I_y(x,y)$$

Corner Detection: Mathematics

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$$E(u, v) \approx [u \ v] \begin{bmatrix} \sum_{x,y} w(x, y) I_x^2(x, y) & \sum_{x,y} w(x, y) I_x(x, y) I_y(x, y) \\ \sum_{x,y} w(x, y) I_x(x, y) I_y(x, y) & \sum_{x,y} w(x, y) I_y^2(x, y) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

Harris detector formulation

- This measure of change can be approximated by:

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Gradient with respect to x,
times gradient with respect to y

↑
Sum over image region – the area
we are checking for corner

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$

Slide credit: Rick Szeliski

What does this matrix reveal?

- First, let's consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

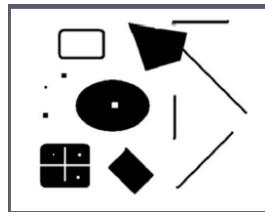
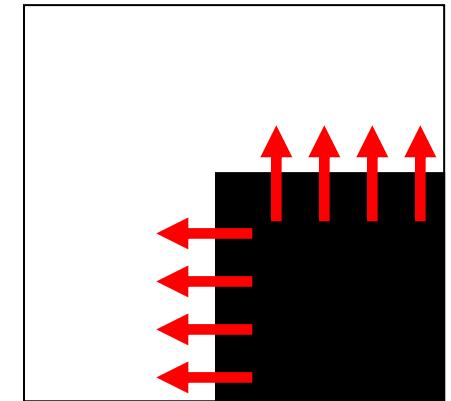
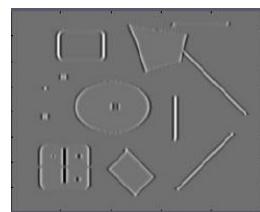
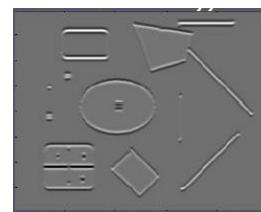


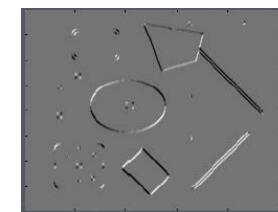
Image I



I_x



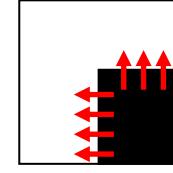
I_y



$I_x I_y$

What does this matrix reveal?

- First, let's consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$


- This means:
 - Dominant gradient directions align with x or y axis
 - If either λ is close to 0, then this is not a corner, so look for locations where both are large.
- What if we have a corner that is not aligned with the image axes?

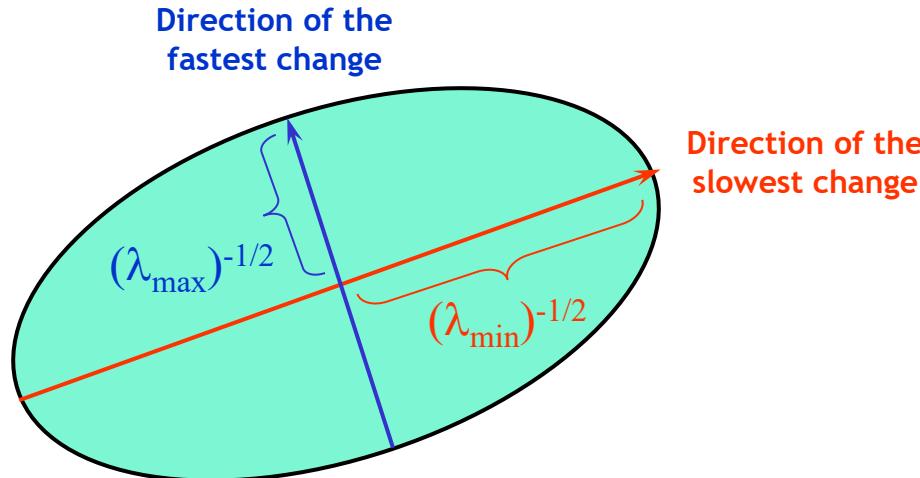
General case

- Since M is symmetric, we have

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

(Eigenvalue decomposition)

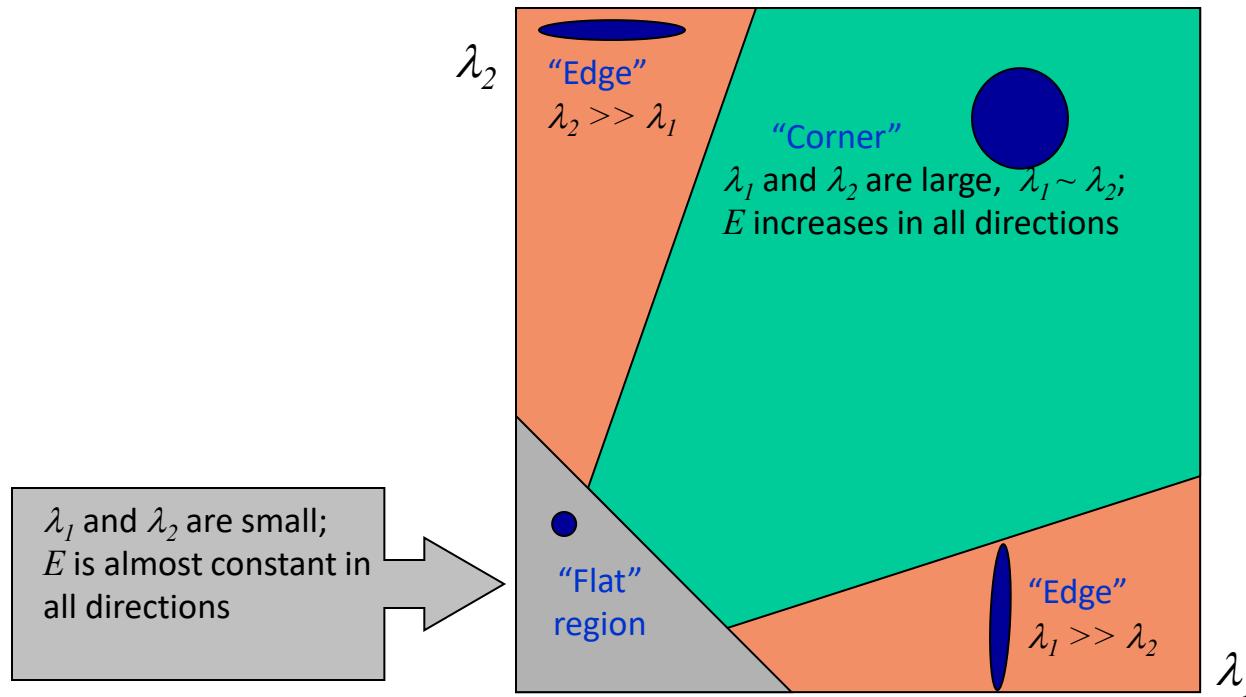
- We can visualize M as an ellipse with axis lengths determined by the eigenvalues and orientation determined by a rotation matrix R



adapted from Darya Frolova, Denis Simakov

Interpreting the eigenvalues

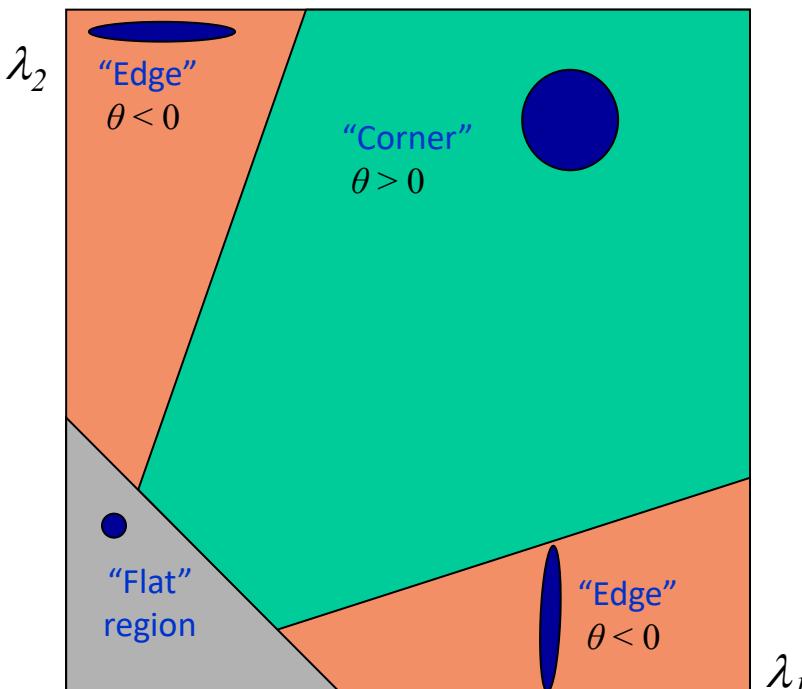
- Classification of image points using eigenvalues of M :



Slide credit: Kristen Grauman

Corner response function

$$\theta = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$



- Fast approximation
 - Avoid computing the eigenvalues
 - α : constant (0.04 to 0.06)

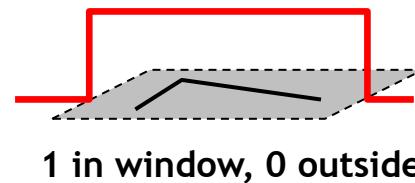
Window Function

$$w(x,y)$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Option 1: uniform window

- Sum over square window

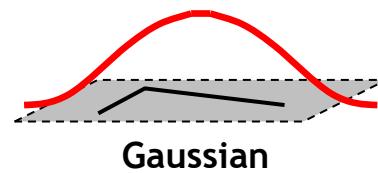


$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Problem: not rotation invariant

- Option 2: Smooth with Gaussian

- Gaussian already performs weighted sum



$$M = g(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Result is rotation invariant

Summary: Harris Detector [Harris88]

- Compute second moment matrix (autocorrelation matrix)

$$M(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

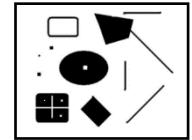
2. Square of derivatives

3. Gaussian filter $g(\sigma_l)$

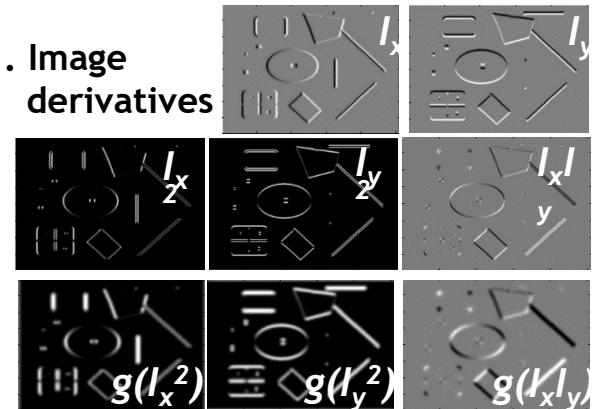
4. Cornerness function - two strong eigenvalues

$$\begin{aligned}\theta &= \det[M(\sigma_I, \sigma_D)] - \alpha[\text{trace}(M(\sigma_I, \sigma_D))]^2 \\ &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2\end{aligned}$$

5. Perform non-maximum suppression



1. Image derivatives



C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)" *Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

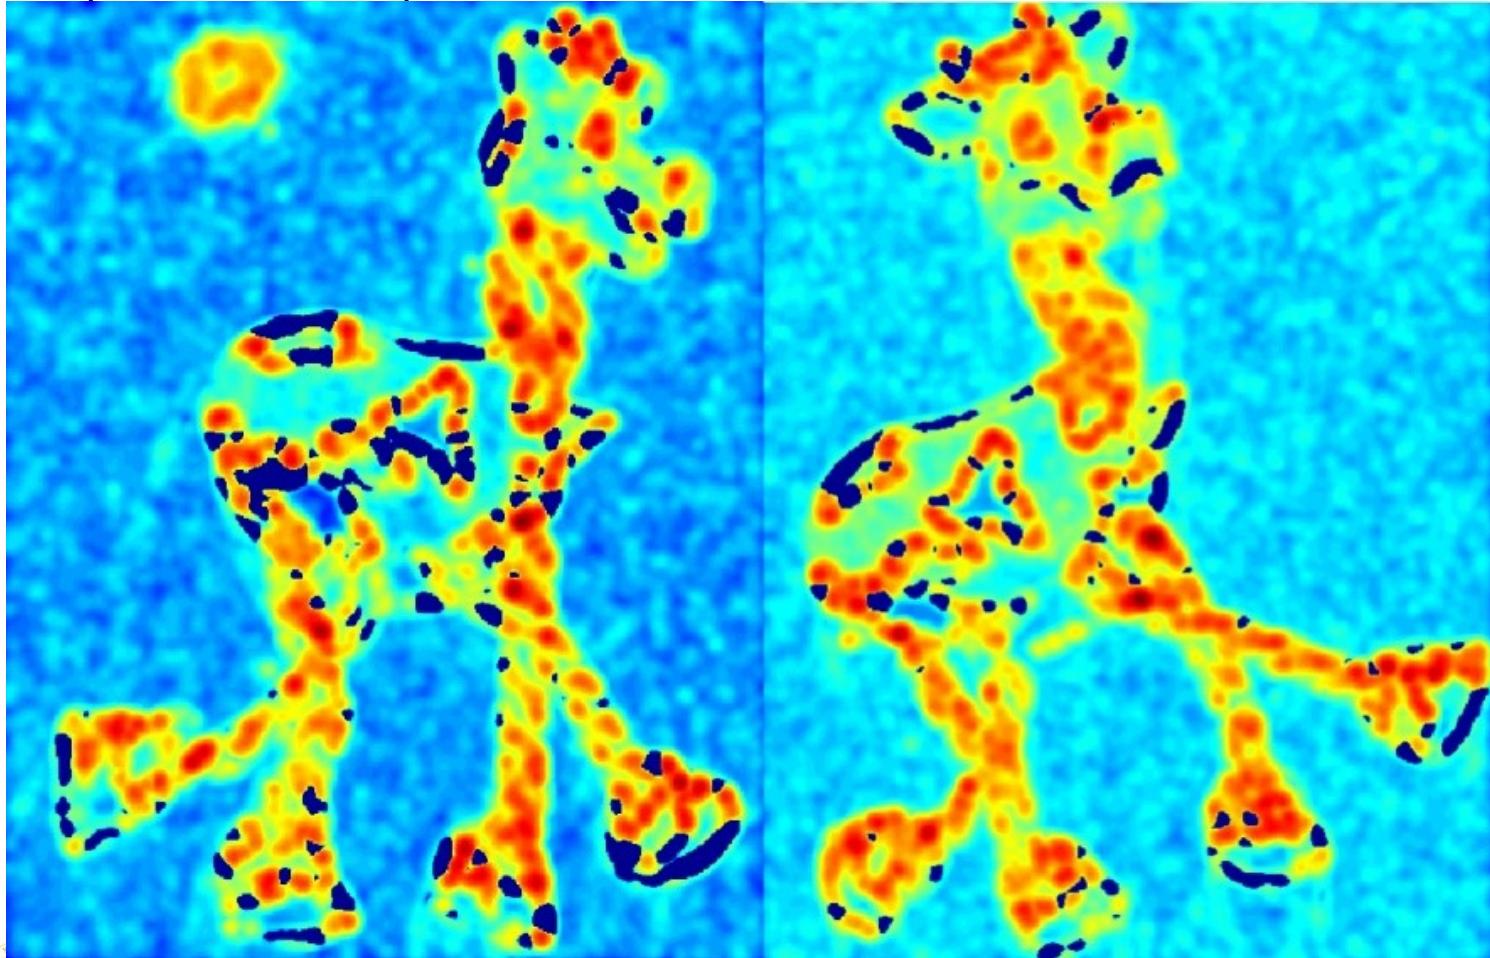
Slide credit: Krystian Mikolajczyk

Harris Detector: Workflow



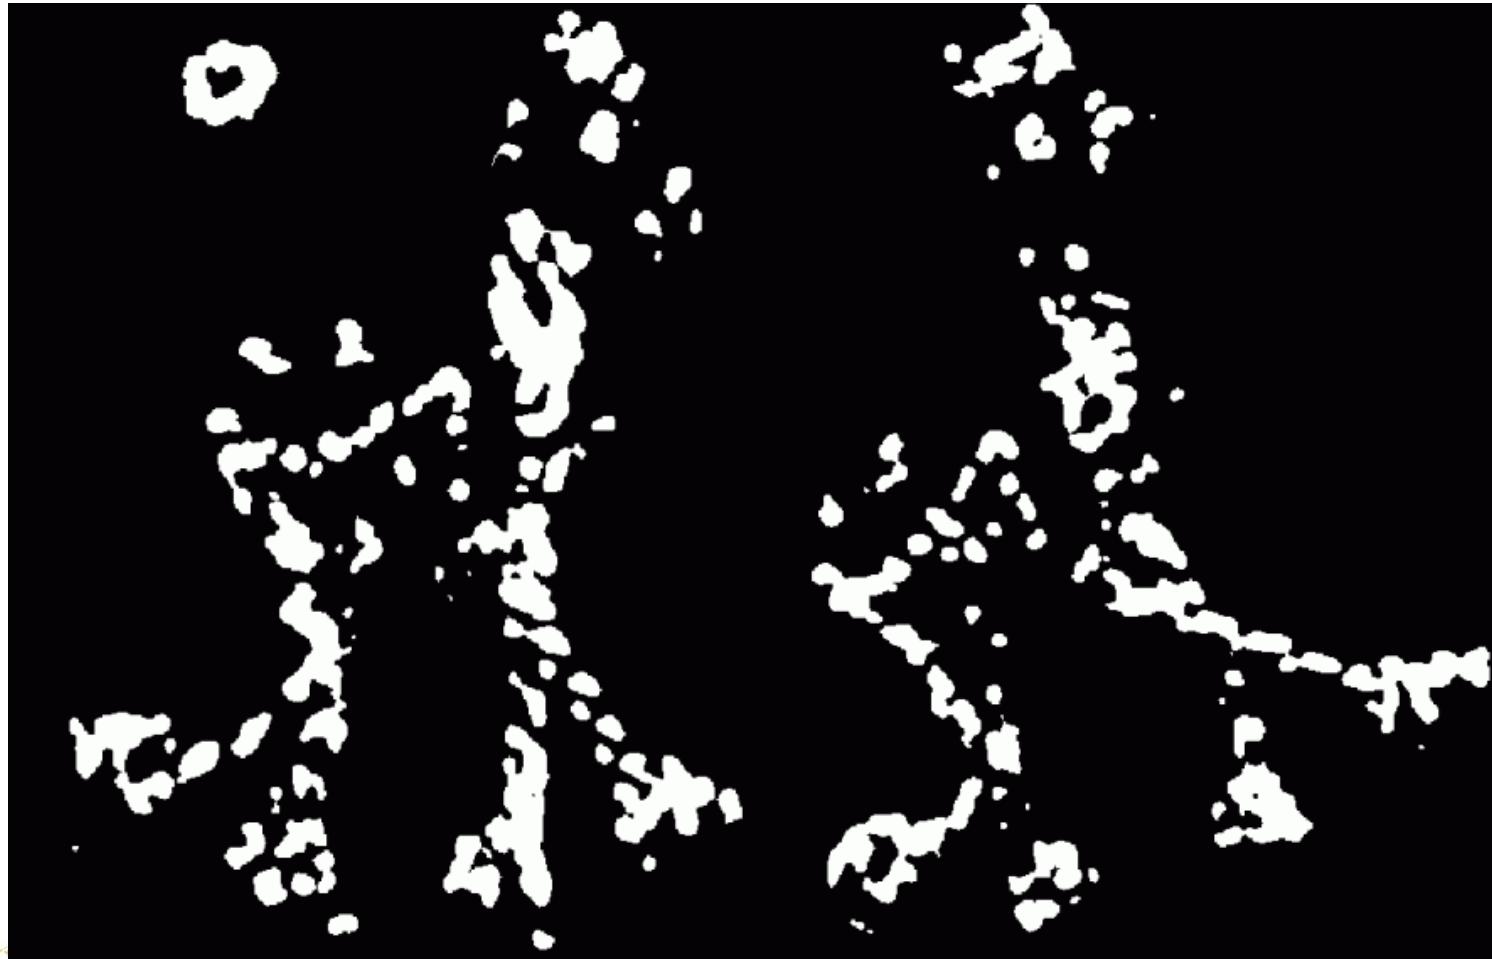
Harris Detector: Workflow

computer corner responses θ



Harris Detector: Workflow

Take points where $\theta > \text{threshold}$



Harris Detector: Workflow

Take only the local maxima of θ , where $\theta > \text{threshold}$

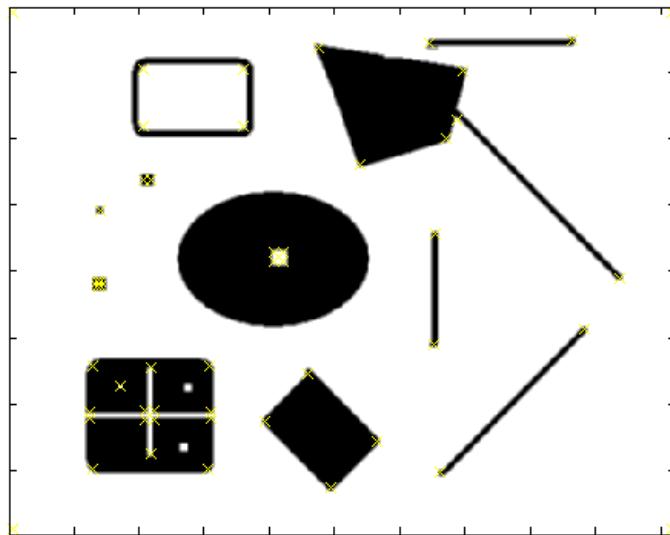


Harris Detector: Workflow

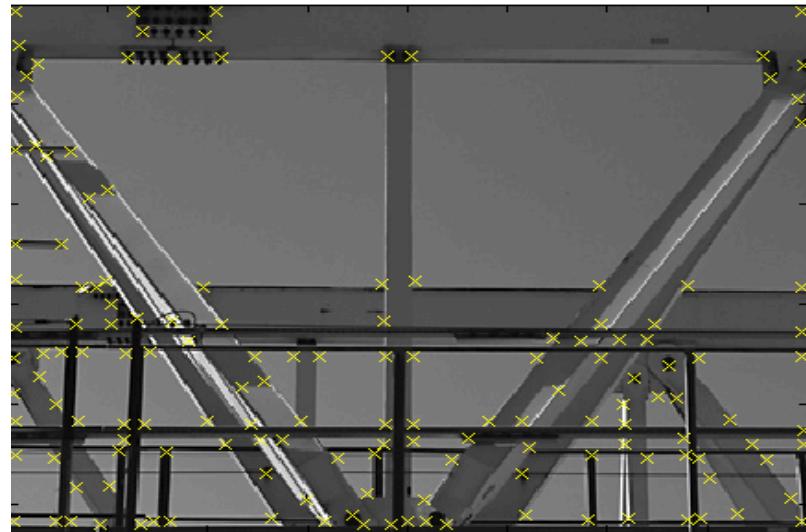
Resulting Harris points



Harris Detector – Responses [Harris88]



Effect: A very precise corner detector.

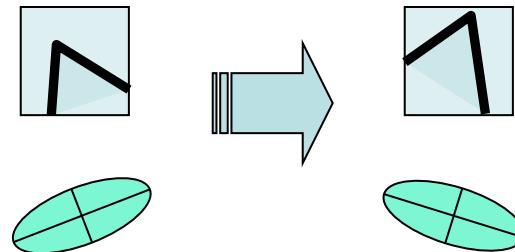


Harris Detector: Properties

- Translation invariance?

Harris Detector: Properties

- Translation invariance
- Rotation invariance?



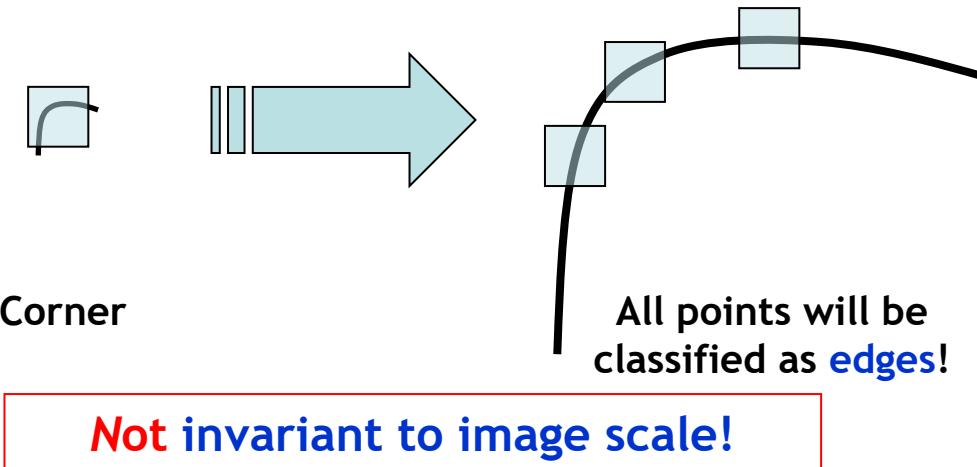
Ellipse rotates but its shape (i.e.
eigenvalues) remains the same

Corner response θ is invariant to image rotation

Slide credit: Kristen Grauman

Harris Detector: Properties

- Translation invariance
- Rotation invariance
- Scale invariance?



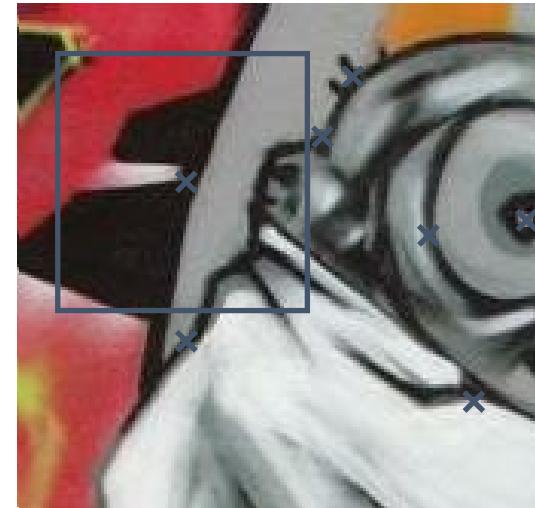
Slide credit: Kristen Grauman

Scale invariance: how to?

- Exhaustive search
- Invariance
- Robustness

Exhaustive search

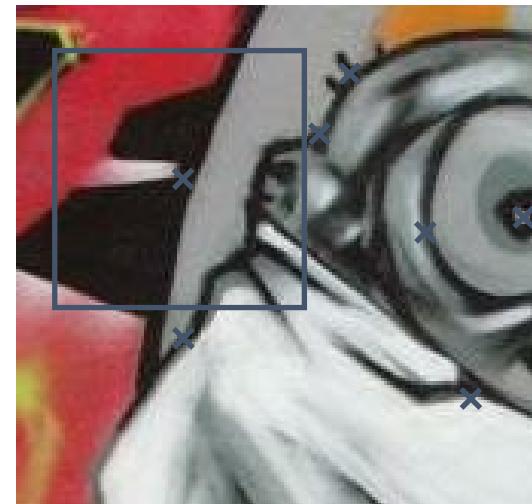
- Multi-scale approach



Slide adapted from T. Tuytelaars ECCV 2006 tutorial

Invariance

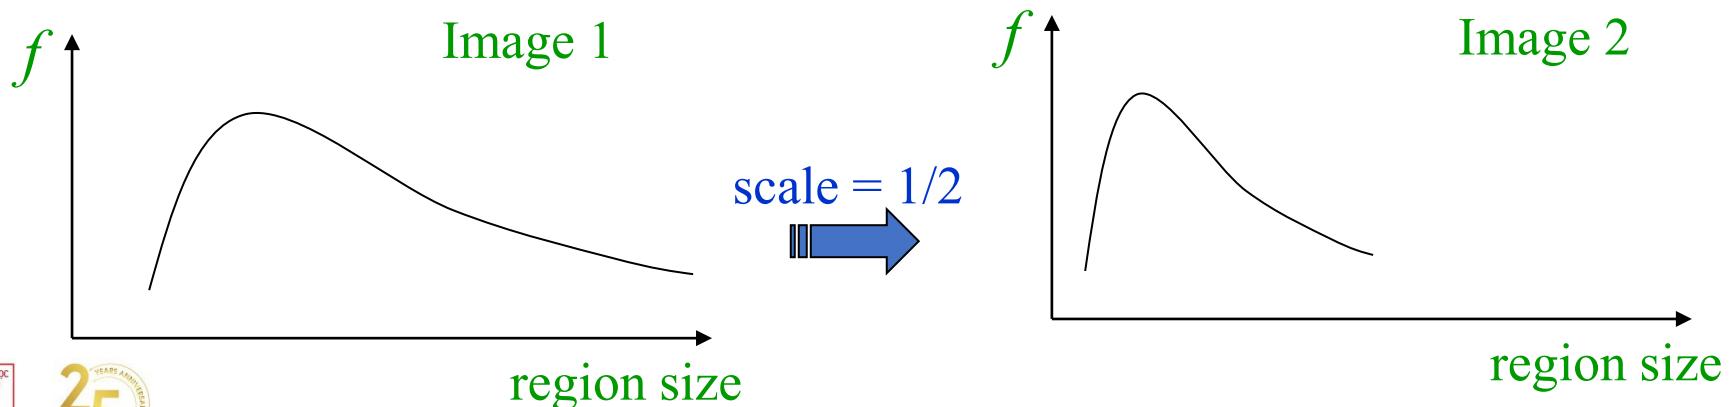
- Extract patch from each image individually



Slide adapted from T. Tuytelaars ECCV 2006 tutorial

Automatic scale selection

- Solution:
 - Design a function on the region, which is “scale invariant” (*the same for corresponding regions, even if they are at different scales*)
Example: average intensity. For corresponding regions (even of different sizes) it will be the same.
 - For a point in one image, we can consider it as a function of region size (patch width)



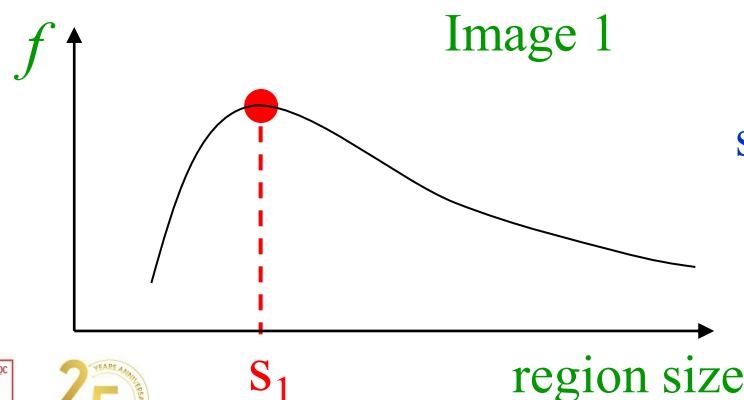
Automatic scale selection

- Common approach:

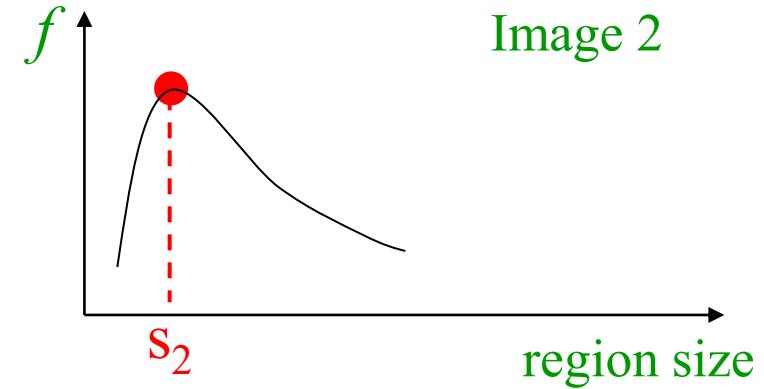
Take a local maximum of this function

Observation: region size, for which the maximum is achieved, should be *invariant* to image scale.

Important: this scale invariant region size is found in each image independently!



scale = 1/2
→



Automatic Scale Selection



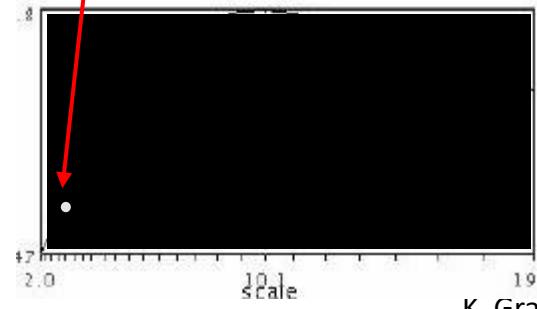
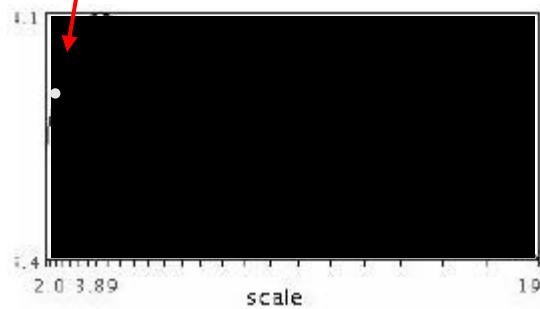
$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

Same operator responses if the patch contains the same image up to scale factor.

K. Grauman, B. Leibe

Example

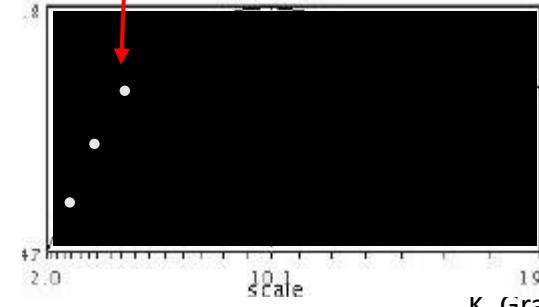
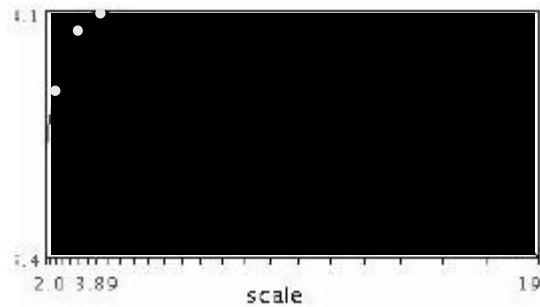
Function responses for increasing scale (scale signature)



K. Grauman, B. Leibe

Example

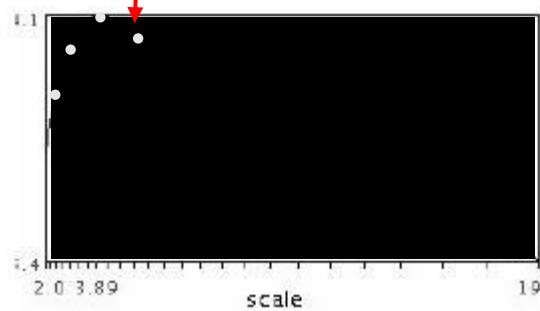
Function responses for increasing scale (scale signature)



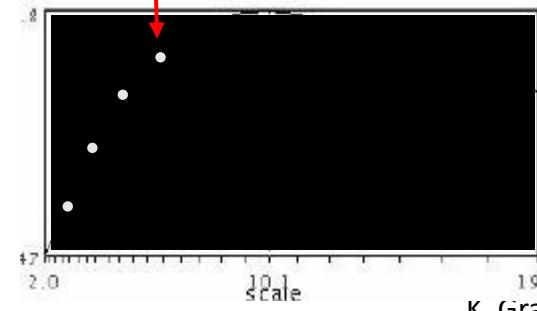
K. Grauman, B. Leibe

Example

Function responses for increasing scale (scale signature)



$$f(I_{i_1 \dots i_m}(x, \sigma))$$

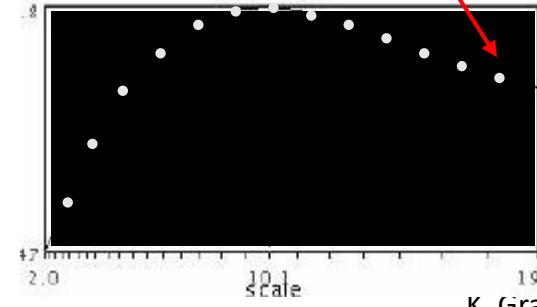
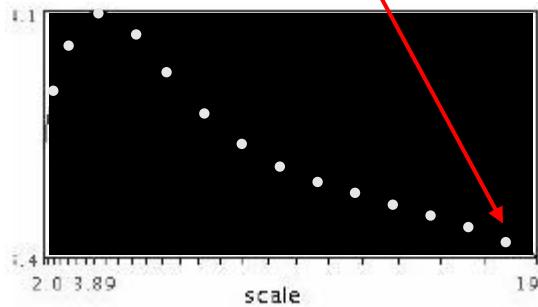


$$f(I_{i_1 \dots i_m}(x', \sigma))$$

K. Grauman, B. Leibe

Example

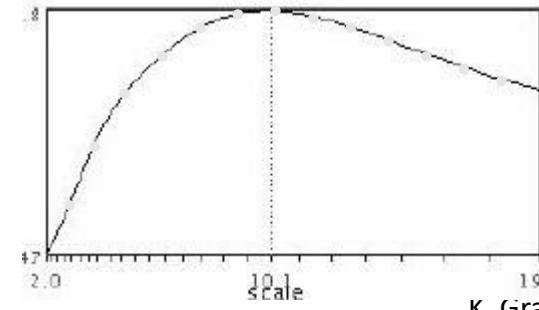
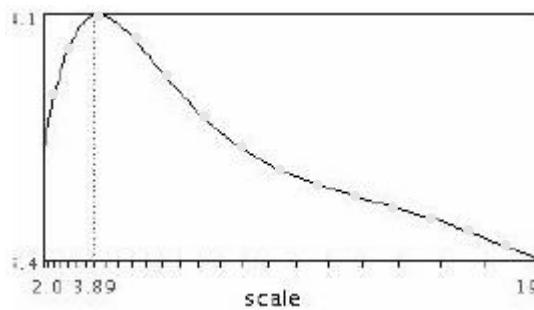
Function responses for increasing scale (scale signature)



K. Grauman, B. Leibe

Example

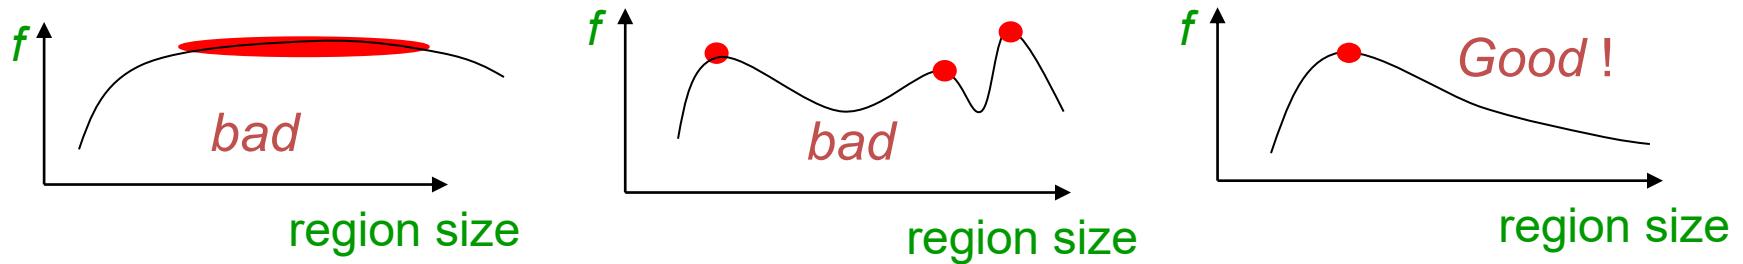
Function responses for increasing scale (scale signature)



K. Grauman, B. Leibe

Scale Invariant Detection

- A “good” function for scale detection:
has one stable sharp peak



- For usual images: a good function would be one which responds to contrast (sharp local intensity change)

What is a useful signature function?

- Functions for determining scale $f = \text{Kernel} * \text{Image}$

Kernels:

$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

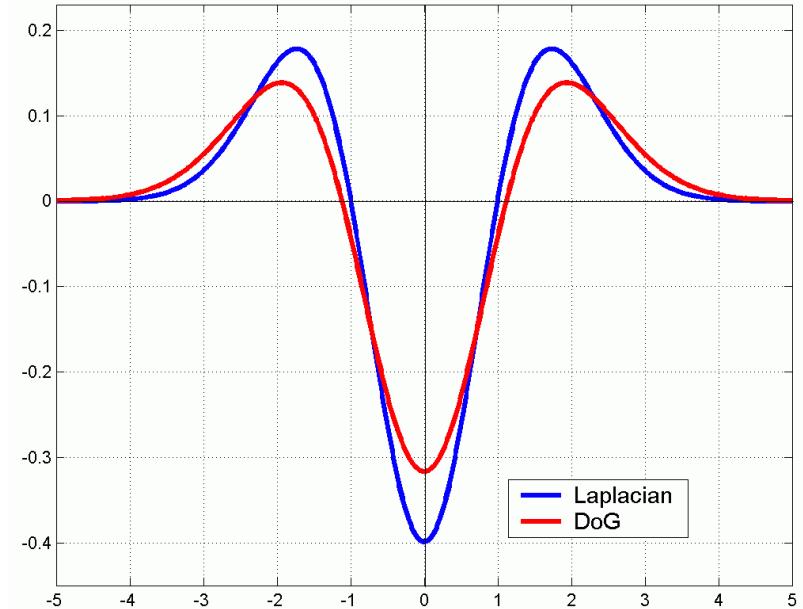
(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)

where Gaussian

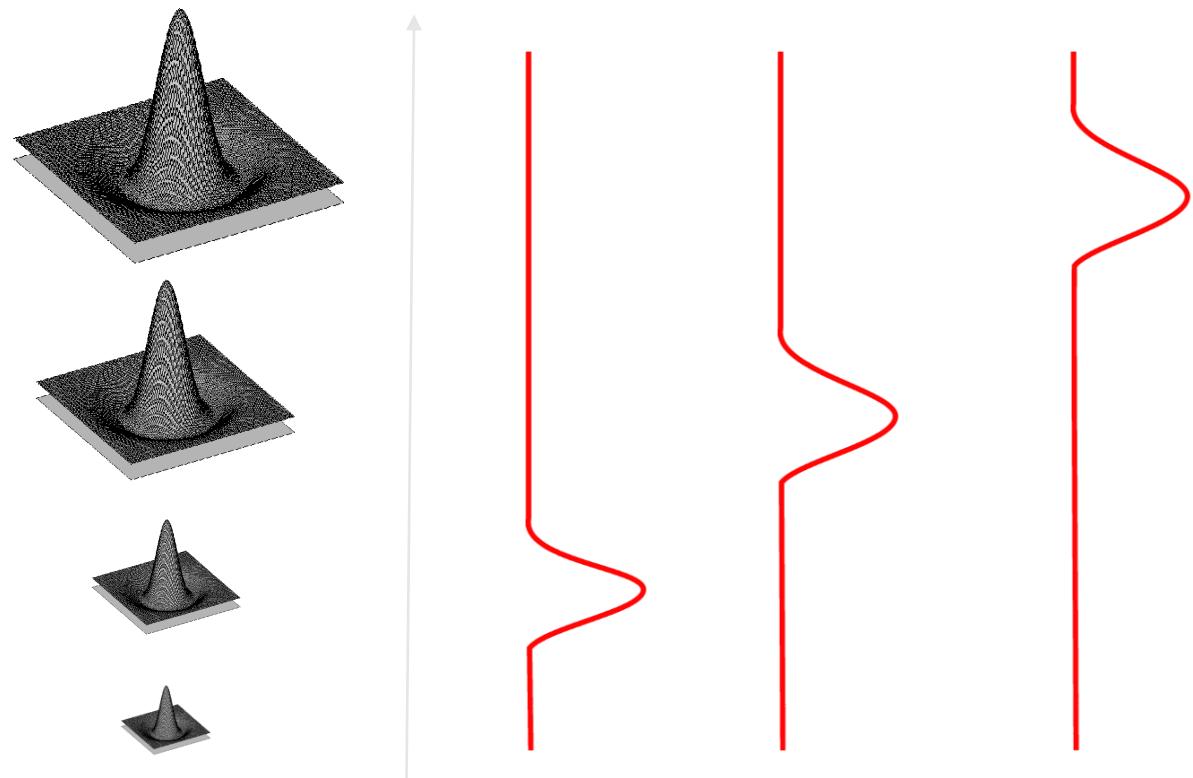
$$G(x, y, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Note: both kernels are invariant to scale and rotation

What is a useful signature function?

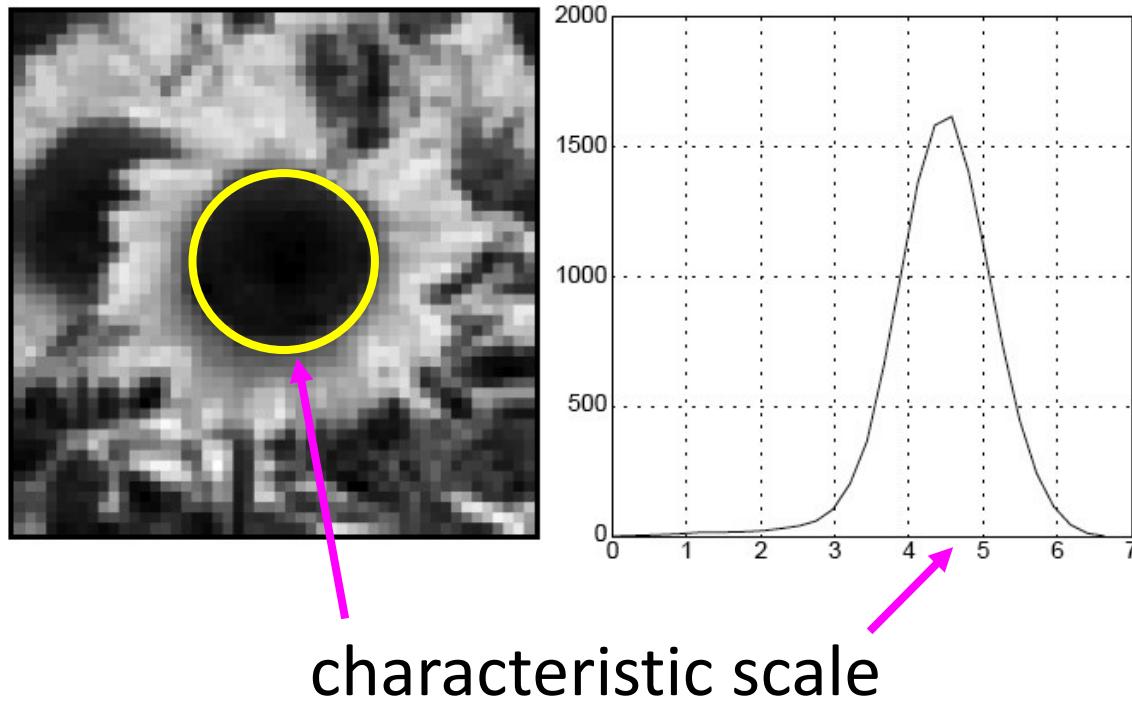
- Laplacian-of-Gaussian = “blob” detector



Source: K. Grauman, B. Leibe

Characteristic scale

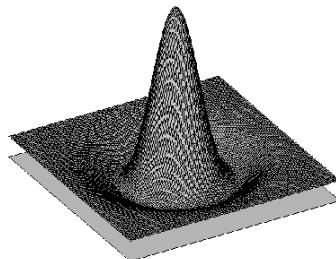
- We define the ***characteristic scale*** as the scale that produces **peak of Laplacian response**



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#) *IJCV* **30** (2): pp 77--116.

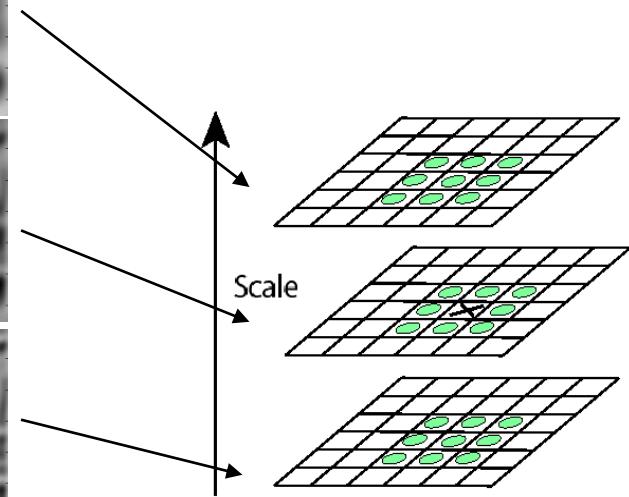
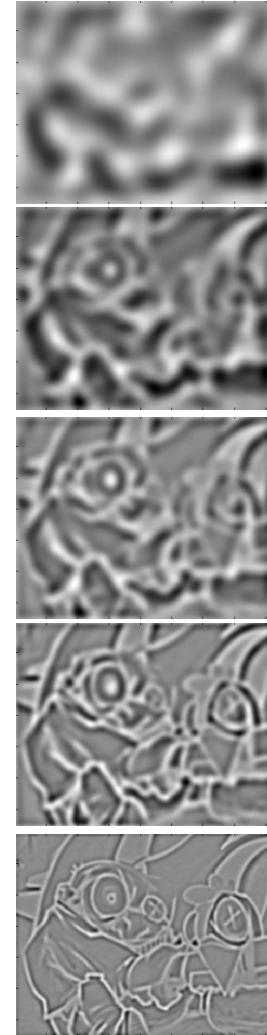
Laplacian-of-Gaussian (LoG)

- Interest points:
Local maxima in scale space of LoG



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^3$$

$$\sigma^5 \\ \sigma^4 \\ \sigma^3 \\ \sigma^2 \\ \sigma$$



⇒ List of
 (x, y, σ)

Example: Scale-space blob detector

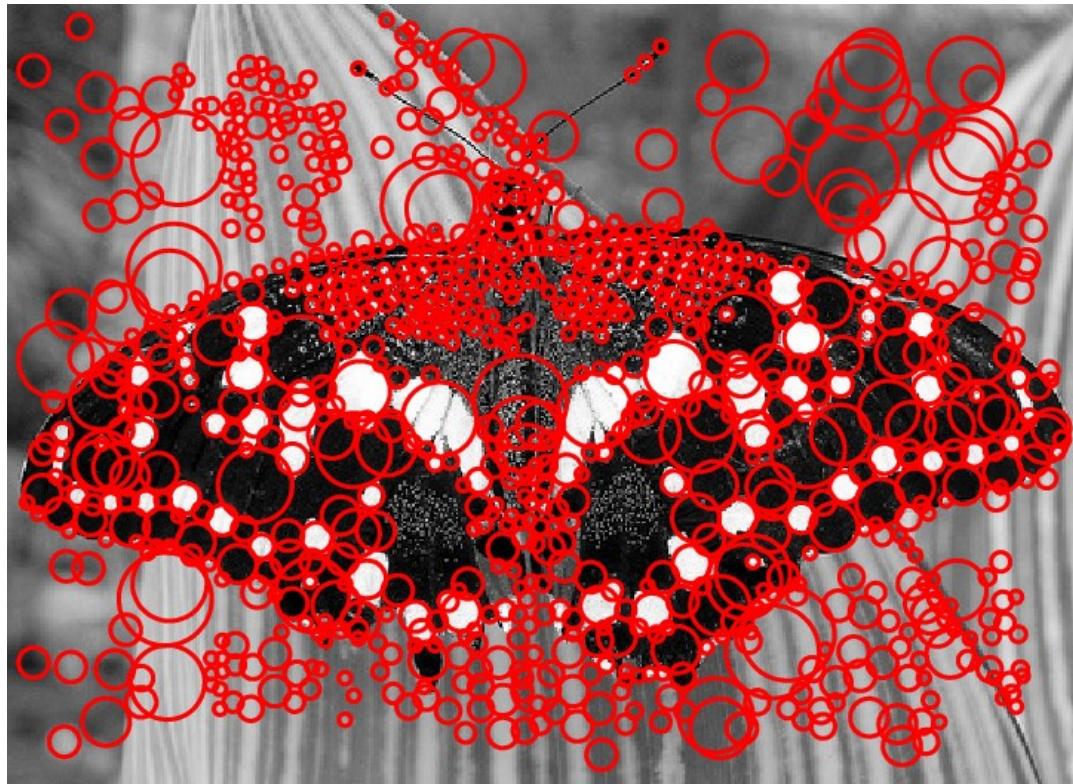


Example: Scale-space blob detector



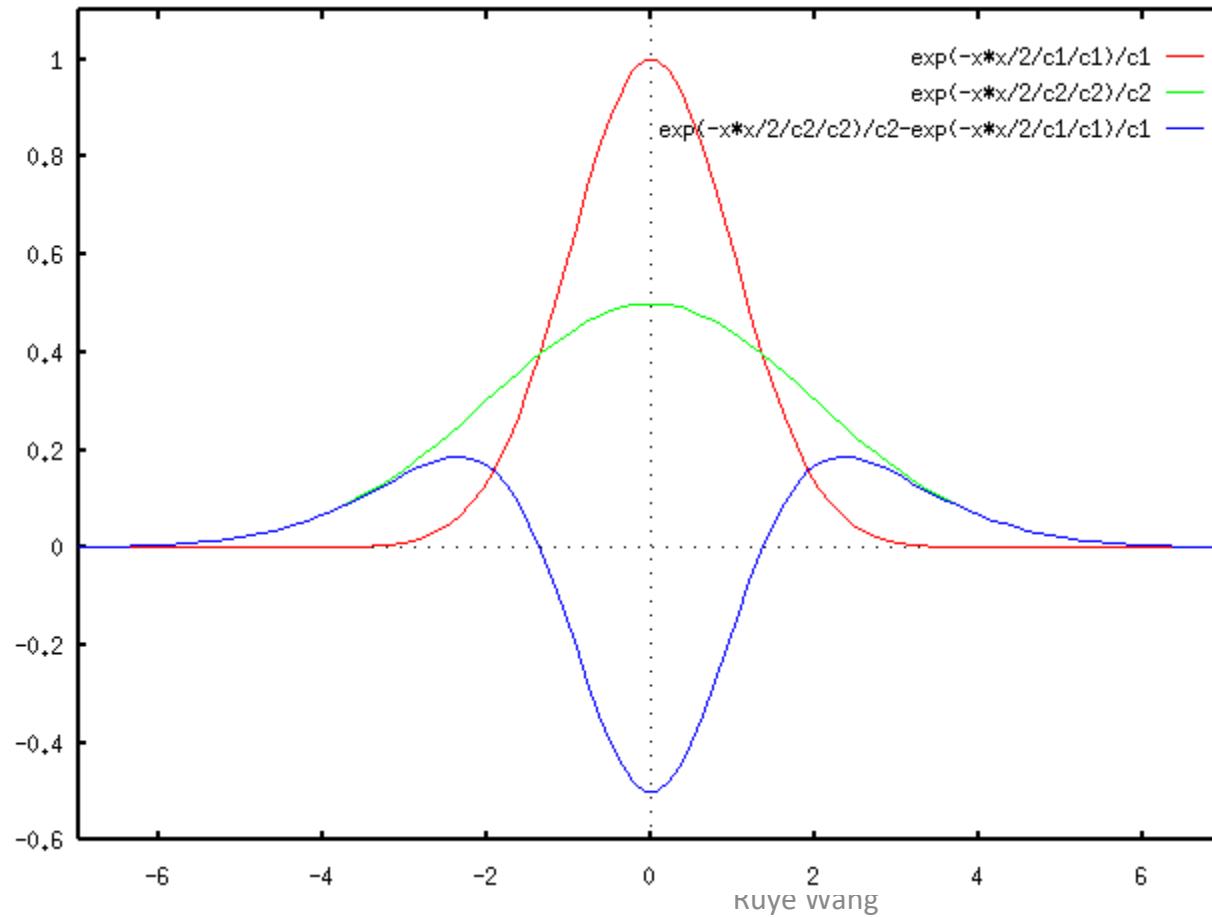
$\sigma = 11.9912$

Example: Scale-space blob detector



Alternative approach

Approximate LoG with Difference-of-Gaussian (DoG).



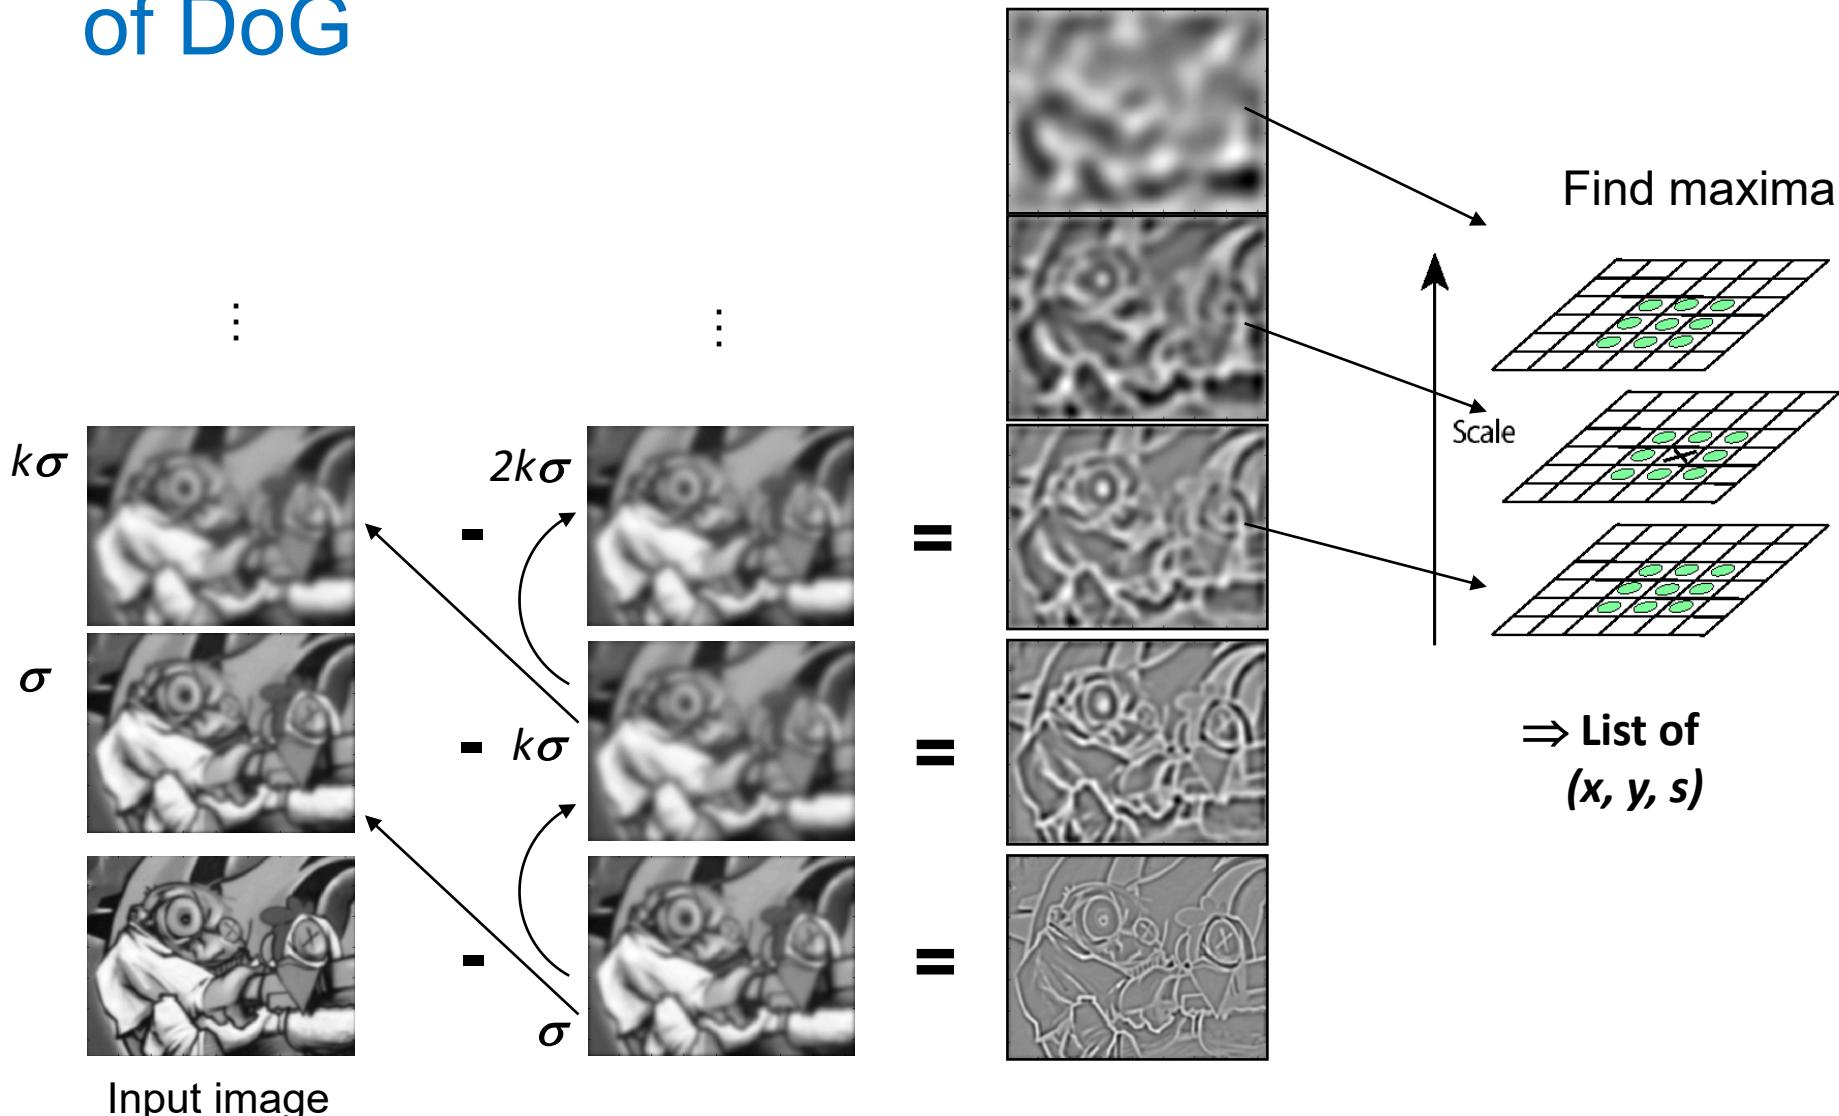
Alternative approach

- Approximate LoG with Difference-of-Gaussian (DoG):
 - 1. Blur image with σ Gaussian kernel
 - 2. Blur image with $k\sigma$ Gaussian kernel
 - 3. Subtract 2. from 1.

Small k gives a closer approximation to LoG, but usually we want to build a scale space quickly out of this. $k = 1.6$ gives an appropriate scale space, $k = \text{sqrt}(2)$



Find local maxima in position-scale space of DoG

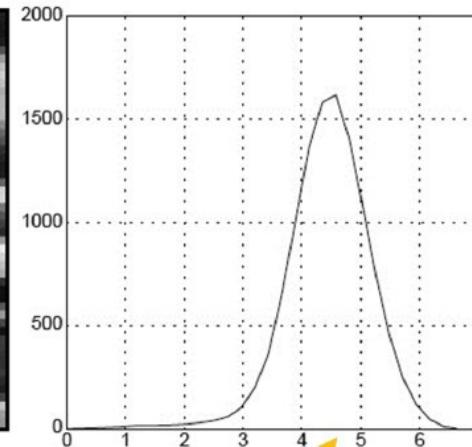
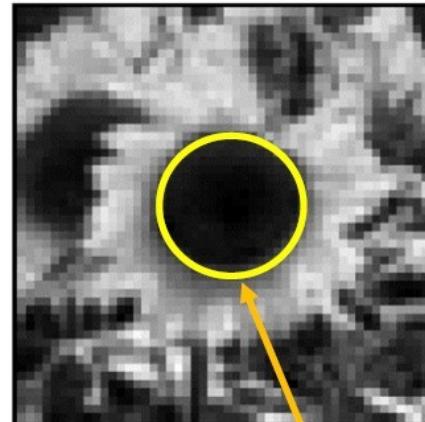
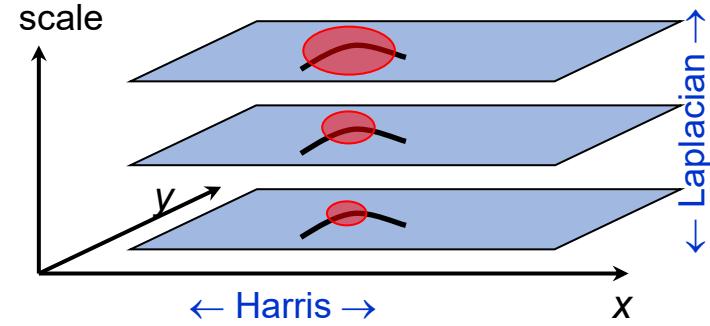


Harris-Laplacian

- **Harris-Laplacian¹**

Find local maximum of:

- Harris corner detector in space
(image coordinates)
- Laplacian in scale



Characteristic scale

¹ K.Mikolajczyk, C.Schmid. "Indexing Based on Scale Invariant Interest Points". ICCV 2001

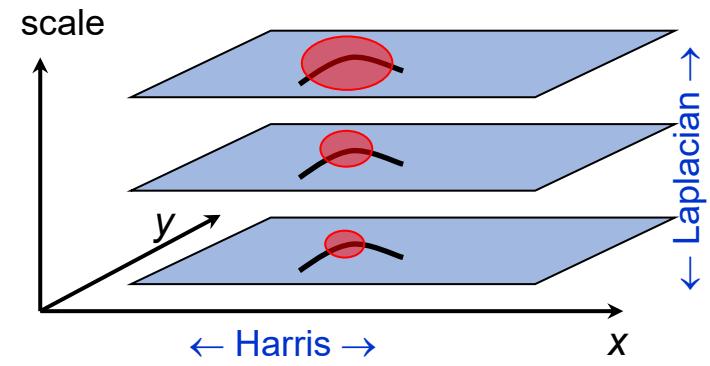
² D.Lowe. "Distinctive Image Features from Scale-Invariant Keypoints". IJCV 2004

Scale Invariant Detectors

- **Harris-Laplacian¹**

Find local maximum of:

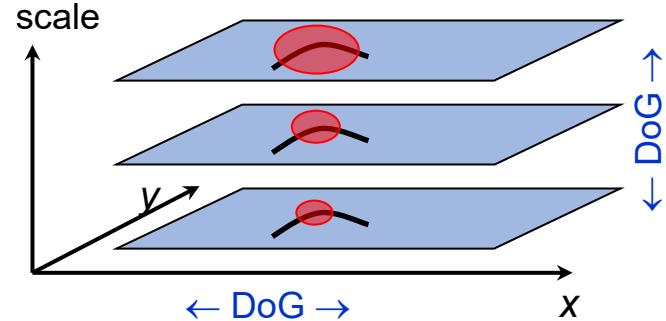
- Harris corner detector in space (image coordinates)
- Laplacian in scale



- **SIFT (D.Lowe)²**

Find local maximum of:

- Difference of Gaussians in space and scale



¹ K.Mikolajczyk, C.Schmid. "Indexing Based on Scale Invariant Interest Points". ICCV 2001

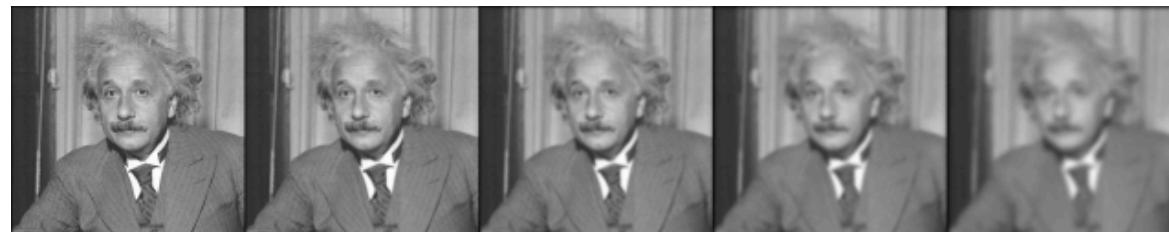
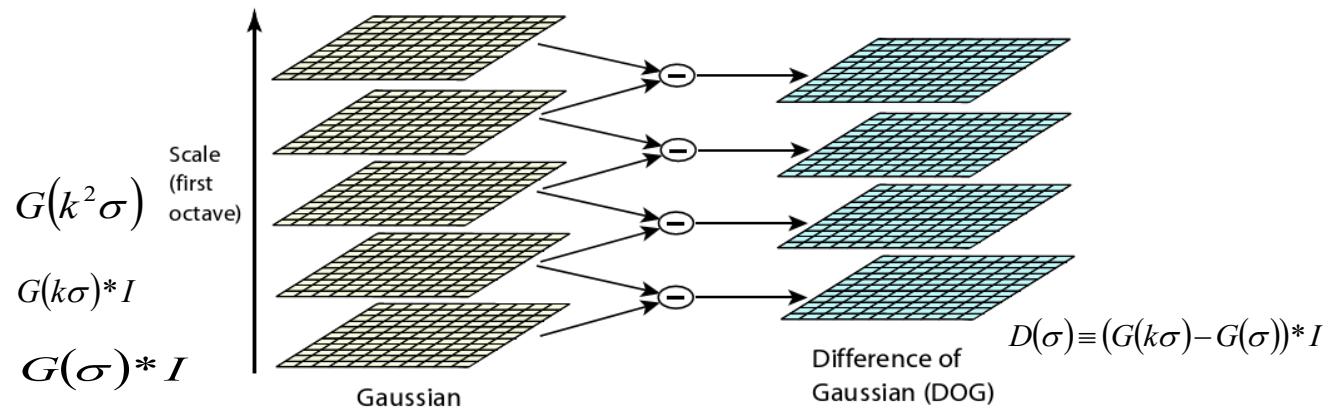
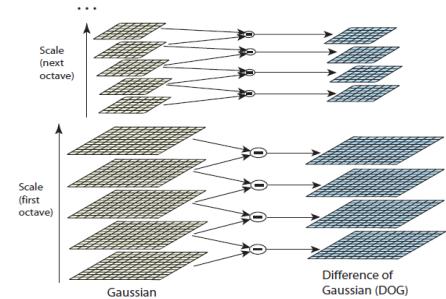
² D.Lowe. "Distinctive Image Features from Scale-Invariant Keypoints". IJCV 2004

DoG (SIFT) keypoint Detector

- DoG at multi-octaves
- Extrema detection in scale space
- Keypoint location
 - Interpolation
 - Removing instable points
- Orientation Assignment

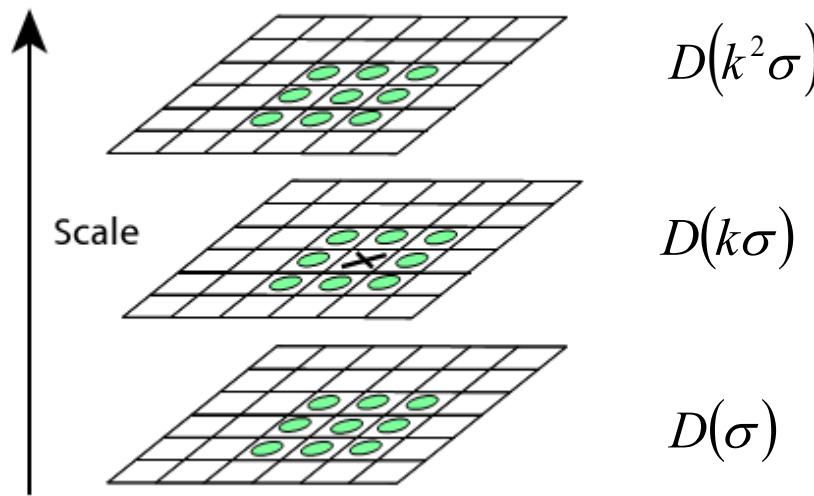
DoG (SIFT) Detector

- DoG at multi-octaves



DoG (SIFT) Detector

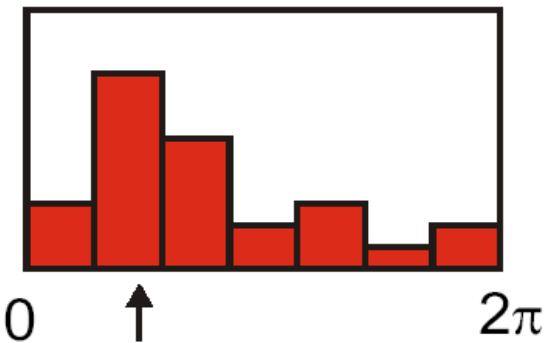
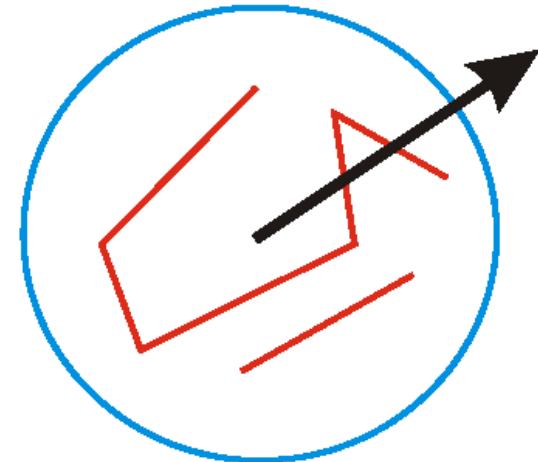
- Scale-Space Extrema Choose all extrema within 3x3x3 neighborhood



X is selected if it is larger or smaller than all 26 neighbors
(its 8 neighbors in the current image and 9 neighbors
each in the scales above and below)

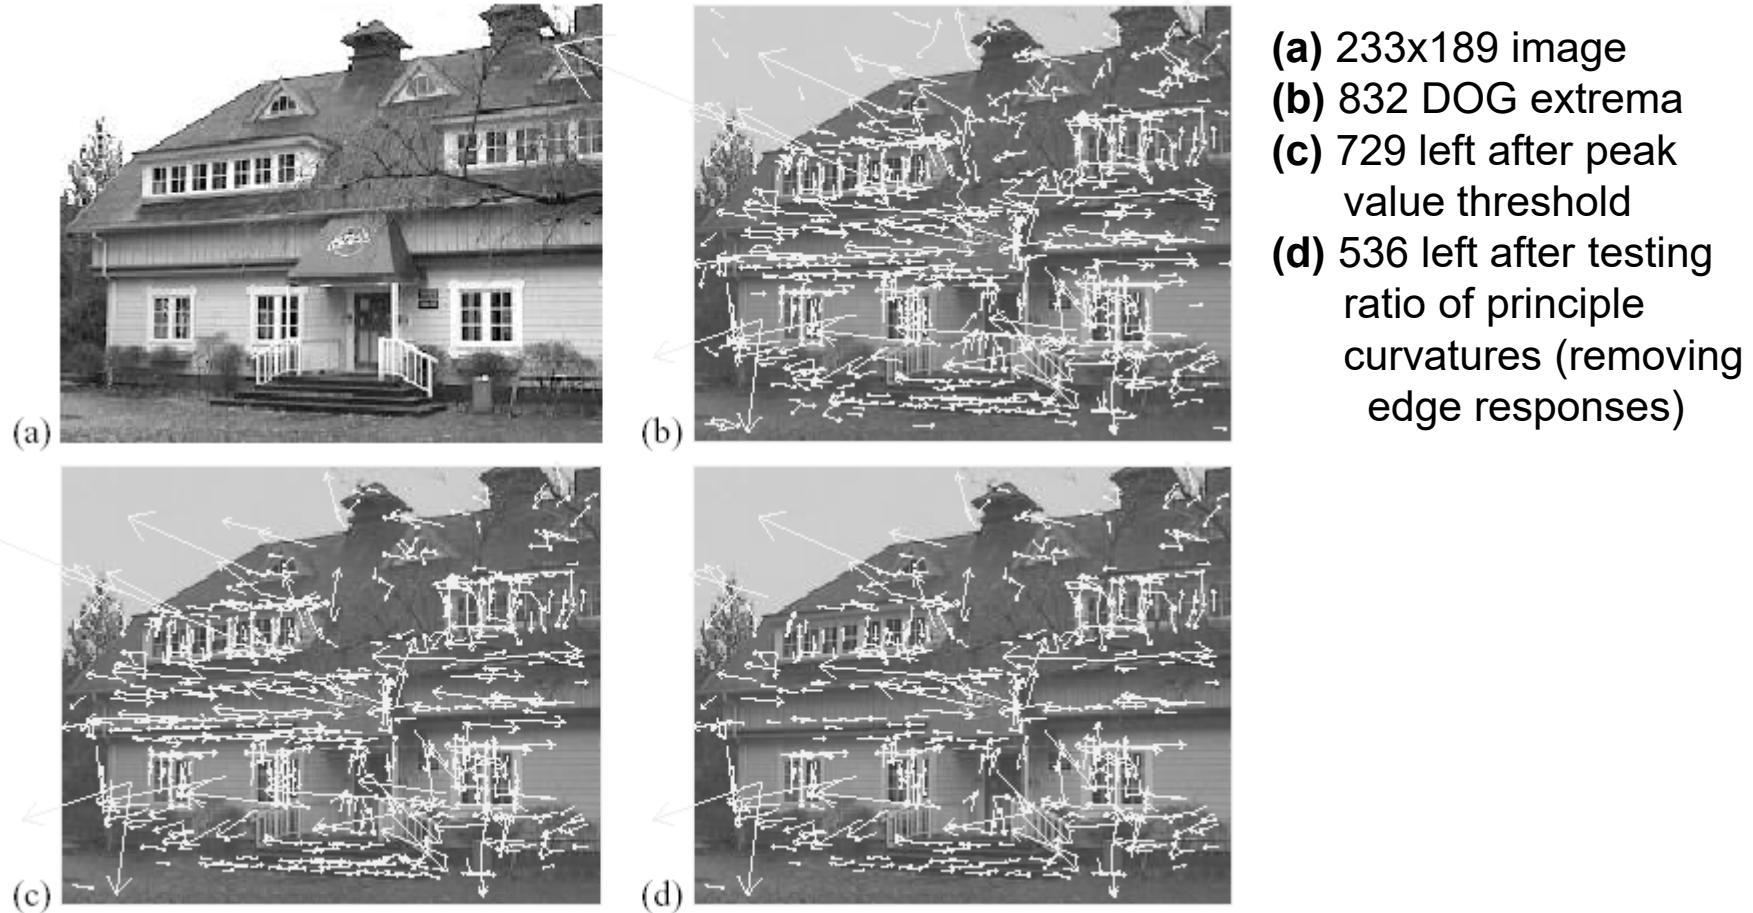
DoG (SIFT) Detector

- Orientation assignment
 - Create histogram of local gradient directions at selected scale
 - Assign canonical orientation at peak of smoothed histogram
- Each key specifies stable 2D coordinates
(x, y, scale, orientation)



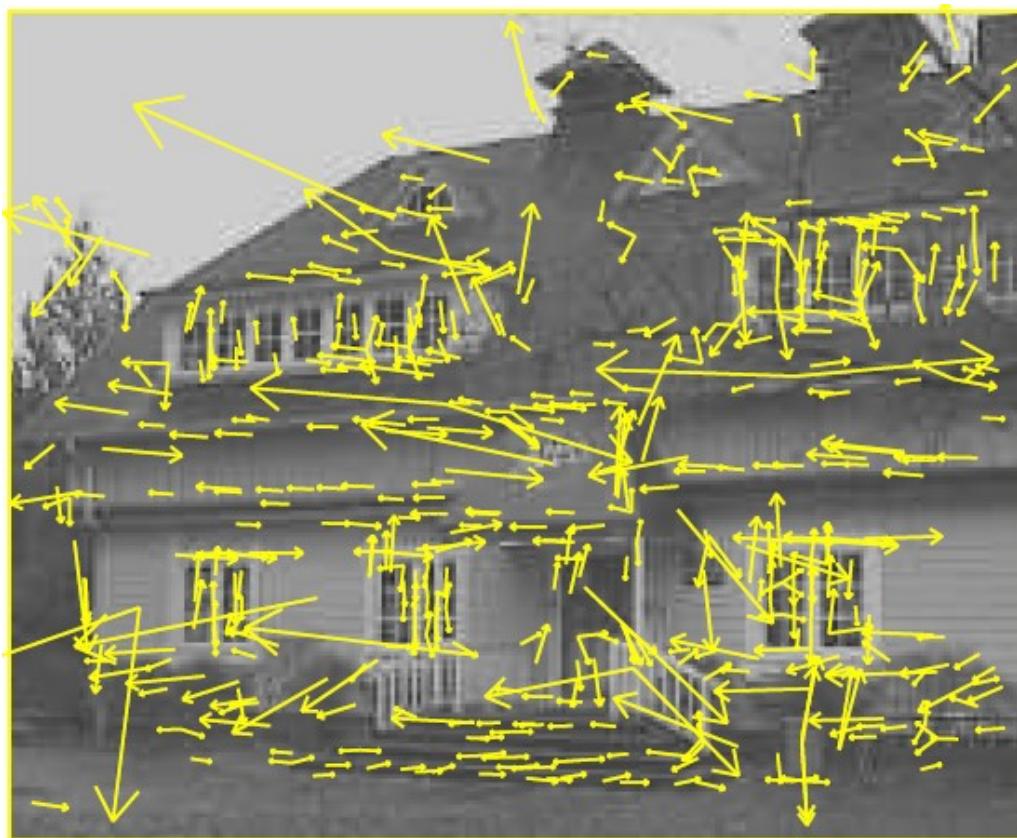
If 2 major orientations, use both.

Example of keypoint detection



- (a) 233x189 image
- (b) 832 DOG extrema
- (c) 729 left after peak value threshold
- (d) 536 left after testing ratio of principle curvatures (removing edge responses)

DoG (SIFT) Detector



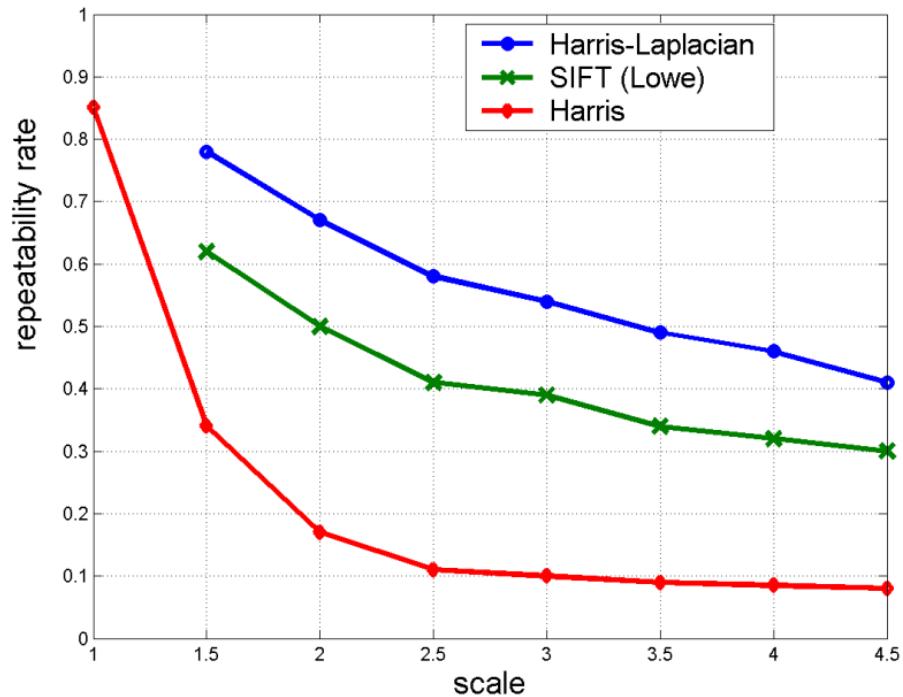
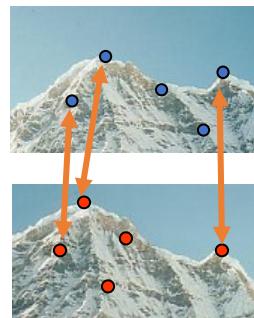
A SIFT keypoint : {**x, y, scale, dominant orientation**}

Scale Invariant Detectors

- Experimental evaluation of detectors w.r.t. scale change

Repeatability rate:

$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



K.Mikolajczyk, C.Schmid. "Indexing Based on Scale Invariant Interest Points". ICCV 2001

Many existing detectors available

- Hessian & Harris [Beaudet '78], [Harris '88]
- Laplacian, DoG [Lindeberg '98], [Lowe '99]
- Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
- Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
- EBR and IBR [Tuytelaars & Van Gool '04]
- MSER [Matas '02]
- Salient Regions [Kadir & Brady '01]
- Others...
- *Those detectors have become a basic building block for many recent applications in Computer Vision.*

Feature extraction

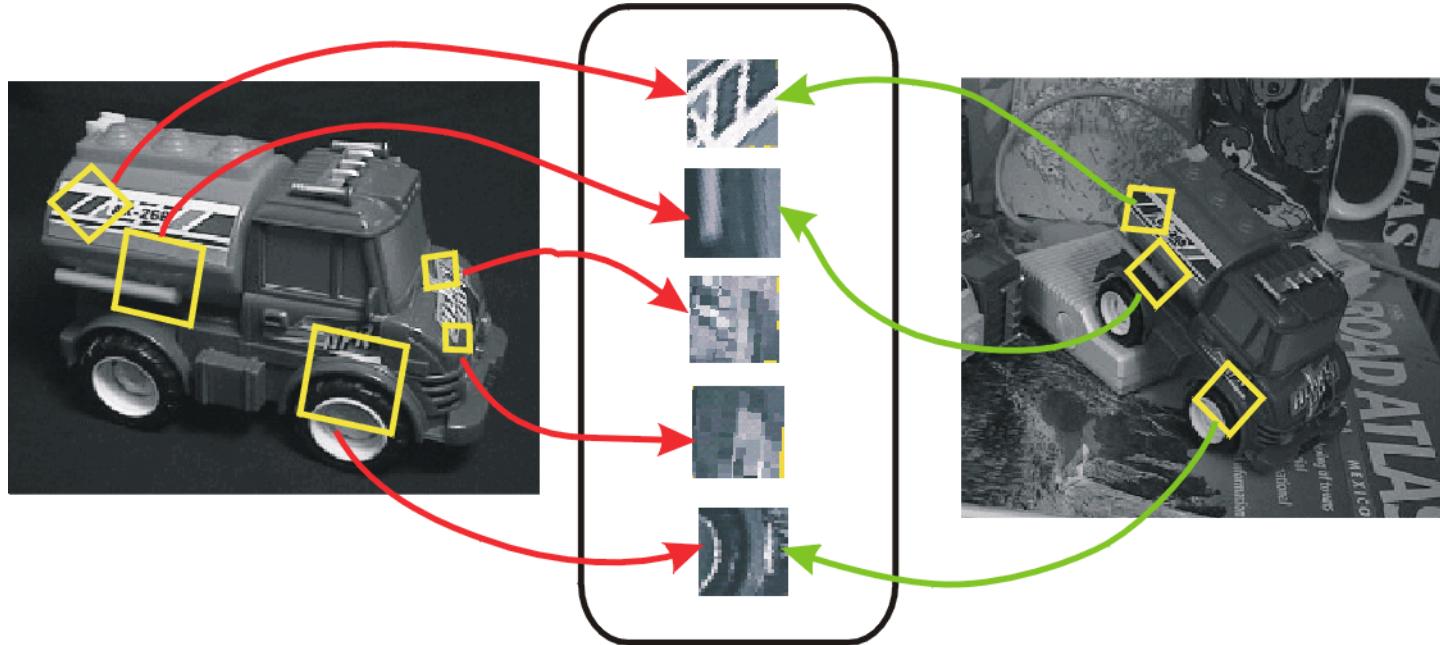
- Global features
- **Local features**
 - Interest point detector
 - **Local descriptor**
 - Matching

Local Descriptor

- Compact, good representation for local information
- Invariant
 - Geometric transformations: rotation, translation, scaling,...
 - Camera view change
 - Illumination
- Examples
 - SIFT, SURF([Speeded Up Robust Features](#)), [PCA-SIFT](#), ...
 - LBP, BRISK, MSER and FREAK, ...

Invariant local features

- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



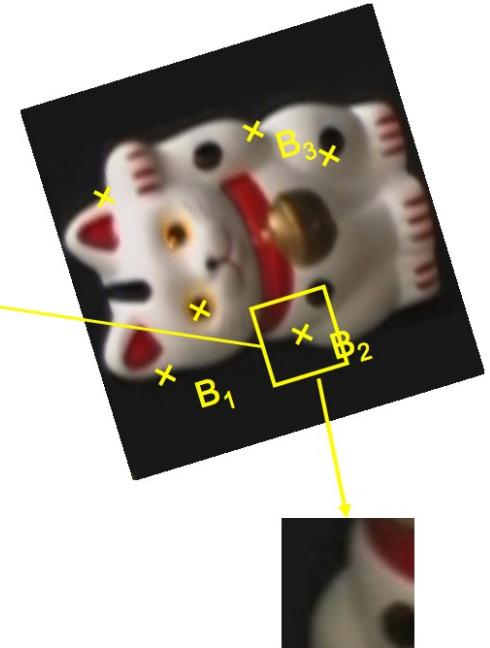
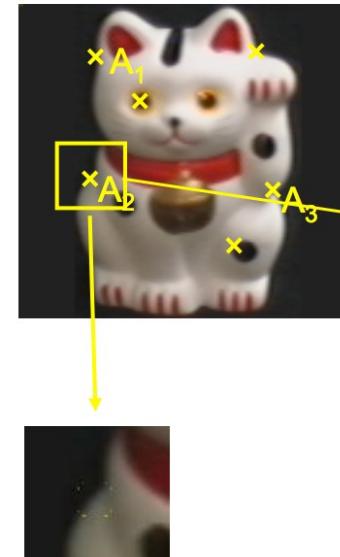
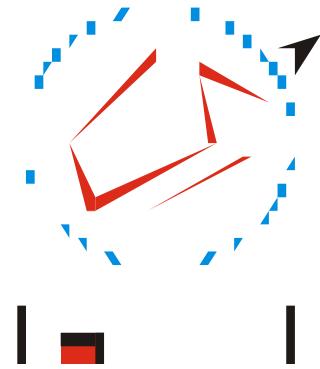
Following slides credit: CVPR 2003 Tutorial on **Recognition and Matching Based on Local Invariant Features** David Lowe

Advantages of invariant local features

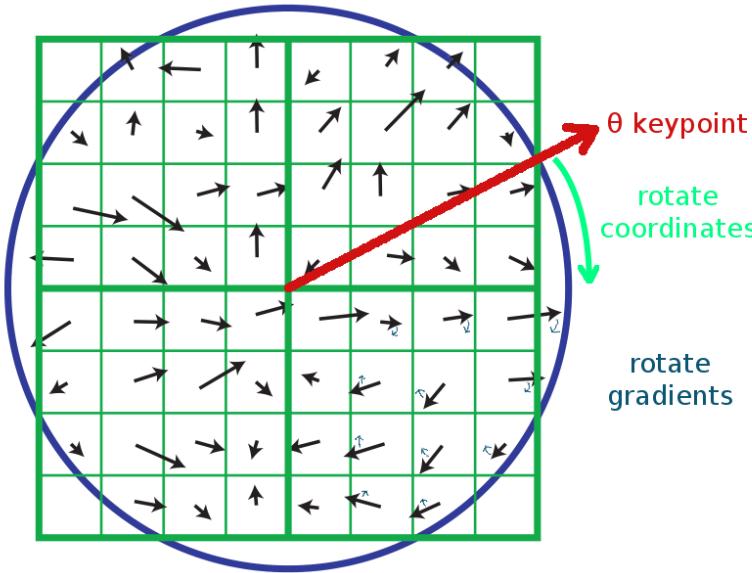
- **Locality:**
 - features are local, so robust to occlusion and clutter (no prior segmentation)
- **Distinctiveness:**
 - individual features can be matched to a large database of objects
- **Quantity:**
 - many features can be generated for even small objects
- **Efficiency:**
 - close to real-time performance
- **Extensibility:**
 - can easily be extended to wide range of differing feature types, with each adding robustness

Becoming rotation invariant

- We are given a keypoint and its scale from DoG
- We will select a **characteristic orientation** for the keypoint (based on the most prominent gradient there)
- We will describe all features **relative** to this orientation



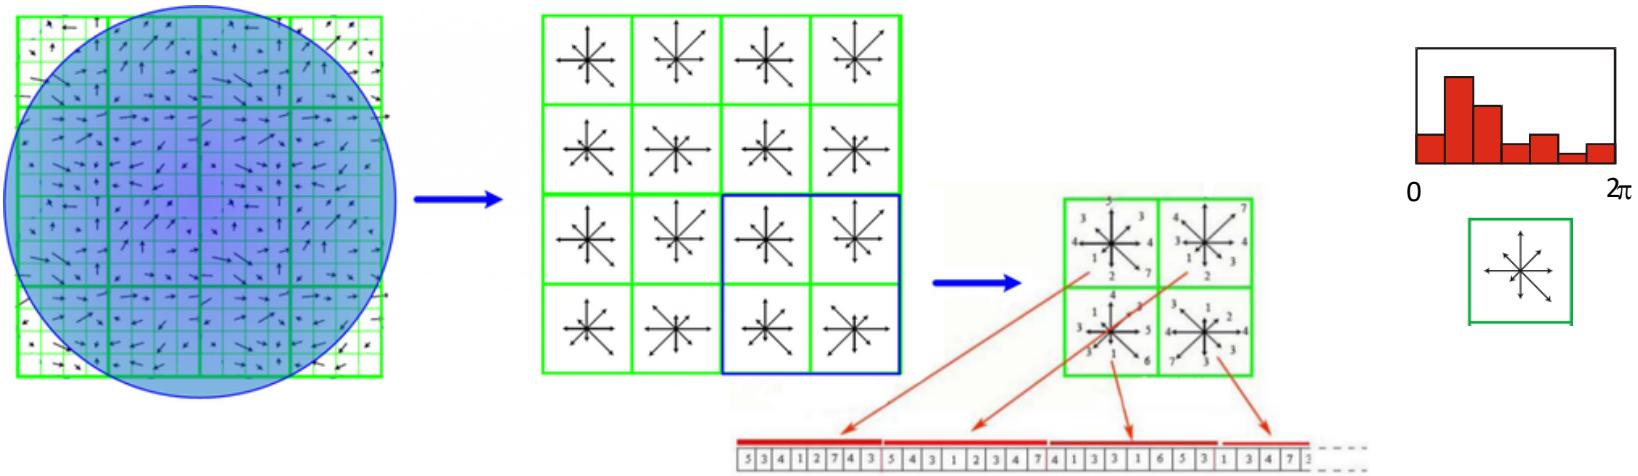
SIFT descriptor formation



- Use the blurred image associated with the keypoint's scale
- Take image gradients over the keypoint neighborhood.
- To become rotation invariant, rotate the gradient **directions AND locations** by (-keypoint orientation)
 - Now we've cancelled out rotation and have gradients expressed at locations **relative** to keypoint orientation θ
 - We could also have just rotated the whole image by $-\theta$, but that would be slower.

Source: [Distinctive Image Features from Scale-Invariant Keypoints](#) – IJCV 2004
http://campar.in.tum.de/twiki/pub/Chair/TeachingWs13TDCV/feature_descriptors.pdf

SIFT descriptor formation

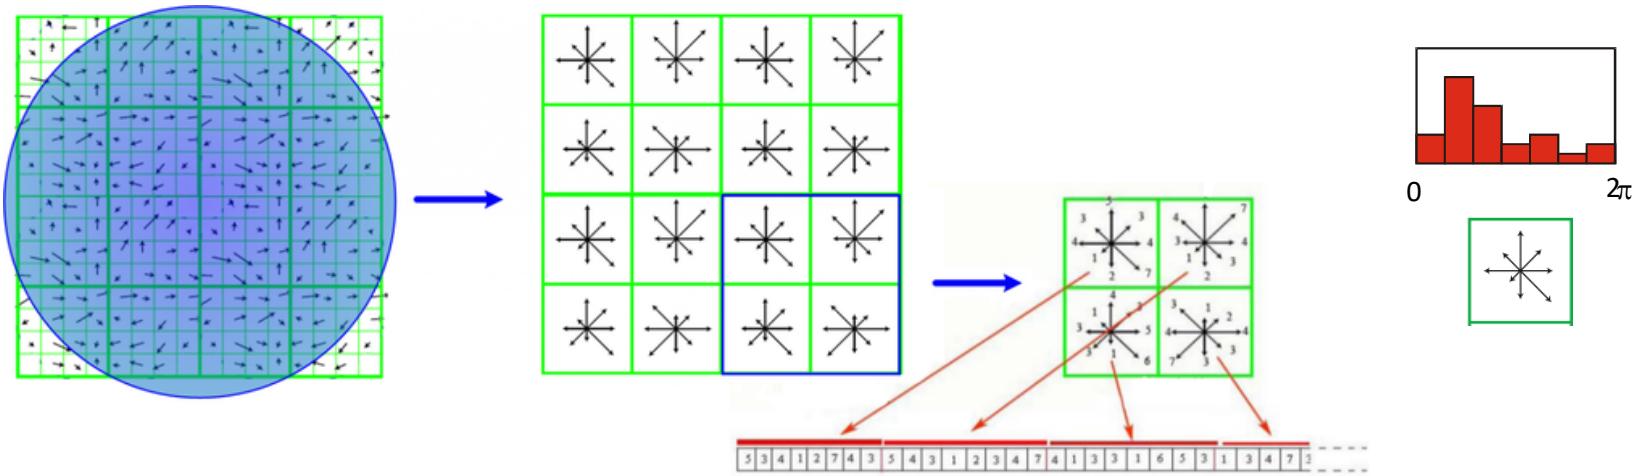


- Using precise gradient locations is fragile. We'd like to allow some “slop” in the image, and still produce a very similar descriptor
- Using Gaussian filter : to **avoid sudden changes** in the descriptor with small changes in the position of the window, and to give **less emphasis to gradients that are far** from the center of the descriptor, as these are most affected by misregistration errors
- Create array of orientation histograms (a 4x4 array is shown)

SIFT: [Distinctive Image Features from Scale-Invariant Keypoints](#) – IJCV 2004

Image: Ashish A Gupta, PhD thesis 2013

SIFT descriptor formation



- Put the rotated gradients into their local orientation histograms
 - A local orientation histogram has n bins (e.g 8 as shown).
 - To avoid all boundary affects in which the descriptor **abruptly changes** as a sample shifts smoothly from being within one histogram to another or from one orientation to another
 - → interpolation is used to distribute the value of each gradient sample into adjacent histogram bins
 - Also, scale **down gradient contributions** for gradients far from the bin center: a weight of $1-d^2$ where d : distance of the sample from the central value of the bin
- The SIFT authors found that best results were with **8 orientation bins per histogram** and **and a 4x4 histogram array** → a SIFT descriptor: vector of 128 values

SIFT descriptor formation

- Adding robustness to **illumination changes**:
 - The descriptor is made of gradients (differences between pixels)
 - → **already invariant to changes in brightness** (e.g. adding 10 to all image pixels yields the exact same descriptor)
 - A higher -contrast photo will increase the magnitude of gradients linearly
 - → to correct for contrast changes, **normalize the vector** (scale to length 1.0)
 - Very large image gradients are usually from unreliable 3D illumination effects (glare, etc)
 - → to reduce their effect, **clamp all values** in the vector to **be ≤ 0.2** (an experimentally tuned value). Then normalize the vector again.
- Result is a vector which is fairly invariant to illumination changes

SIFT

- Extraordinarily robust matching technique
 - Can handle changes in viewpoint: up to about 60 degree out of plane rotation
 - Can handle significant changes in illumination
 - Sometimes even day vs. night (below)
 - Fast and efficient—can run in real time



Steve Seitz

Sensitivity to number of histogram orientations

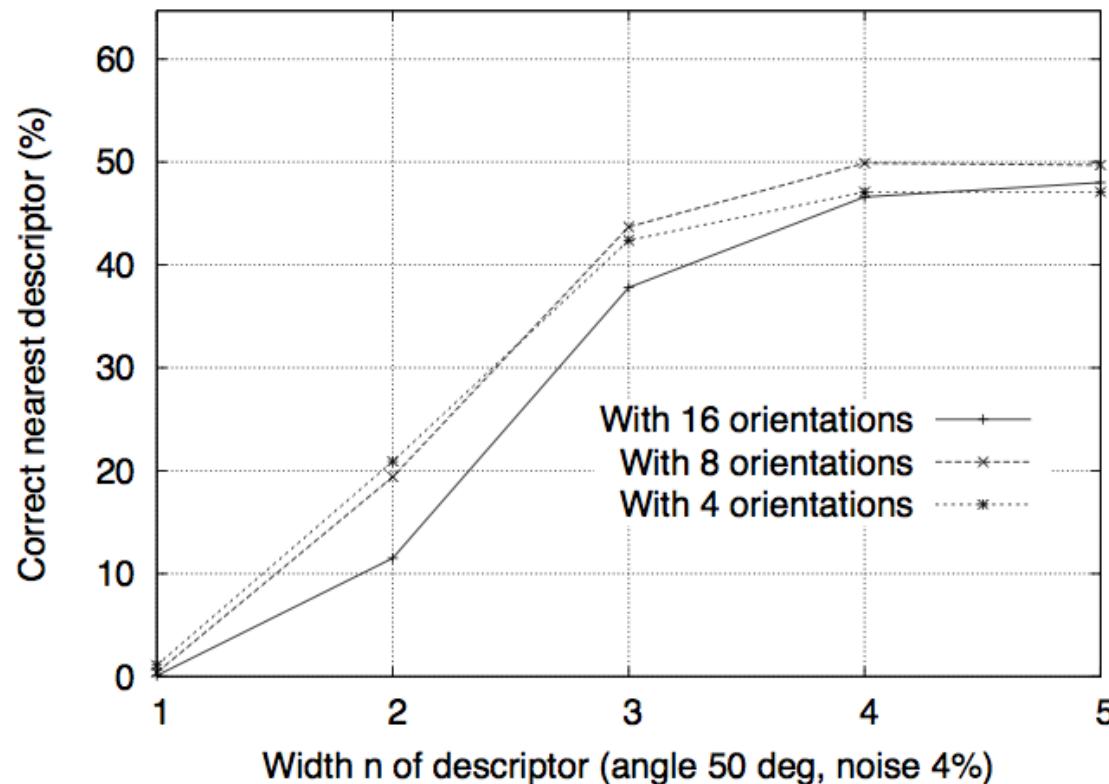
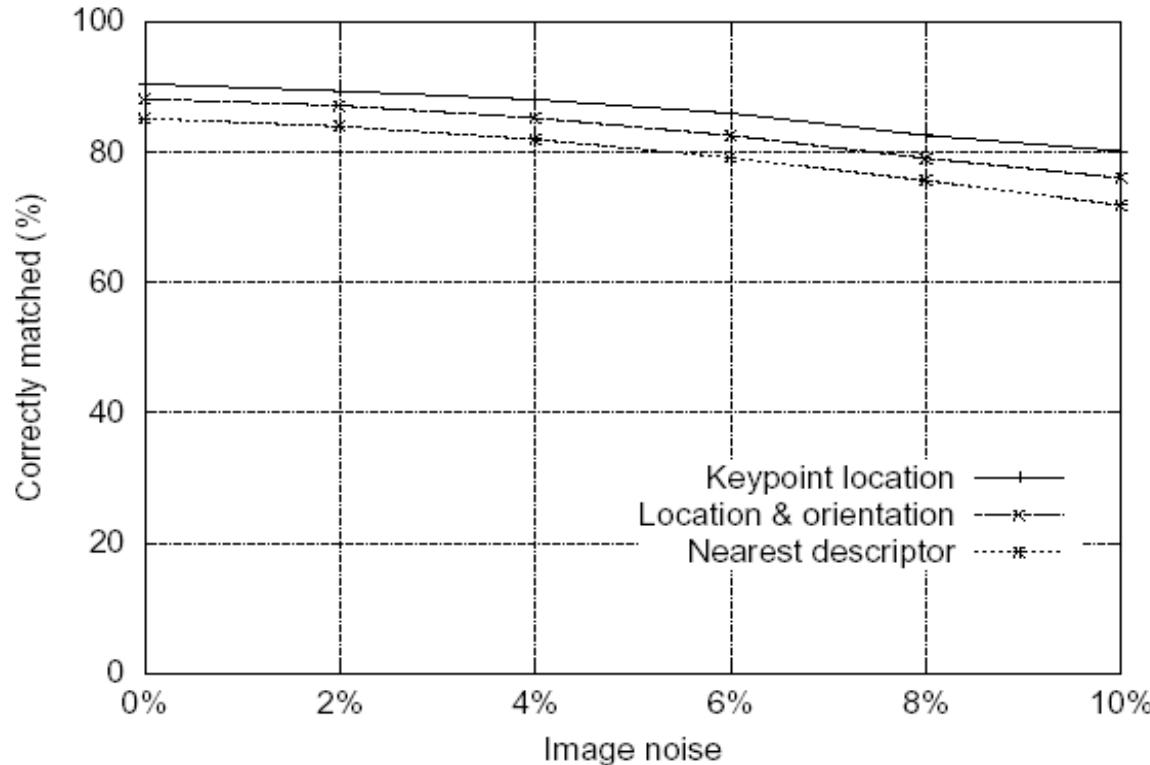


Figure 8: This graph shows the percent of keypoints giving the correct match to a database of 40,000 keypoints as a function of width of the $n \times n$ keypoint descriptor and the number of orientations in each histogram. The graph is computed for images with affine viewpoint change of 50 degrees and addition of 4% noise.

Feature stability to noise

- Match features after random change in image scale & orientation, with differing levels of image noise
- Find nearest neighbor in database of 30,000 features



David G. Lowe, "Distinctive image features from scale-invariant keypoints," IJCV, 60, 2 (2004), pp. 91-110
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Feature stability to affine change

- Match features after random change in image scale & orientation, with 2% image noise, and affine distortion
- Find nearest neighbor in database of 30,000 features

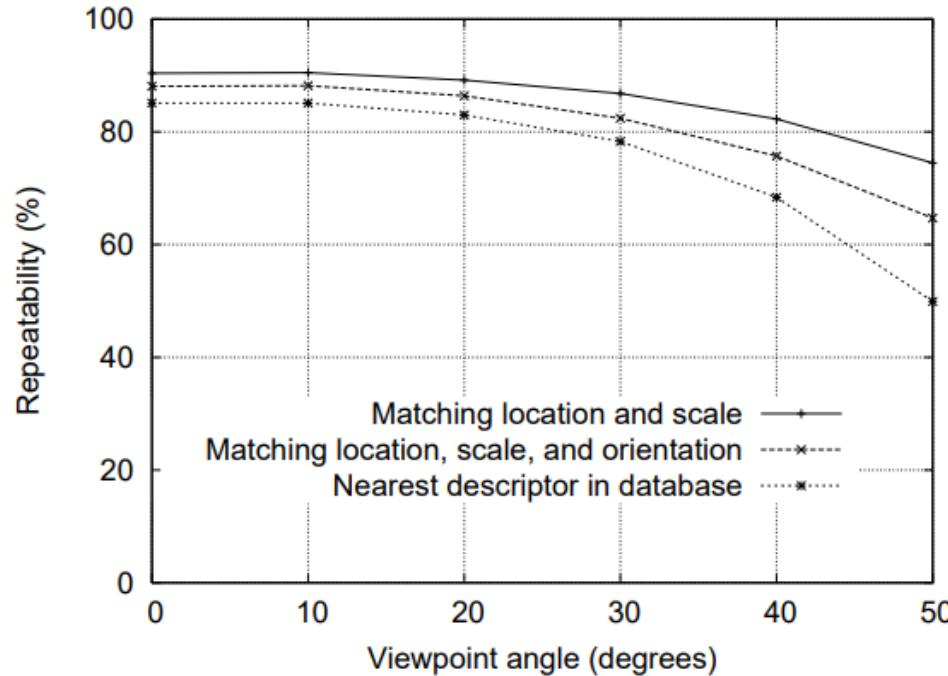
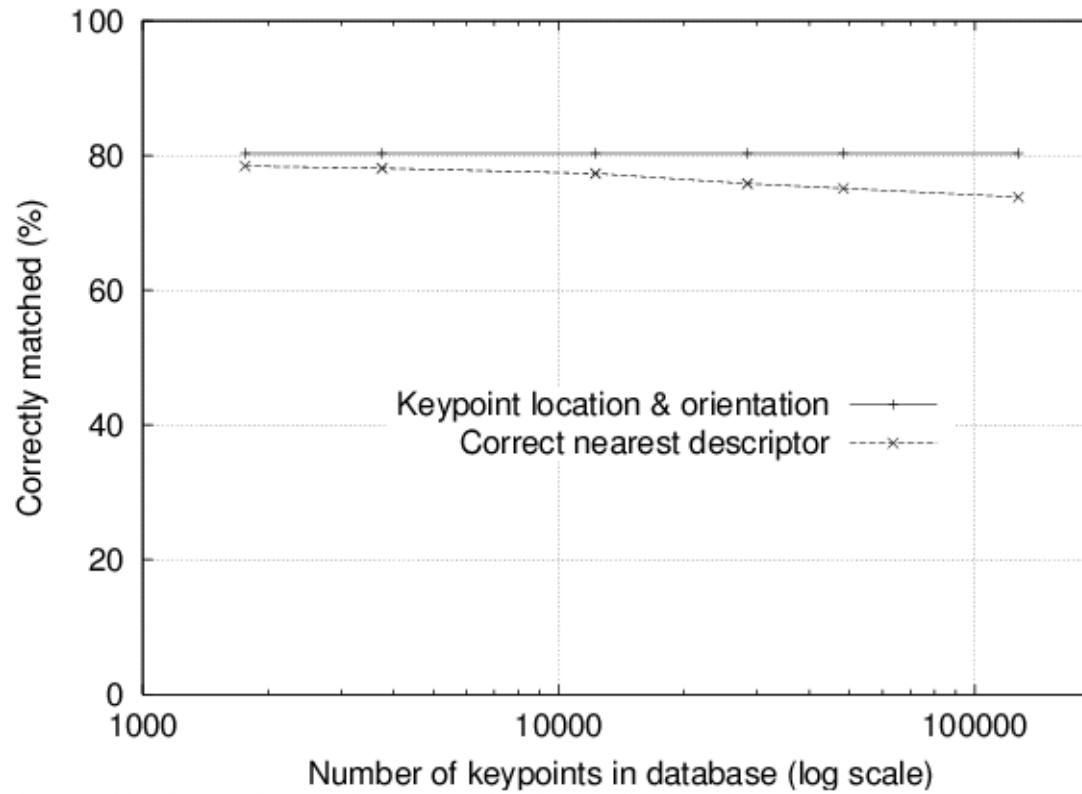


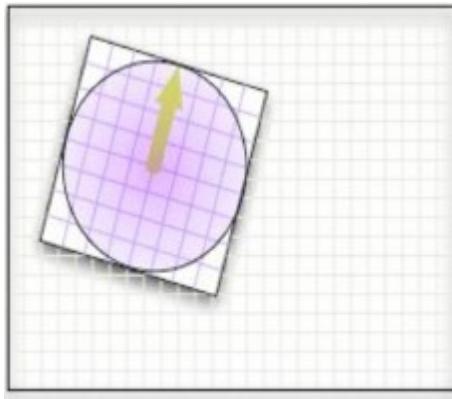
Figure 9: This graph shows the stability of detection for keypoint location, orientation, and final matching to a database as a function of affine distortion. The degree of affine distortion is expressed in terms of the equivalent viewpoint rotation in depth for a planar surface.

Distinctiveness of features

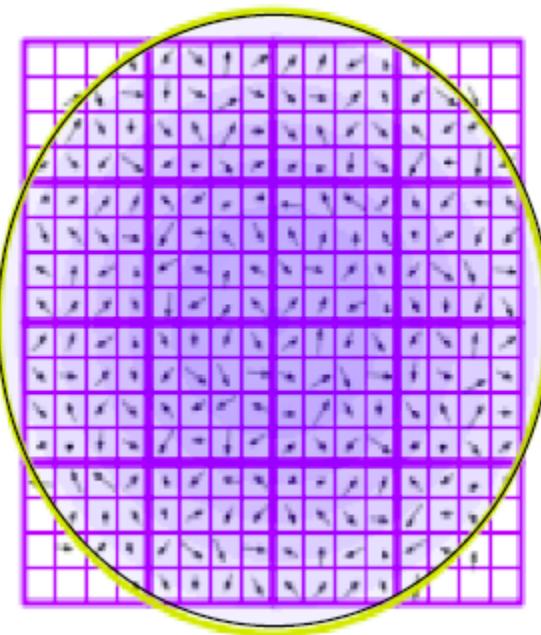
- Vary size of database of features, with 30 degree affine change, 2% image noise
- Measure % correct for single nearest neighbor match



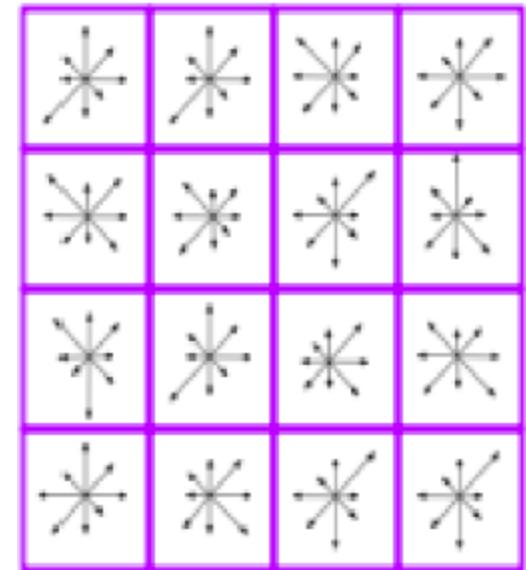
SIFT Keypoint Descriptor: summary



Blur the image
using the scale of
the keypoint
(scale invariance)



Compute gradients in
respect to the keypoint
orientation(rotation
invariance)



Compute orientation
histogram in 8
directions over 4x4
sample regions

Other detectors and descriptors

Popular features: SURF, HOG, SIFT

http://campar.in.tum.de/twiki/pub/Chair/TeachingWs13TDCV/feature_descriptors.pdf

Summary some local features:

http://www.cse.iitm.ac.in/~vplab/courses/CV_DIP/PDF/Feature_Detectors_and_Descriptors.pdf

Feature extraction

- Global features
- **Local features**
 - Interest point detector
 - Local descriptor
 - Matching

Feature matching

Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
 - Use L1, L2, cosine, Mahalanobis,... distance
2. Test all the features in I_2 , find the one with min distance

OpenCV:

- Brute force matching
- Flann Matching: Fast Library for Approximate Nearest Neighbors
[Muja and Lowe, 2009]

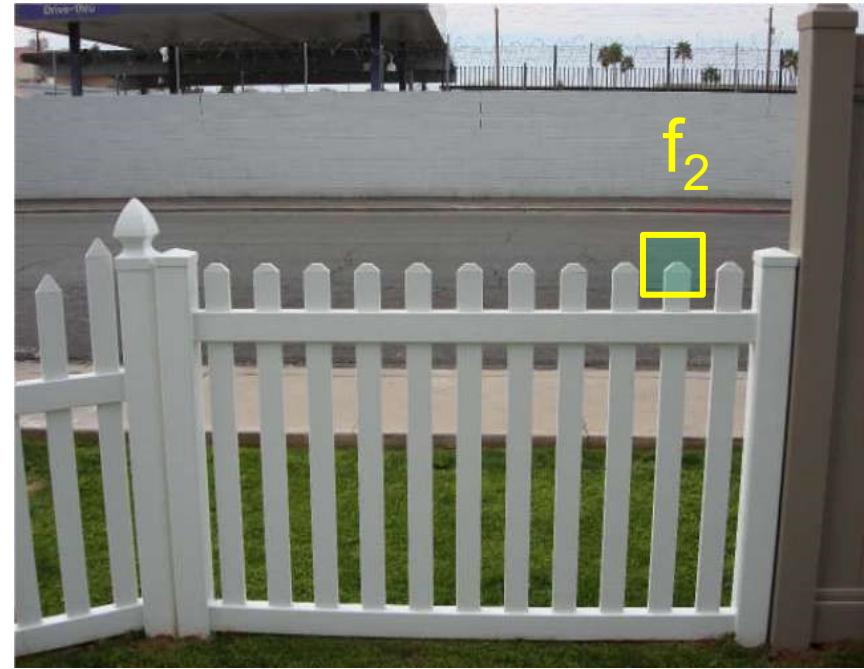
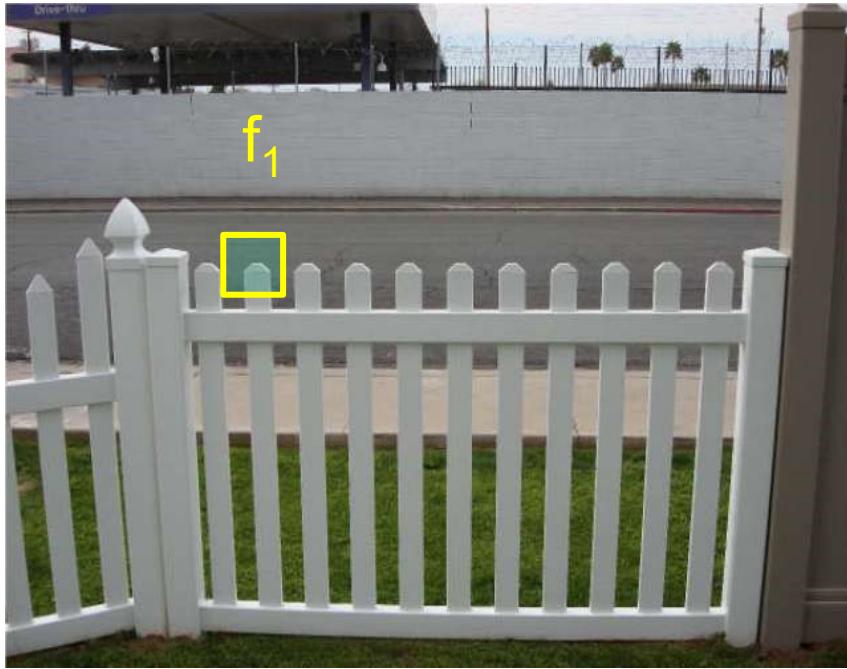
Marius Muja and David G Lowe. Fast approximate nearest neighbors with automatic algorithm configuration. In VISAPP (1), pages 331–340, 2009

Feature matching

- How to define the difference between two features f_1, f_2 ?
 - Simple approach: **use only distance value $d(f_1, f_2)$**
 - → can give good score to very ambiguous matches
 - Better approaches: add **additional constraints**
 - Ratio of distance
 - Spatial constraints between neighborhood pixels
 - Fitting the transformation, then refine the matches (RANSAC)

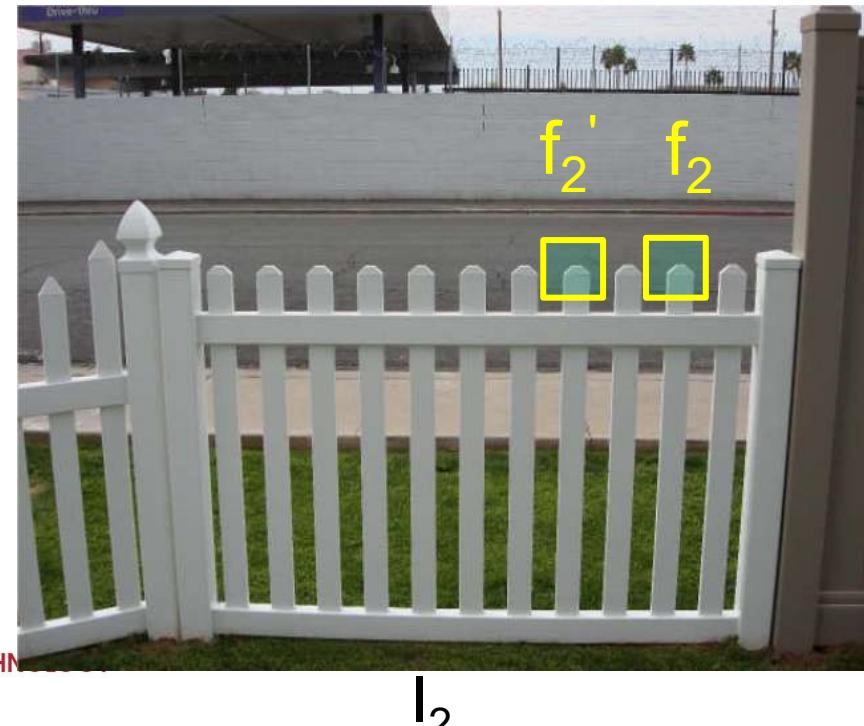
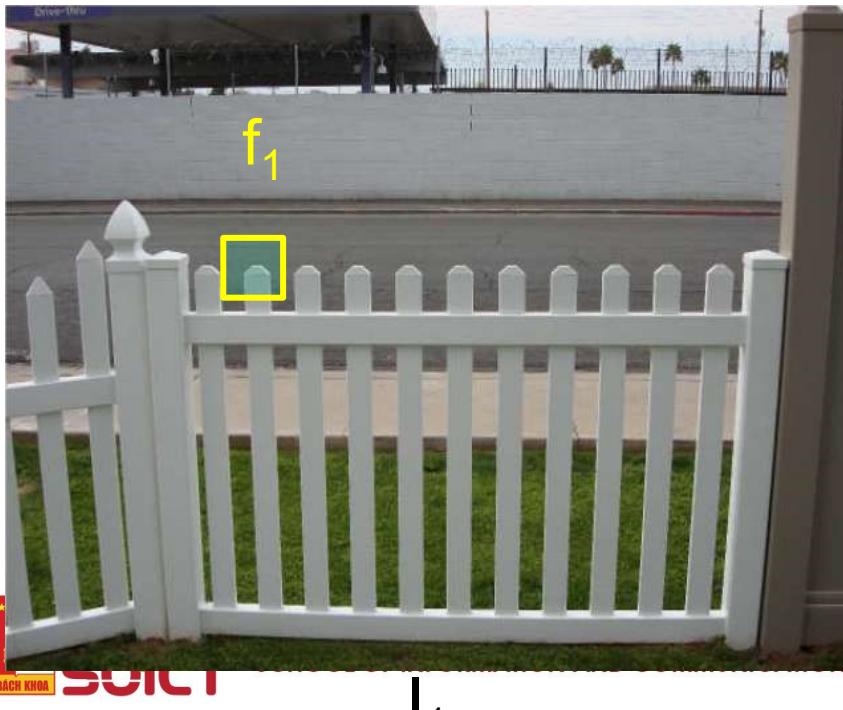
Feature matching

- Simple approach: use distance value $d(f_1, f_2)$
→ can give good score to very ambiguous matches



Feature matching

- Better approaches: **ratio of distance = $d(f_1, f_2) / d(f_1, f_2')$**
 - f_2 is best match to f_1 in I_2 ;
 - f_2' is 2nd best SSD match to f_1 in I_2
 - An ambiguous/bad match will have ratio close 1
 - Look for unique matches which have low ratio



Ratio of distances reliable for matching

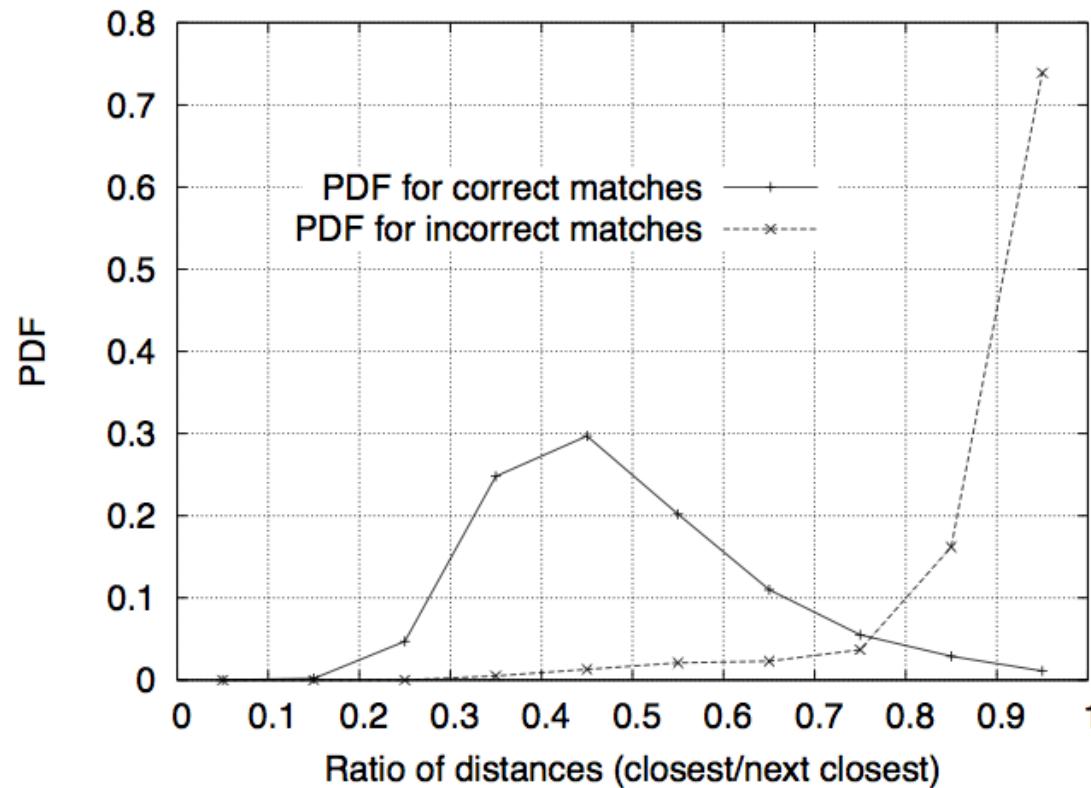
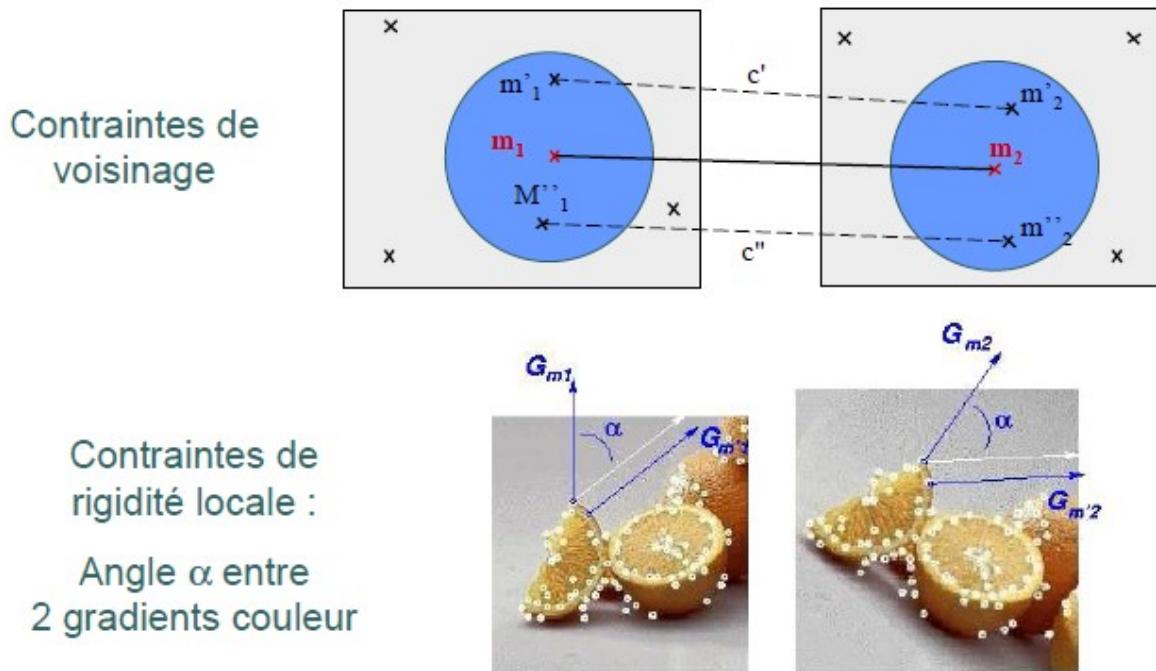


Figure 11: The probability that a match is correct can be determined by taking the ratio of distance from the closest neighbor to the distance of the second closest. Using a database of 40,000 keypoints, the solid line shows the PDF of this ratio for correct matches, while the dotted line is for matches that were incorrect.

Feature matching

- Better approaches: Spatial constraints between neighborhood pixels



Source: from slides of Valérie Gouet-Brunet

Feature matching

- Better approaches: fitting the transformation (RANSAC alg.)
 - Fitting 2D transformation matrix

- Six variables

- Each point give two equations
 - → at least three points

- Least squares

$$H = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ 0 & 0 & 1 \end{bmatrix}$$

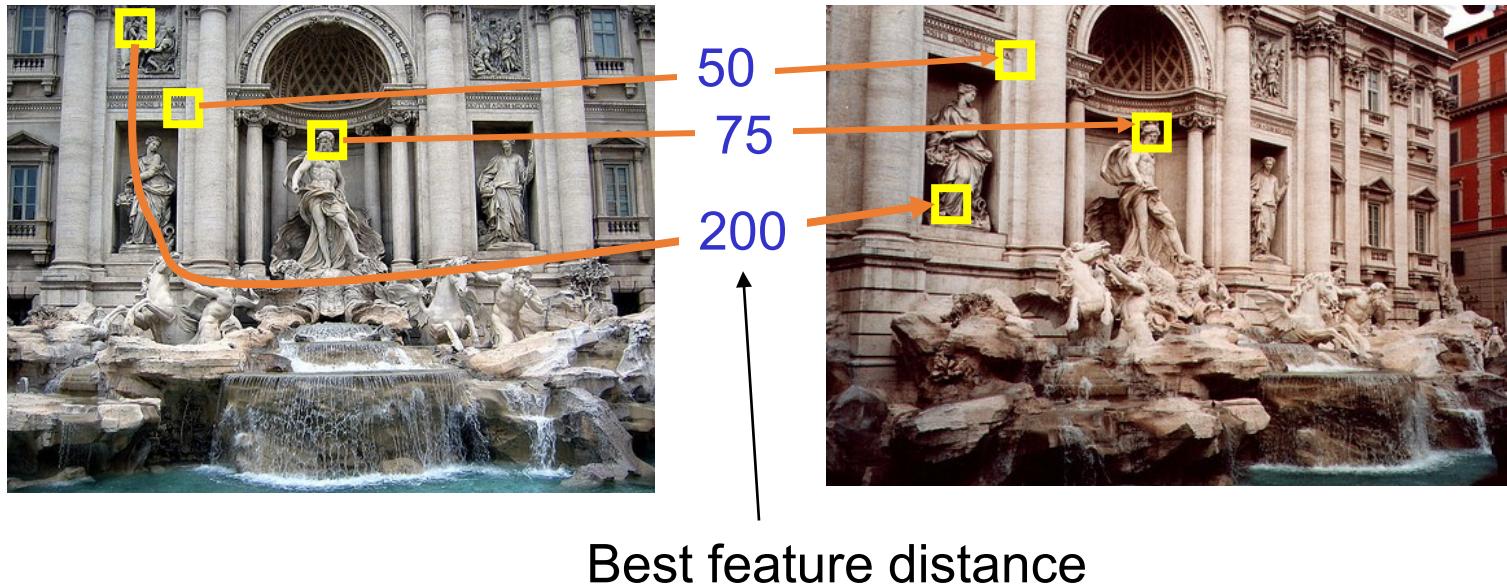
- RANSAC: refinement of matches

- Compute error:

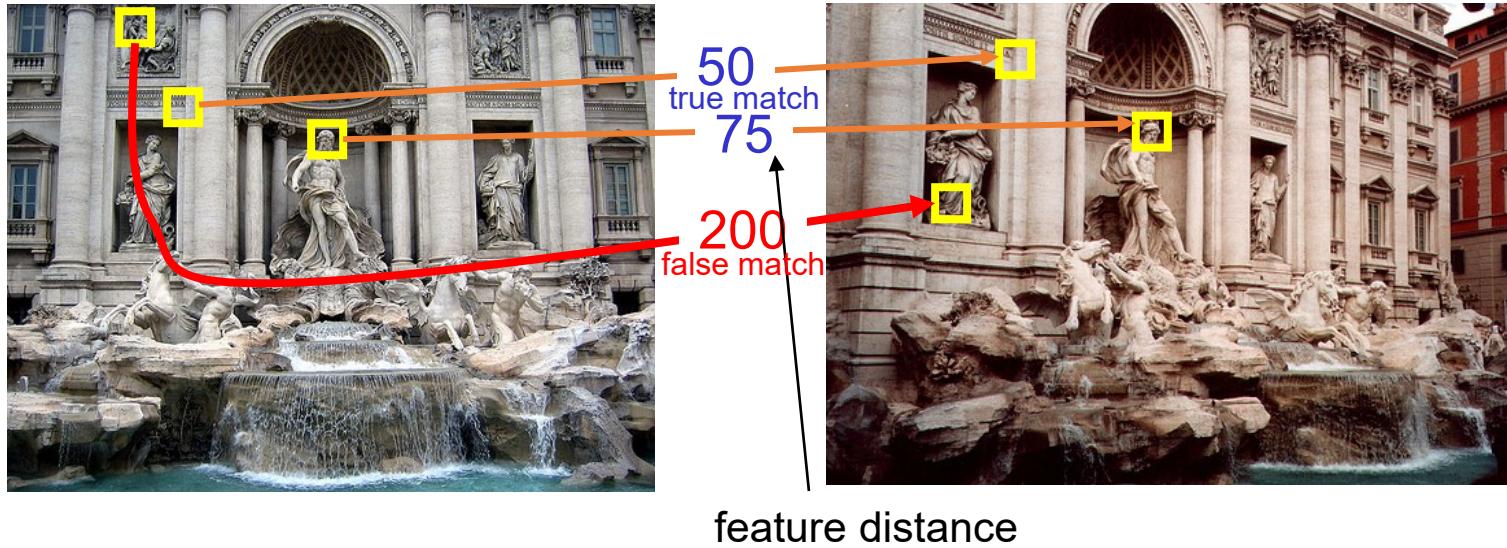
$$\left\| \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} - H \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} \right\|_2$$

Evaluating the results

How can we measure the performance of a feature matcher?



True/false positives



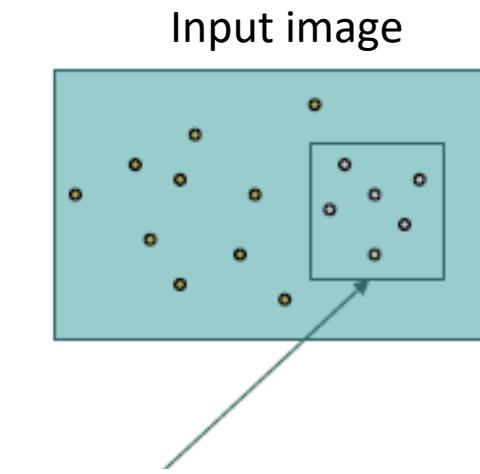
The distance threshold affects performance

- **True positives** = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- **False positives** = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Image matching

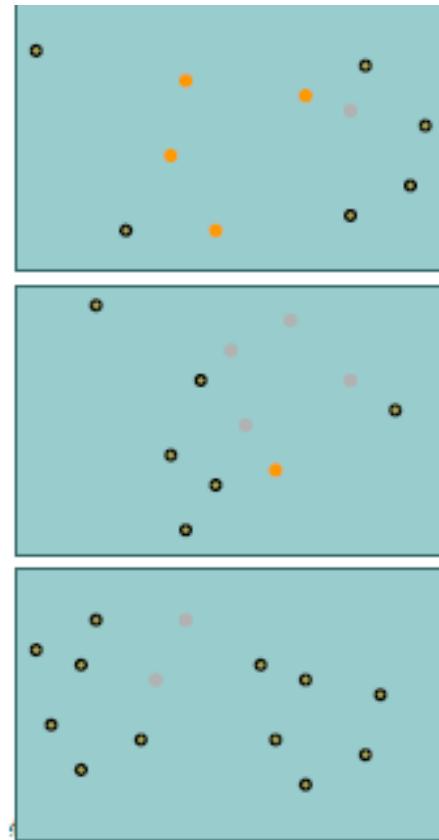
- How to define the distance between 2 images I_1, I_2 ?
 - Using global features: easy
$$d(I_1, I_2) = d(\text{feature of } I_1, \text{feature of } I_2)$$
 - Using local features:
 - Voting strategy
 - Solving an optimization problem (time consuming)
 - Building a "global" feature from local features: BoW (bag-of-words, bag-of-features), VLAD, ..

Voting strategy



Selected region
for query

Images in a database



- Distance faible $\leq \varepsilon$
- Distance plus élevée $\leq \varepsilon$
- Distance $> \varepsilon$



The similarity between 2 images is based on **the number of matches**

Optimization problem

- *Transportation problem*

$I_1 : \{(r_i, w_i), i=1, N\}$ Provider

$I_2 : \{(r'_j, w_j), j=1, M\}$ Consommier

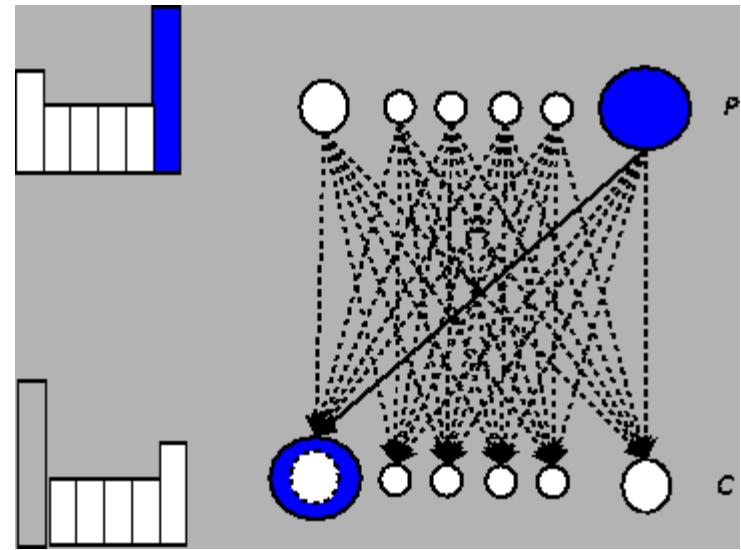
$d(I_1, I_2) = ???$

$$d(I_1, I_2) = \min \sum_i \sum_j f_{ij} \times d(r_i, r'_j)$$

$$f_{ij} \geq 0$$

$$\sum_i f_{ij} \leq w_j, \sum_j f_{ij} \leq w_i$$

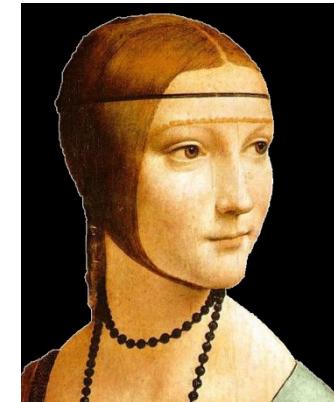
$$\sum_i \sum_j f_{ij} = \min(\sum_i w_i, \sum_j w_j)$$



$$d_{EMD}(I_1, I_2) = \frac{\sum_i \sum_j f_{ij}^* \times d(r_i, r'_j)}{\sum_i \sum_j f_{ij}^*}$$

Bag-of-words

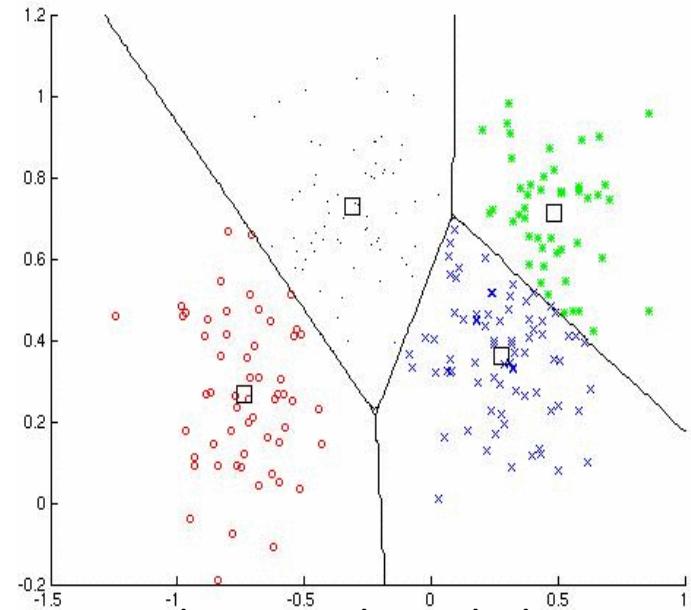
- Local feature ~~ a word
- An image ~~ a document
- Apply a technique for textual document representation: vector model



Visual Vocabulary ...



1. Extracting local features from a set of images

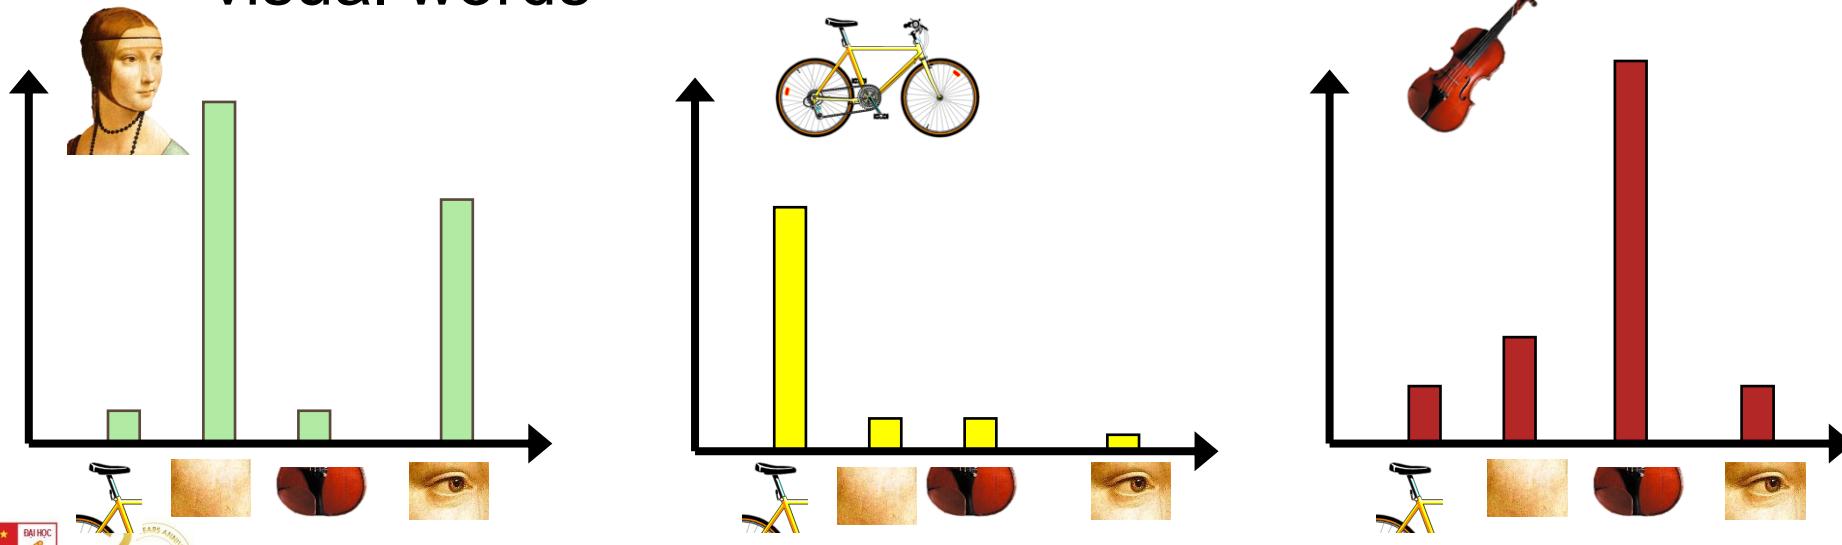


2. Building visual vocabulary (dictionary) using a clustering method

3. An image is represented by a bag of words
→ can be represented by tf.idf vector

Bag of words: outline

1. Extract features
2. Learn “visual vocabulary”
3. Quantize features using visual vocabulary
4. Represent images by frequencies of “visual words”



Applications

Object detection/recognition/search



Sivic and Zisserman, 2003



Rothganger et al. 2003



Lowe 2002

Object detection/recognition

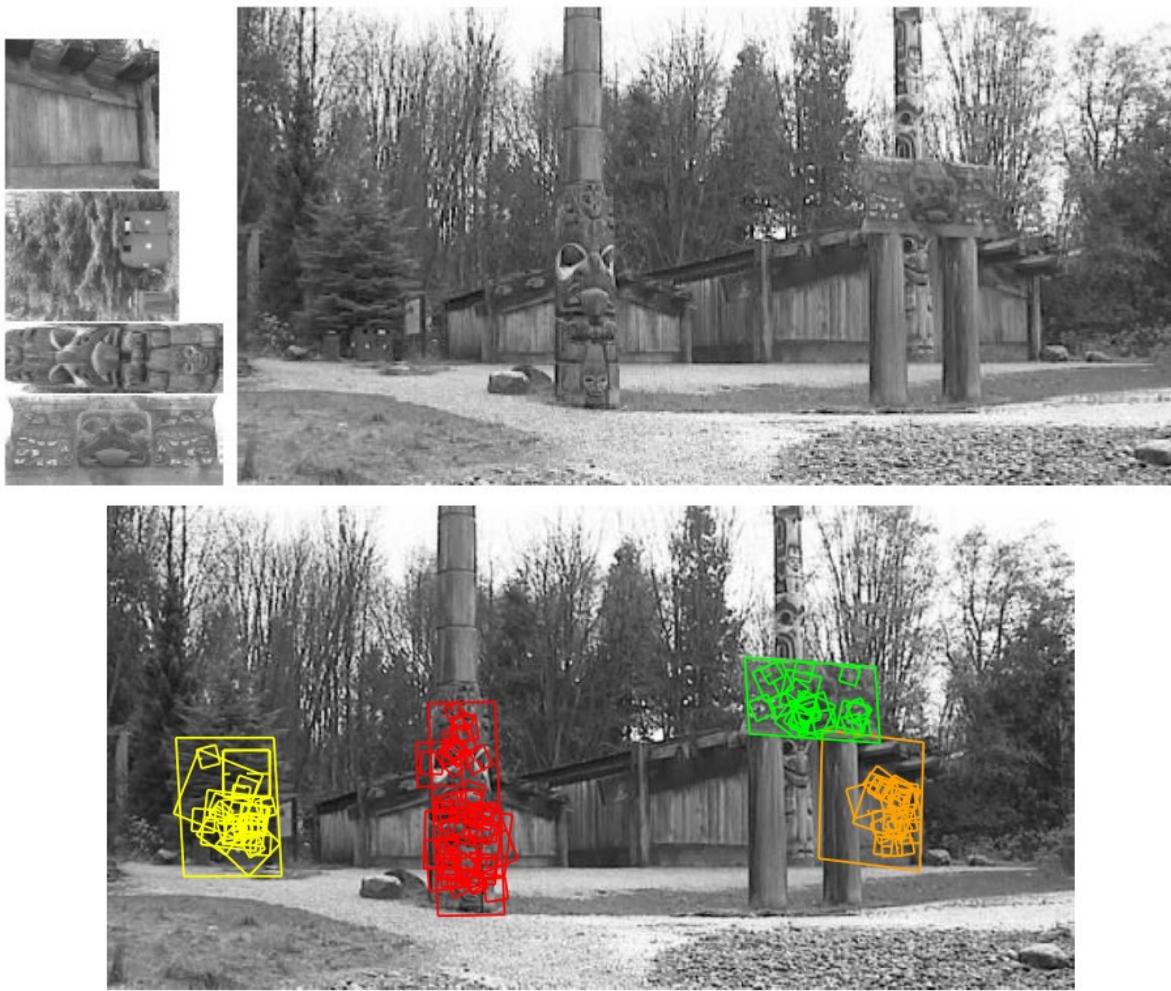


Figure 13: This example shows location recognition within a complex scene. The training images for locations are shown at the upper left and the 640x315 pixel test image taken from a different viewpoint is on the upper right. The recognized regions are shown on the lower image, with keypoints shown as squares and an outer parallelogram showing the boundaries of the training images under the affine transform used for recognition.

Application: Image Panoramas



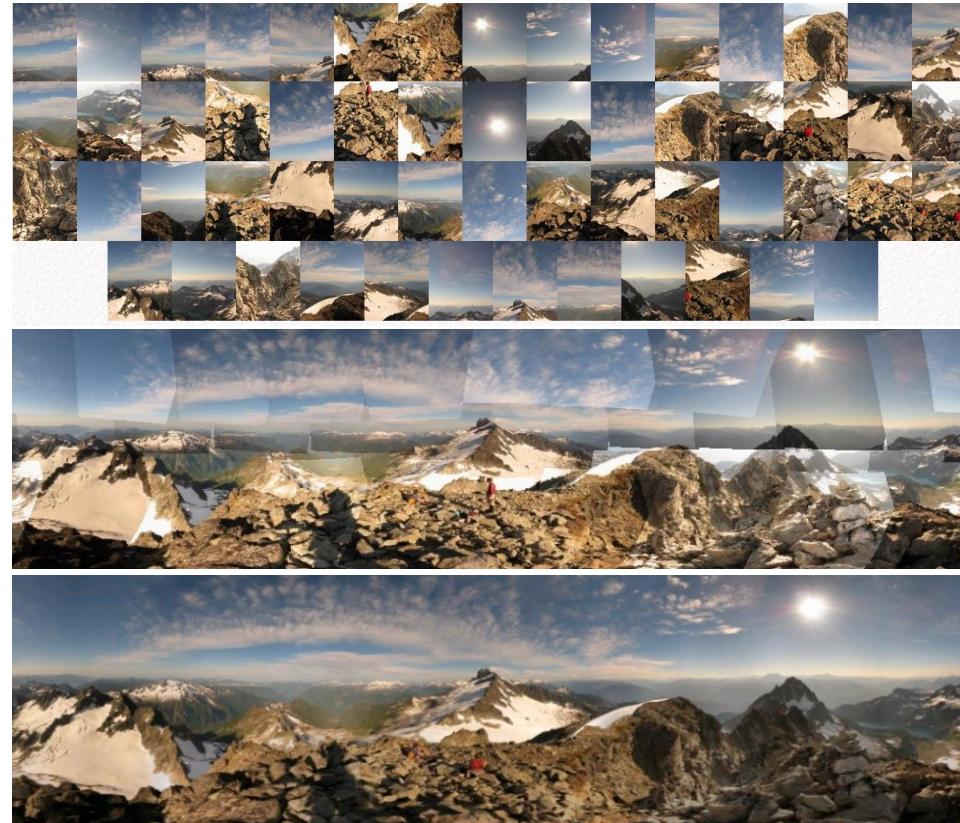
Slide credit: Darya Frolova, Denis Simakov

Application: Image Panoramas



- Procedure:
 - Detect feature points in both images
 - Find corresponding pairs
 - Use these pairs to align the images

Automatic mosaicing



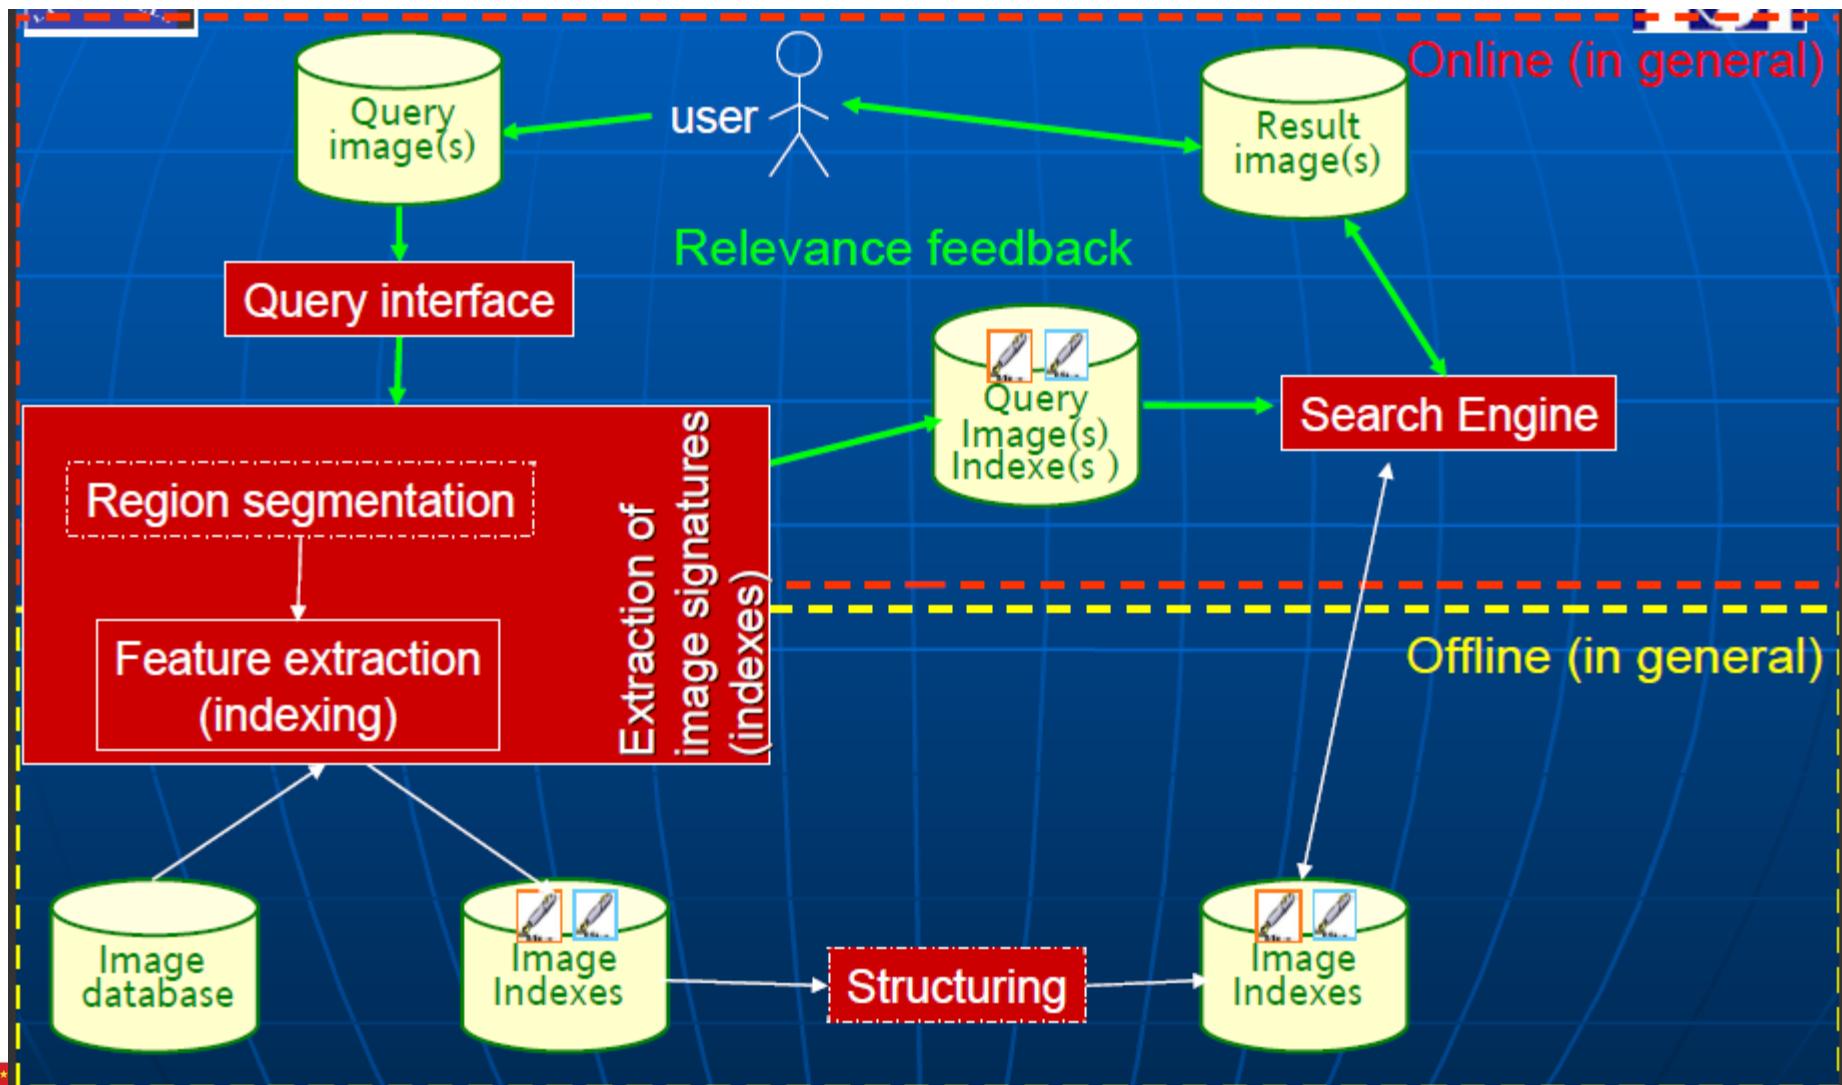
<http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

Wide baseline stereo



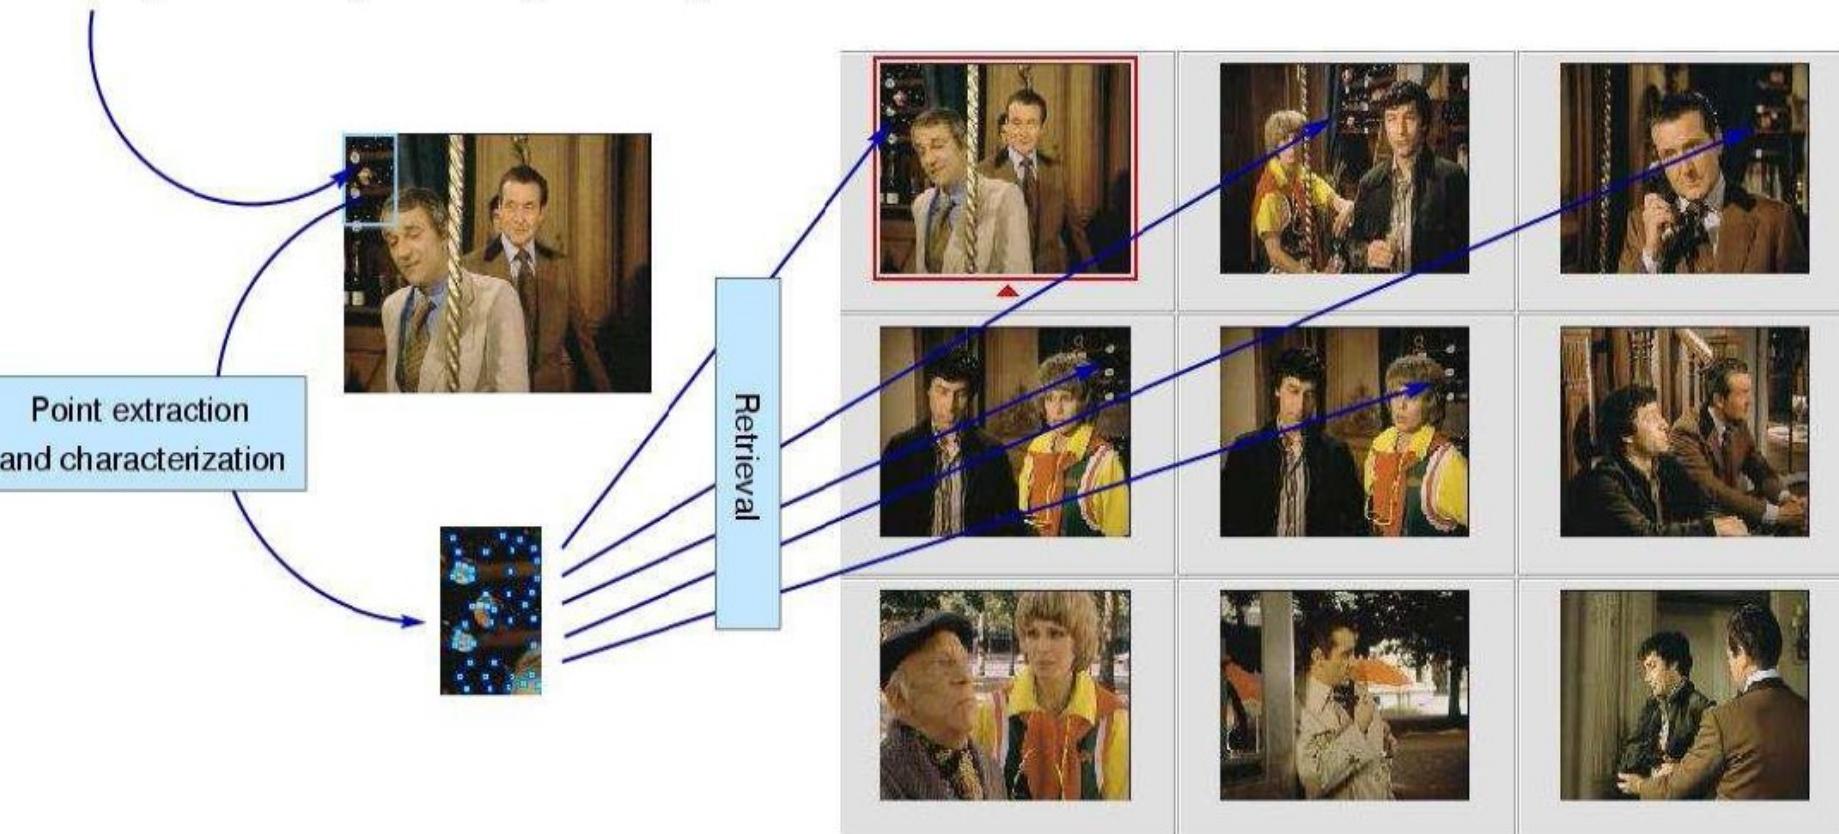
[Image from T. Tuytelaars ECCV 2006 tutorial]

CBIR (Content-based image retrieval)

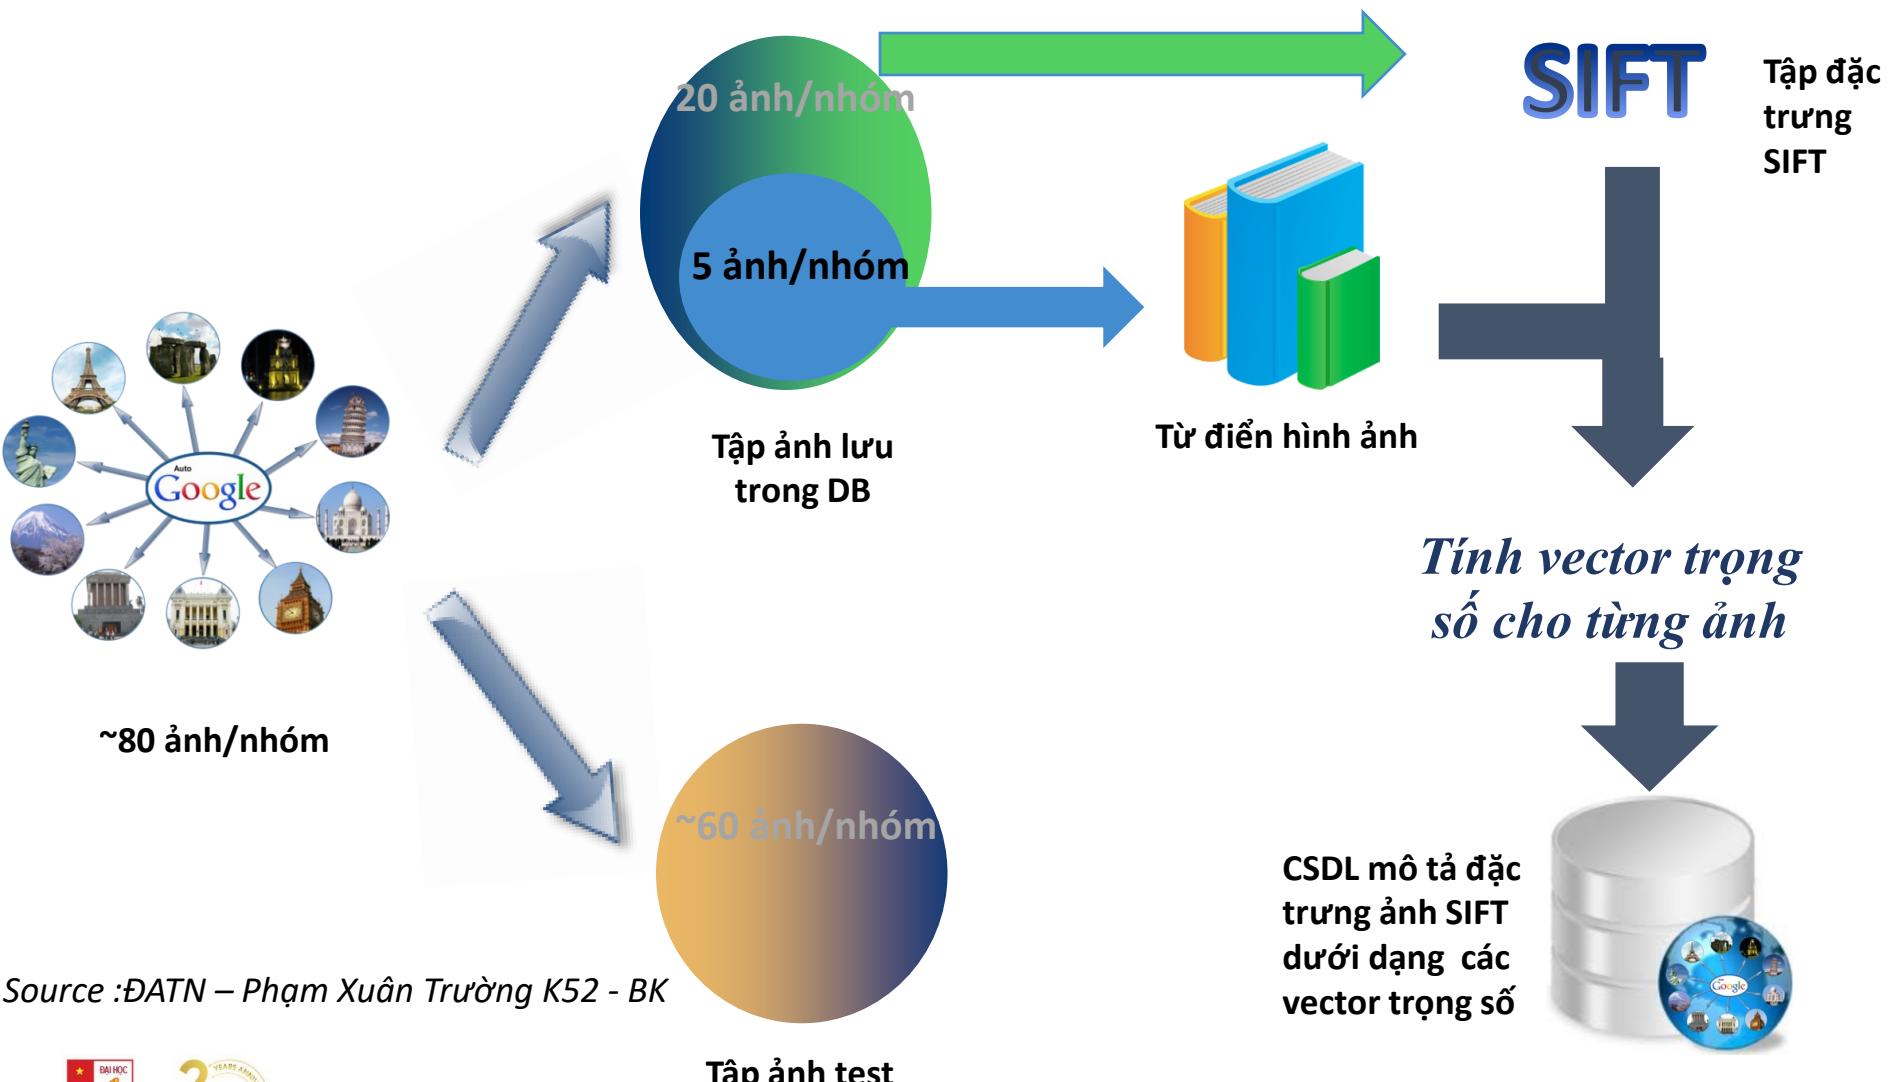


CBIR: partial retrieval

The query : "I am looking for the images involving the room where there is this wine storeroom".

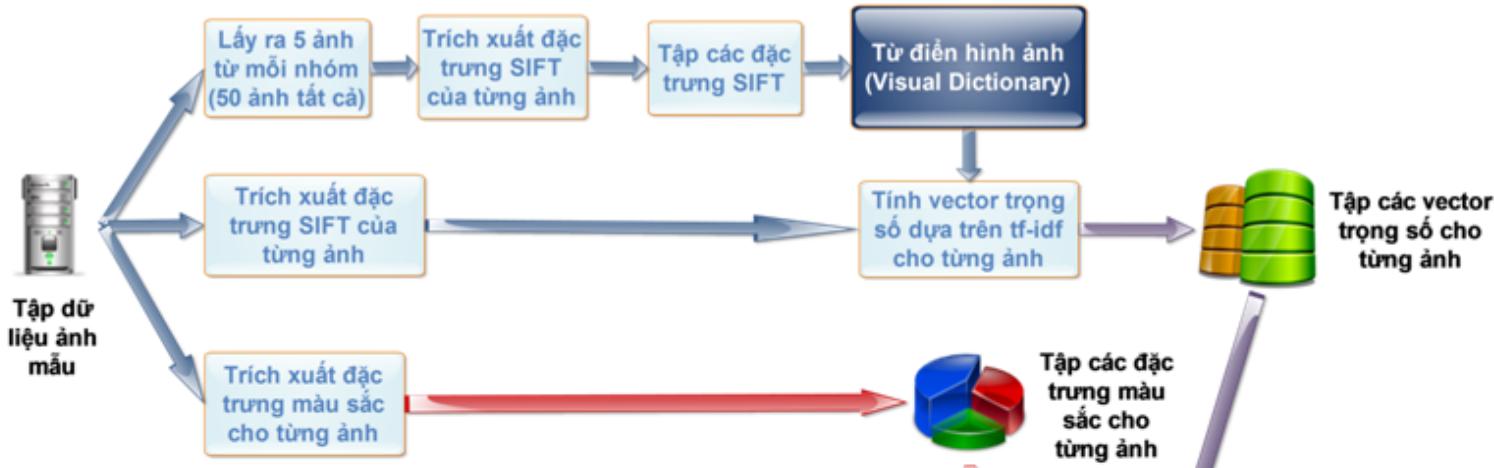


CBIR: BoW with SIFT + histogram



Source :ĐATN – Phạm Xuân Trường K52 - BK

CBIR: BoW with SIFT + histogram

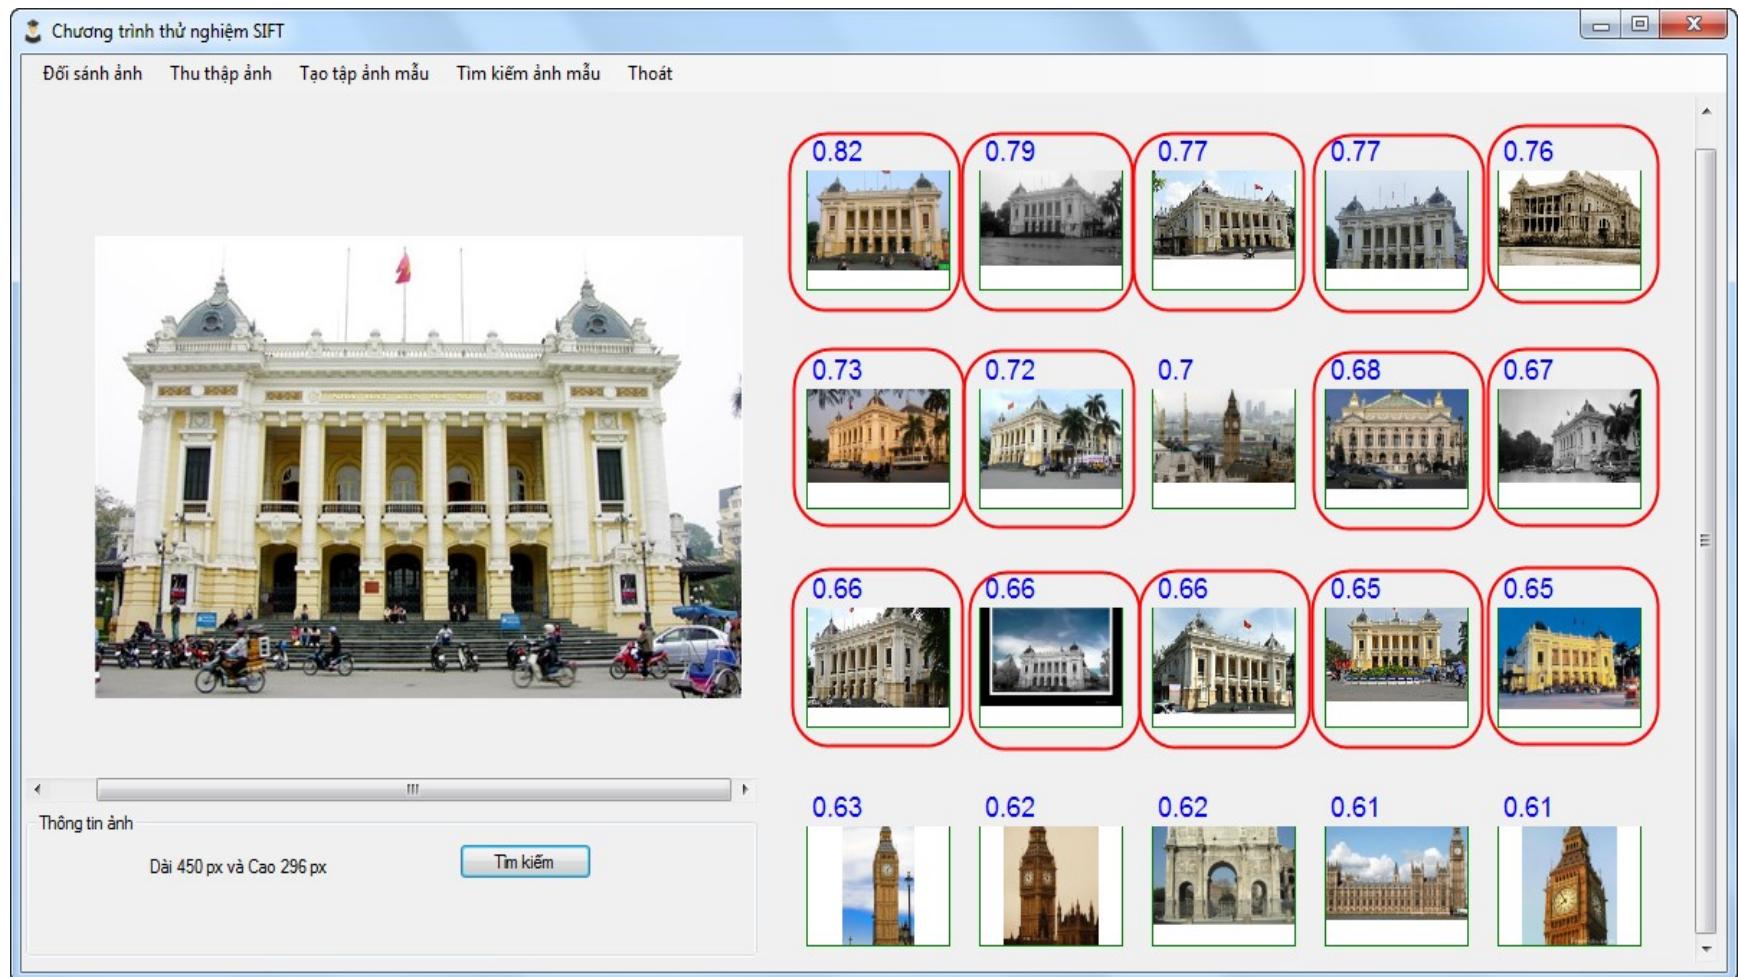


Offline



Source :Đồ án tốt nghiệp – Phạm Xuân Trường K52 - BK

CBIR: BoW with SIFT + histogram



Source :Đồ án tốt nghiệp – Phạm Xuân Trường K52 - BK

References

- Lecture 5,6: CS231 - Juan Carlos Niebles and Ranjay Krishna, Stanford Vision and Learning Lab
- Vision par Ordinateur, Alain Boucher, IFI



25
YEARS ANNIVERSARY
SOICT

VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

**Thank you for
your attention!**

